

# FRI Soundness Above the Johnson Radius via Threshold Halving

Raullen Chai<sup>\*1</sup> and Xinxin Fan<sup>†1</sup>

<sup>1</sup>IoTeX Network

April 29, 2026

## Abstract

The classical 1:2 proximity-gap inequality of Rothblum–Vadhan–Wigderson [31] (STOC 2013) holds for any linear code, in any characteristic, at any distance. We observe that this inequality, applied above the Johnson radius, places FRI’s effective distance into BCIKS’s zero-loss regime after a single round — the algebraic identity is  $\delta/2 < (1-\rho)/2 \leq 1 - \sqrt{\rho} = \delta_J$  for every  $\delta \in (\delta_J, 1-\rho)$  (Remark 1) — where the proven proximity gap of BCIKS [7] supplies the round-by-round bound. The result is the first unconditional FRI soundness theorem in the open zone  $\delta \in (\delta_J, 1-\rho)$ , at the cost of a  $\sim 2\times$  query overhead: for  $\text{RS}[F, L, k]$  with  $n = |L|$ ,  $\rho = k/n$ ,  $k = 2^m$ , and  $L$  admitting a fixed-point-free involution, FRI with  $R = m$  rounds and  $q$  base-level queries (each opening one full path) satisfies

$$\varepsilon_{\text{FRI}} \leq \frac{nR}{|F|} + \left(1 - \frac{\delta}{2}\right)^q.$$

We further show that this  $2\times$  factor is intrinsic to the standard two-bad- $\gamma$  correlated-agreement (CA) framework at FRI scale: the equal-threshold CA upper bound is  $\varepsilon_{\text{ca}}(C, \delta, \delta) \leq \binom{n}{w}/|F|$  unconditionally for  $w = \lfloor \delta n \rfloor$  (Theorem 8), with a matching tightness construction at the equality case  $\delta = (n - k - 1)/n$  for  $|F| > \binom{n}{w}^2$  (Proposition 10; see Remark 12 for scope). The bound is vacuous for  $n \gg \log |F|$ . As a structural complement, we give the first  $p$ -dependent moment bounds on the list size at all codimension excesses ( $\mathbb{E}[M] = \binom{n}{w}/p^c$  with Poisson dispersion, Corollary 30), an explicit obstruction to the worst-case  $M = 0$  form of OP2 at FRI scale (Proposition 35), and a conjectural linear- $n$  phase-diagram bound at  $c \geq 3$  (Conjecture 43, supported by exhaustive search up to  $n = 40$ ). Coverage: theorem-level for multiplicative FRI in odd characteristic (Theorem 20), additive FRI in characteristic 2 for the plain FRI protocol (Theorem 64), circle FRI / Stwo (Theorem 72), and STIR (Theorem 48); the WHIR analogue is conjectural (Conjecture 54; the proximity-layer substitution is rigorous via Lemma 55, but the iteration-level induction of WHIR [11, §5] would need to be re-run with the substituted folding term to upgrade the conjecture to a theorem) and the characteristic-2 STIR/WHIR analogue is sketched in Remark 66; both are listed as follow-up open tasks. Binius / binary-tower-field STIR/WHIR therefore requires the sketch to be filled in, which we identify as the main remaining char-2 task in §9.<sup>1</sup>

---

<sup>\*</sup>raullen@iotex.io

<sup>†</sup>xinxin@iotex.io

<sup>1</sup>The half-threshold CA theorem (Theorem 6) and the equal-threshold CA upper bound (Theorem 8) are formalized in Lean 4 with Mathlib (zero `sorry`, zero project-level axioms beyond Mathlib’s standard ones); the half-threshold theorem also includes a paper-form  $|F|$ -probability corollary. Building blocks for the FRI coupling, char-2 (additive

**Keywords:** Reed–Solomon codes, FRI protocol, proximity gap, correlated agreement, Johnson radius, STARK soundness, interactive oracle proofs.

**MSC 2020:** 94B35 (primary), 11T71, 68Q10.

# 1 Introduction

## 1.1 The problem

**Why this matters for Ethereum.** Ethereum’s medium-term scaling and post-quantum roadmap rests on succinct proofs: zkEVMS, zkVMs, and a planned post-quantum signature-aggregation layer. EthProofs<sup>2</sup> alone tracks  $\sim 30$  such systems in active use or near-term rollout. Almost all of them are STARK-based, and almost all STARKs are built on a single low-level cryptographic component: the FRI proximity test (or its successors STIR and WHIR). FRI’s soundness is what stops a malicious prover from convincing the chain — and through it, every smart contract and bridge — of a statement that is in fact false. To meet performance budgets (proof size, prover memory, verifier gas), many deployed systems run FRI in the so-called *open zone above the Johnson radius*: the parameter regime where, until now, the soundness analysis depended on an *unproved* conjecture (the “correlated agreement above the Johnson radius” conjecture, originating in Ben-Sasson–Carmon–Ishai–Kopparty–Saraf [7] and articulated as an open problem by Arnon–Boneh–Fenzi [4]). In late 2025 the strongest form of that conjecture — correlated agreement up to capacity — was disproved [5, 12]. This left every above-Johnson soundness argument used in production technically *unproven*: a positive security number quoted in audit reports, but with no theorem behind it in the relevant regime.

This paper closes that gap. We give an unconditional soundness theorem in the full open zone, with no change to protocol structure — only the verifier’s query parameter is recalibrated. Concretely, the system designer pays approximately one extra query round (§1.6) and, in return, gets a proven soundness bound whose validity does not rely on the equal-threshold conjecture or its refinements (including any sharper version of the Crites–Stewart construction). Operators of SP1, RISC Zero, Plonky3, Stwo, and similar systems can replace the conjectural baseline they currently assume with a theorem; this paper supplies the theorem and quantifies the recalibration cost (Table 1, §1.6, §9).

**Formal statement of the prize and the prior gap.** The Ethereum Foundation Proximity Prize<sup>3</sup> [1] targets the Reed–Solomon proximity gap conjectures of [4] that control the soundness of FRI [7], STIR [10], and WHIR [11]. BCHKS [6] proved zero-loss proximity gaps for  $\delta \leq \delta_J$  with  $O(n)/|F|$  per-round error (equivalently  $O(Rn)/|F|$  over  $R$  rounds, as in Table 1). Goyal–Guruswami [8] achieved optimal proximity gaps up to capacity for subspace-design codes and for *random* plain RS codes; Jeronimo–Liu–Rajpal [9] do the same for *folded* RS codes. For plain RS codes above the Johnson radius on a *fixed* multiplicative-coset evaluation domain  $L$  — and therefore neither random nor folded — no positive result existed.

---

coupling and recovery), and batch-CA (per-coordinate union) chains are present in the library; the concrete FRI-domain instantiations, a faithful BCIKS ’20 Theorem 1.2 import, and the WHIR / circle / Fiat–Shamir chains are still on the roadmap — see the companion repository’s `paper1/lean/STATUS.md` for the per-theorem state. Lean sources, Python verification scripts, and reproduction commands are in the companion repository [39]. The practical context and concrete bit-security tables are in §1 and §9.

<sup>2</sup><https://ethproofs.org>

<sup>3</sup><https://proximityprize.org>

The baseline is the conjecture  $\varepsilon_{\text{ca}}(C, \delta, \delta) = O(1)/|F|$  for all  $\delta < 1-\rho$  [4], disproved up-to-capacity by Crites–Stewart [5] (ePrint 2025/2046, November 2025); independently, Kambiré [12] (arXiv 2604.09724, 2026) gives a near-capacity prime-field counterexample at distance  $O(1/\log n)$  below capacity. For the intermediate zone ( $\delta_J < \delta < 1-\rho$ ): the bound  $\varepsilon_{\text{ca}}(C, \delta, \delta) \leq \binom{n}{w}/|F|$  (Theorem 8, tight by Proposition 10) is formally  $O(1)/|F|$  but with constant exponential in the code length, making it vacuous at FRI scale. Without the contribution of this paper, the above-Johnson FRI setting on Ethereum has zero proven soundness in the open zone.

## 1.2 What we deliver

**The structural observation.** The half-threshold correlated-agreement bound (Theorem 6) is *not new*: it is the classical 1:2 proximity-gap inequality of Rothblum–Vadhan–Wigderson [31] (STOC 2013), holding for any linear code over any field at any distance. What we contribute is the observation that, applied above the Johnson radius, this classical inequality has a previously-unnoticed consequence: a single round of threshold halving lowers the effective distance from  $\delta$  to  $\delta/2$ , and the chain  $\delta/2 < (1-\rho)/2 \leq 1-\sqrt{\rho} = \delta_J$  (Remark 1) places every subsequent round into BCIKS’s zero-loss regime, where the zero-loss BCIKS proximity gap [7] locks the distance for the rest of the protocol. The bridge is one line of algebra; converting it into an end-to-end deployment theorem requires three further structural results, each new in this paper: a matched-pair equal-threshold CA upper and lower bound proving that the  $2\times$  overhead is *intrinsic* to any CA-based argument at FRI scale (Theorem 8, Proposition 10); the first  $p$ -dependent moment bounds and sub-Poisson tail on the list size at all codimension excesses (§7); and theorem-level coverage of every FRI variant currently deployed — multiplicative, additive char-2, circle, and STIR (Theorems 64, 72, 48). Each is detailed in §1.4.

**The bound.** For FRI (Theorem 20; Theorem 64 for  $\text{char}(F) = 2$ ) on plain Reed–Solomon codes over any finite field, with  $n = |L|$ ,  $\rho = k/n$ ,  $k = 2^m$ , domain  $L$  admitting a fixed-point-free involution,  $R = m$  rounds, and  $q$  base-level queries (each opening one full path), for every  $\delta$  in the open zone ( $\delta_J, 1-\rho$ ):

$$\varepsilon_{\text{FRI}} \leq \frac{nR}{|F|} + \left(1 - \frac{\delta}{2}\right)^q. \quad (1)$$

This is an unconditional theorem in a regime where, prior to this paper, above-Johnson FRI relied on the now-disproved zero-loss CA conjecture [4, 5, 12]. No change to the FRI protocol structure, prover messages, or verifier checks — only the soundness analysis and the query parameter  $q$  it calibrates. Increasing  $q$  is still a real implementation-parameter change: it increases Merkle openings, consistency checks, proof bytes, and wrapper-verifier work. The same bound holds for FRI in characteristic 2 via additive folding (Appendix B); a batch-CA theorem extends it to STIR (Theorem 48) and WHIR (Corollary 50, Conjecture 54) in the odd-characteristic setting of those formal statements (the characteristic-2 STIR/WHIR extension is sketched in Remark 66); non-interactive knowledge soundness follows via round-by-round soundness and Fiat–Shamir compilation (Appendix D).

**Why it works.** Prior results [7, 6] proved soundness only at or below the Johnson radius, where zero-loss CA ( $\delta_{\text{hd}} = \delta$ ) holds. Extending to the open zone above Johnson was blocked by the equal-threshold CA conjecture [4], which requires  $M = 0$  on all RS-compatible flats — known to be false in its strongest form (Proposition 35). Threshold halving sidesteps the conjecture entirely: by accepting a  $\sim 2\times$  query overhead — which we further show is intrinsic to any CA-based proof

at current code lengths (Theorem 8, Remark 14) — it yields unconditional soundness in the full open zone, immune to any future open-zone counterexample.

**Why this wasn’t proven earlier.** The bridge  $\delta/2 < (1-\rho)/2 \leq \delta_J$  is one line of algebra and could in principle have been written down any time after BCIKS ’20. It was not, for two reasons. First, the community was tracking the conjectural up-to-capacity bound  $\varepsilon_{\text{ca}}(C, \delta, \delta) = O(1)/|F|$  [4], against which the  $2\times$  query overhead of threshold halving was unattractive. Second, the equal-threshold CA upper bound  $\binom{n}{w}/|F|$  (Theorem 8) and its matching tightness (Proposition 10) — the structural barrier showing that the  $2\times$  overhead cannot be reduced inside the CA framework, a fact not in the prior literature — left open the possibility that a sharper equal-threshold bound was within reach without paying the halving toll. After Crites–Stewart and Kambiré disproved the strongest up-to-capacity form in late 2025 [5, 12], every above-Johnson FRI deployment was left with neither an unconditional bound nor a definitive negative result on the equal-threshold side. This paper closes both: the unconditional theorem in the open zone, and the structural barrier that explains why the  $2\times$  price is the right one.

### 1.3 Coverage of current FRI variants

The theorem applies to the following settings, all operating above the Johnson radius:

- **Odd characteristic.** SP1 (BabyBear<sub>4</sub>, rate 1/2), RISC Zero (BabyBear, rate 1/4), Plonky3 (BabyBear / Koala-Bear, Polygon zkEVM), and comparable multiplicative-FRI zkVM systems.
- **Circle FRI (covered).** Stwo (Starkware, Mersenne31) uses circle FRI [22] on the unit circle  $\{(x, y) \in \mathbb{F}_p^2 : x^2 + y^2 = 1\}$ , whose folding structure differs from the multiplicative and additive cases (the involution is  $(x, y) \mapsto (x, -y)$ , not negation on a subgroup). Appendix C proves the required coupling lemma for circle codes and establishes FRI soundness (Theorem 72) with the analogous  $NR/|F| + (1-\delta/2)^q$  bound, where  $N = |D|$  is the size of the circle evaluation domain (the analogue of  $n = |L|$  in the multiplicative case).
- **Characteristic 2.** Systems using additive FRI on  $\mathbb{F}_2$ -linear subspaces (binary tower fields  $\mathbb{F}_{2^{128}}$ ). Our Theorem 64 covers the basic additive fold  $s(x) = x^2 + \beta x$ . Binius [21] uses a more elaborate block-decomposition/Gao-Mateer recursion that is compatible with — but distinct from — this basic fold.

Protocol coverage:

- **FRI (proved).** Theorem 20 ( $\text{char} \neq 2$ ) and Theorem 64 ( $\text{char} = 2$ ).
- **STIR (theorem-level in odd characteristic).** Theorem 48 proves STIR soundness above Johnson by applying the half-threshold CA bound (Theorem 6) to the round- $i$  batched function  $q_i + \gamma_i \cdot f^{(i-1)} \upharpoonright_{L^{(i)}}$  together with the BCIKS proximity gap at later rounds; the characteristic-2 analogue is sketched in Remark 66.
- **WHIR (conjectural in odd characteristic).** Corollary 50 proves our CA bound holds for any *linear* subcode of  $\text{RS}_k$ , and end-to-end WHIR soundness with sumcheck composition is conjectured in Conjecture 54; the characteristic-2 analogue is sketched in Remark 66.

The equal-threshold bound  $\binom{n}{w}/|F|$  is vacuous at FRI scale (§1.7); our half-threshold bound is optimal within the standard two-bad- $\gamma$  CA framework, and whether a non-CA argument or one tracking  $\geq 3$  bad  $\gamma$ ’s could beat the 1:2 ratio is left open in §8.

## 1.4 Technical contributions

- (i) **FRI soundness theorem above Johnson** (Theorem 20; Theorem 64 for characteristic 2; Theorem 72 for circle FRI, with  $N = |D|$  replacing  $n = |L|$ ): the bound  $nR/|F| + (1-\delta/2)^q$  in the open zone  $(\delta_J, 1-\rho)$ , via the structural observation (2) that  $\delta/2 < (1-\rho)/2 \leq \delta_J$  bridges the classical RVW13 inequality [31] (Theorem 6) to BCIKS’s zero-loss regime [7]. The commit-phase term  $O(nR/|F|)$  matches BCHKS up to constants; the contribution is that this is *proved unconditionally*, not conjectured. An  $O(1)$  round-1 proximity gap (Theorem 16,  $\leq 1$  bad  $\alpha$ ) is the local consequence used in the proof.
- (ii) **Equal-threshold CA: a new exact upper bound paired with a new tightness witness** (Theorem 8, Proposition 10, Remark 14): we prove the unconditional upper bound  $\varepsilon_{\text{ca}}(C, \delta, \delta) \leq \binom{n}{w}/|F|$  for  $w = \lfloor \delta n \rfloor$  via an injection of bad coefficient values into  $\binom{[n]}{n-w}$ , and we construct a matching lower-bound witness at the equality case  $\delta = (n - k - 1)/n$  for  $|F| > \binom{n}{w}^2$  (Proposition 10; see Remark 12 for the  $|F|$ -asymptotic scope). To our knowledge, neither side of this pair has been stated in the prior literature. Together they establish a structural barrier within the standard two-bad- $\gamma$  CA framework: at FRI scale  $n \gg \log |F|$  the upper bound is exponentially larger than 1 and therefore vacuous against any meaningful security target, yet over large fields it is attained at the equality case, certifying that the  $2\times$  query overhead of (i) cannot be removed within this framework. Whether arguments tracking  $\geq 3$  bad  $\gamma$ ’s or non-CA proximity proofs can do so is left open (§8).
- (iii) **List-size phase diagram and OP2 obstruction** (§7): the first  $p$ -dependent moment bounds on the list size —  $\mathbb{E}[M] = \binom{n}{w}/p^c$ ,  $\text{Var} \approx \mathbb{E}$  at every codimension excess  $c \geq 1$  (Theorem 25, Theorem 28, Corollary 30) — via an error-locator normal-vector identity (Lemma 23); a sub-Poisson Bernstein tail at  $c=2$  via a Möbius-induced negative-correlation argument (Theorem 41, Lemma 39); a codimension phase diagram (§7.6) with a linear- $n$  conjecture  $M_{\text{true}} \leq \lfloor (2D-1)/c \rfloor$  at  $c \geq 3$  (Conjecture 43, empirically verified through  $n = 40$ , the large- $c$  regime relevant to FRI). Explicit obstruction:  $M_{\text{max}} \geq 2$  at FRI parameters for all  $p$  (Proposition 35), refuting the strongest  $M = 0$  form of OP2.
- (iv) **Extensions** (Appendices B, C, A, D): characteristic-2 FRI via additive subspace-polynomial folding (Theorem 64); circle FRI for Stwo’s Mersenne31 setting (Theorem 72, with Lemma 70); a batch-CA theorem (Theorem 46) yielding STIR (Theorem 48) and WHIR (Conjecture 54) soundness above Johnson in odd characteristic, with the characteristic-2 analogue sketched in Remark 66; round-by-round soundness (Theorem 75), Fiat–Shamir-compiled non-interactive FRI (Theorem 76), and DEEP-FRI knowledge soundness (Theorem 78, Corollary 79).

## 1.5 Comparison

## 1.6 Concrete cost

**Against what is actually proved today**, our bound provides positive proven security at every  $\delta$  in the open zone where BCHKS provides none: above the Johnson radius, BCHKS gives zero proven soundness, whereas we give  $\varepsilon_{\text{FRI}} \leq nR/|F| + (1 - \delta/2)^q$ . There is no proved-vs-proved regression; there is only added coverage.

Against the *conjectured* zero-loss CA baseline ( $\varepsilon_{\text{ca}}(C, \delta, \delta) = O(1)/|F|$  above Johnson per ABF [4]; the only currently *proved* equal-threshold bound is the unconditional upper bound  $\varepsilon_{\text{ca}}(C, \delta, \delta) \leq \binom{n}{w}/|F|$  of Theorem 8, exponential in  $n$  at FRI scale — see Remark 14 and §8,

| Result            | Regime                       | FRI commit error                                                | Loss       | Ref.       |
|-------------------|------------------------------|-----------------------------------------------------------------|------------|------------|
| BCIKS [7]         | $\delta < \delta_J$          | $O(Rn^2)/ F $                                                   | 0          | FOCS'20    |
| BCHKS [6]         | $\delta \leq \delta_J$       | $O(Rn)/ F $                                                     | 0          | 2025       |
| Goyal–Gur. [8]    | $\delta < 1-\rho$            | $O(R)/ F $<br>( <i>random plain RS; subspace-design codes</i> ) | 0          | 2025       |
| <b>This paper</b> | $\delta_J < \delta < 1-\rho$ | $\mathbf{O(nR)}/ F $                                            | $\delta/2$ | <b>new</b> |

Table 1: FRI commit-phase soundness for **plain fixed- $L$**  RS (the fixed-domain setting). Goyal–Guruswami applies to random plain RS and subspace-design codes, not to this fixed- $L$  setting; circle FRI is covered in Appendix C.

| Field ( $ F $ )                     | Commit<br>(both) | Query ( $q=128, \delta=0.4$ ) |                                | Proved total |
|-------------------------------------|------------------|-------------------------------|--------------------------------|--------------|
|                                     |                  | Conj. $(1-\delta)^q$          | <b>Proved</b> $(1-\delta/2)^q$ |              |
| BabyBear base ( $2^{31}$ )          | 7                | $94$ ( <i>unproved</i> )      | 41                             | <b>7</b>     |
| BabyBear <sub>4</sub> ( $2^{124}$ ) | 100              | $94$ ( <i>unproved</i> )      | 41                             | <b>41</b>    |
| Goldilocks ( $2^{64}$ )             | 40               | $94$ ( <i>unproved</i> )      | 41                             | <b>40</b>    |

Table 2: Interactive FRI security (bits),  $n = 2^{20}$ ,  $\rho = 1/2$ ,  $R = 20$ . The commit-phase term matches BCHKS [6]. Prior to this work, the 94-bit query-phase figure was *conjectured*; we *prove* 41 bits at the same  $q = 128$ . Proved total =  $\min(\text{commit}, \text{proved query})$ : on BabyBear<sub>4</sub> the query phase is the bottleneck and our proved total is 41 bits; on Goldilocks and base-field BabyBear the commit phase dominates. Increasing to  $q \approx 398$  achieves 128-bit *interactive* query-phase security; BabyBear<sub>4</sub> remains capped at  $\sim 100$  total interactive bits by the commit term. Table 3 accounts separately for Fiat–Shamir’s non-interactive  $Q$ -loss. Circle FRI (Stwo/Mersenne31) achieves the same bound (Theorem 72).

item 1 — and the strongest up-to-capacity form was independently disproved [5, 12]), the half-threshold exacts a precise toll on query complexity. To achieve  $\lambda$ -bit security from the query term alone:

$$q_{\text{zero-loss}} = \frac{\lambda}{\log_2(1/(1-\delta))}, \quad q_{\text{ours}} = \frac{\lambda}{\log_2(1/(1-\delta/2))}.$$

The overhead ratio  $q_{\text{ours}}/q_{\text{zero-loss}} = \log_2(1/(1-\delta))/\log_2(1/(1-\delta/2))$ :

| $\delta$ | Regime                   | $q_{\text{ours}}/q_{\text{zero-loss}}$ | Status of zero-loss                       |
|----------|--------------------------|----------------------------------------|-------------------------------------------|
| 0.293    | Johnson ( $\rho = 1/2$ ) | $2.2\times$                            | $\binom{n}{w}/ F $ (Remark 14)            |
| 0.40     | Intermediate             | $2.3\times$                            | $\binom{n}{w}/ F $ , vacuous at FRI scale |
| 0.45     | Near capacity            | $2.4\times$                            | Disproved at capacity [5, 12]             |

The overhead is paid only in queries, and only against a baseline that is itself unproven. The commit-phase term  $O(nR)/|F|$  matches BCHKS up to constants; we do not claim a Pareto improvement over BCHKS in the Johnson regime. The contribution is coverage of the open zone, where BCHKS provides no bound. Section 9 quantifies both terms in concrete parameters.

## 1.7 The correlated agreement framework and its limits

Every known FRI soundness proof — BBHR [20], BCIKS [7], BCHKS [6], Goyal–Guruswami [8], and the present paper — follows the same two-step strategy:

1. **Correlated agreement (CA) lemma:** if  $(f_1, f_2)$  has joint distance  $> \delta$  from  $C^{\equiv 2}$ , then for random  $\gamma$ , the fold  $f_1 + \gamma f_2$  stays far from  $C$  with high probability.
2. **Coupling:** the FRI decomposition  $f \mapsto (f_{\text{even}}, f_{\text{odd}})$  maps distance from  $\text{RS}_k$  to joint distance from  $\text{RS}_{k/2}^{\equiv 2}$ , linking the CA lemma to the protocol.

We call this the **CA framework**. The query-phase soundness of FRI is controlled by the *conclusion threshold*  $\delta_{\text{hd}}$  of the CA lemma: each query catches a far fold with probability  $\delta_{\text{hd}}$ , giving soundness  $(1 - \delta_{\text{hd}})^q$  after  $q$  queries. The ratio  $\delta_{\text{hd}}/\delta$  determines the *query overhead*: at ratio 1:1 (zero loss), queries are optimal; at ratio 1: $r$ , the overhead is  $\sim r \times$ .

**The standard two-bad- $\gamma$  CA framework is settled at the 1:2 ratio.** The  $\delta_{\text{hd}} = \delta/2$  bound (ratio 1:2) was established by Rothblum–Vadhan–Wigderson [31] in 2013 for arbitrary linear codes (Theorem 6). At equal threshold  $\delta_{\text{hd}} = \delta$ :  $\varepsilon_{\text{ca}} \leq \binom{n}{w}/|F|$  with constant exponential in the code length (Remark 14), vacuous against any meaningful security target at FRI scale. The equal-threshold upper bound of Theorem 8 is exponential in  $n$ , suggesting that significantly tighter CA bounds at FRI scale via the two-bad- $\gamma$  argument will require additional structure beyond linearity; we are not aware of any two-bad- $\gamma$  CA argument that beats the 1:2 ratio at deployment scale.

**What lies beyond?** Reducing the overhead below  $2 \times$  would require either a CA argument tracking  $\geq 3$  bad  $\gamma$ 's (open, see §8) or a proof that does not factor through a CA lemma at all. No such argument exists in the public literature. We discuss concrete candidates — character sums on multiplicative subgroups, cross-correlation techniques — in Section 10.

## 1.8 Scope and limitations

To prevent misreading: this paper (i) proves an unconditional FRI soundness theorem above Johnson with a  $\sim 2 \times$  query overhead that is optimal within the CA framework (§1.7); (ii) is *not* a protocol change — FRI/STIR/WHIR are unchanged; (iii) is *not* free — the  $\sim 2 \times$  overhead vs. the zero-loss baseline is the price of an unconditional proof.

**Structural assumptions.** The  $k = 2^m$  condition is necessary: for odd  $k \geq 3$  the monomial  $x^k$  gives a counterexample with  $\Delta(x^k, \text{RS}_k) = 1 - \rho$  yet the FRI fold is always a codeword (Remark 5). Standard FRI uses  $k = 2^m$  already. The second-moment bound (Theorem 25) holds for all  $k$  but at codimension  $c = 1$ ; it is not on the critical path to the main FRI theorem. There is *no* characteristic restriction: the main body handles  $\text{char}(F) \neq 2$  via multiplicative folding; Appendix B extends to  $\text{char}(F) = 2$  via additive folding, with the identical  $nR/|F| + (1 - \delta/2)^q$  bound. **There is also no rate restriction:** the main theorem holds for every rate  $\rho = k/n$  with  $k = 2^m$  (the numerical tables in §1.6 use  $\rho = 1/2$  as the BabyBear/Goldilocks reference; rates such as  $\rho = 1/4$  or  $\rho = 1/8$ , common in zkVM parameter choices for soundness amplification, satisfy the same bound with  $R = m$  rounds and any  $\delta \in (\delta_J(\rho), 1 - \rho)$ ).

## 1.9 Relation to the EF Proximity Prize

The Ethereum Foundation Proximity Prize [1] (May 2026) frames the open problem as two challenges. The *Grand mutual correlated agreement (MCA) Challenge*: determine the largest  $\delta^*$  such that  $\varepsilon_{\text{mca}}(C, \delta^*) \leq \varepsilon^*$  with  $\varepsilon^* = 2^{-128}$  on Reed–Solomon codes at rates  $\rho \in \{1/2, 1/4, 1/8, 1/16\}$ . The *Grand List Decoding Challenge*: same target on  $|\Lambda(C^{\equiv m}, \delta^*)| \leq \varepsilon^*|F|$  for constant  $m$ . Both subsume the ABF open problems [4]: OP1 (zero-loss CA above Johnson) feeds the MCA target, OP2 (sub- $O(n)$  list-size route) feeds the LD target, OP3 (improved CA constants) sharpens both.

The table below is the claim map used in the rest of the paper. It separates theorem-level results from partial progress and from the original zero-loss targets that remain open. Throughout

the table, “above Johnson” means  $\delta \in (\delta_J, 1-\rho)$  (the open zone, strictly above the Johnson radius and strictly below capacity); the upper limit is implicit in the open-zone framing of §2.

| Prize target                                                                      | Paper result                                                                                                                                                                                                                                                                                           | Status                                                                                               | Reference                                                           | Remaining gap                                                               |
|-----------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------|-----------------------------------------------------------------------------|
| Grand MCA: equal-threshold CA above Johnson                                       | Unconditional upper bound $\varepsilon_{ca}(C, \delta, \delta) \leq \binom{n}{w}/ F $ (Theorem 8); matching tightness construction at the equality-case $w = n-k-1$ for $ F  > \binom{n}{w}^2$ (Proposition 10). Constant is exponential in $n$ at FRI scale                                           | Refuted in the asymptotic large- $n$ form; bound and (equality-case) tightness now in the literature | Theorem 8; Proposition 10; Remark 14                                | General- $w$ tightness; non-CA proof beating the $2\times$ overhead         |
| Grand MCA: half-threshold CA, all rates incl. $\rho = 1/16$                       | Unconditional $\varepsilon_{ca}(C, \delta/2, \delta) \leq 1/ F $ , rate-uniform; yields theorem-level FRI/STIR soundness and conjecture-level WHIR soundness (proximity-substitution proved) with common query term $(1 - \delta/2)^q$ and protocol-specific commit-phase terms, at no protocol change | Solved as theorem (rate-uniform) for FRI/STIR; conjectural for WHIR                                  | Theorem 6; Theorems 20, 64, 72, 48; Conjecture 54                   | $\sim 2\times$ query overhead vs. the conjectured zero-loss baseline        |
| Grand LD: $ \Lambda(C^{\equiv m}, \delta^*)  \leq \varepsilon^* F $ above Johnson | First $p$ -dependent moment bounds; sub-Poisson tail at $c = 2$ with high probability over $(s_1, s_2)$ ; codimension phase diagram with linear- $n$ conjecture at $c \geq 3$ (the large- $c$ regime relevant to FRI); explicit obstruction to worst-case $M = 0$                                      | Partially addressed; strongest $M = 0$ form refuted; phase diagram conjectural at $c \geq 3$         | Theorem 25; Corollary 30; Theorem 41; Conjecture 43; Proposition 35 | Worst-case $c = 2$ extension; rigorous proof of Conjecture 43 at $c \geq 3$ |
| ABF OP3: improved CA constants                                                    | Half-threshold CA constant is $1/ F $ and commit phase matches $O(nR)/ F $ up to constants                                                                                                                                                                                                             | Partially addressed                                                                                  | Theorem 6; Theorem 20                                               | Does not prove zero-loss CA                                                 |
| Original up-to-capacity MCA / zero-loss proximity gap                             | Not proved; capacity-level variants are now known false in related settings                                                                                                                                                                                                                            | Not solved here                                                                                      | Crites–Stewart [5]; Kambiré [12]; §8                                | Requires different techniques or revised conjectures                        |

**Interpretation.** The paper does not claim the original zero-loss proximity gap. Its prize-facing contribution is narrower and unconditional: FRI/STIR get theorem-level soundness above Johnson with a half-threshold CA safety net at *every* rate  $\rho \in \{1/2, 1/4, 1/8, 1/16\}$  (and WHIR gets conjecture-level soundness via the same chain, with the proximity-layer substitution rigorously proved), while the equal-threshold CA route is shown to be vacuous at FRI scale. The remaining route to better-than- $2\times$  query overhead must therefore avoid the CA framework or prove a strictly stronger non-CA proximity statement.

**Companion work in preparation.** A pair of companion papers in preparation addresses non-CA proximity statements complementary to this paper’s CA-framework boundary [2, 3]; concrete result statements await those papers’ completion and are not asserted here. See §10 for a brief pointer.



## 1.10 Guide to the proof

The paper has two logically independent tracks: the **mainline** (§3–§6), which proves the FRI soundness theorem in five steps, and the **structural track** (§7), which establishes the first  $p$ -dependent list-size bounds and resolves Open Problem 2 (negatively). The mainline is entirely self-contained; the structural track is not on the critical path to any soundness theorem.

### Mainline: FRI soundness in five steps.

- Step 1. Half-threshold CA** (Theorem 6, §3). For any linear code  $C$ : if  $(f_1, f_2)$  has joint distance  $> \delta$  from  $C^{\perp=2}$ , at most one  $\gamma \in F$  makes  $f_1 + \gamma f_2$  within distance  $\delta/2$  of  $C$ . *Input*: linearity of  $C$  only (RVW13 [31]).
- Step 2. Even/odd coupling** (Lemma 15, §5). For even  $k$ : the distance from  $\text{RS}_k$  is at most the joint distance from  $\text{RS}_{k/2}^{\perp=2}$ . *Input*: the isomorphism  $\text{RS}_k \xrightarrow{\sim} \text{RS}_{k/2}^{\perp=2}$  (requires  $k$  even; Remark 5).
- Step 3. Proximity gap** (Theorem 16, §5). At most 1 value of  $\alpha$  makes the FRI fold  $(\delta/2)$ -close to  $\text{RS}_{k/2}$ . *Input*: Steps 1 + 2 composed.
- Step 4. Distance locking** (Remark 21, §6). After round 1, the effective distance is  $\delta/2 < (1-\rho)/2 \leq \delta_J$  (Remark 1), inside BCIKS’s zero-loss regime. The BCIKS proximity gap [7] (Theorem 1.2, Eq. (1.1)) at threshold  $\gamma = \delta/2 < \delta_J$  gives  $\leq n$  bad  $\alpha$ ’s per round; distance stays at  $\delta/2$  for all subsequent rounds. *Input*: the chain (2) (always true for  $\delta < 1-\rho$ ) + BCIKS [7] (external, proved).
- Step 5. FRI soundness** (Theorem 20, §6). Strategy A (honest fold): multilinear Schwartz–Zippel gives  $\leq R/|F|$ . Strategy B (deviate): Steps 3–4 bound bad  $\alpha$ ’s to  $\leq 1$  (round 1) and  $\leq n$  (rounds  $\geq 2$ ) per round and the consistency check catches a  $(\delta/2)$ -far deviation with probability  $\geq 1 - (1-\delta/2)^q$ . Union:  $nR/|F| + (1-\delta/2)^q$ . *Input*: Steps 3 + 4 + Schwartz–Zippel.

The mainline uses *no* result from §7. The entire proof rests on three external ingredients: linearity of  $\text{RS}_k$  (trivial), the BCIKS proximity gap at  $\delta \leq \delta_J$  (proved [7]), and the Schwartz–Zippel lemma.

**Extensions (from the same CA bound).** Theorem 6 is characteristic-independent, so three extensions branch directly from Step 1 (see Figure 1):

- *STIR*: Batch CA (Theorem 46)  $\rightarrow$  STIR soundness (Theorem 48), Appendix A.
- *WHIR*: Sumcheck composition + half-threshold CA (Conjecture 54), Appendix A.
- *Characteristic 2*: Additive coupling (Lemma 62)  $\rightarrow$  FRI soundness in characteristic 2 (Theorem 64), Appendix B.
- *Circle FRI*: Circle coupling (Lemma 70)  $\rightarrow$  FRI soundness on unit circle (Theorem 72), Appendix C.

**Structural track (§7): independent of the soundness proof.** Section 7 studies the per-word list size  $M$  (how many codewords lie within distance  $w$  of a received word), aiming at Open Problem 2 (§8). Four results:

1. *First  $p$ -dependent moment bound.*  $\mathbb{E}[M] = \binom{n}{w}/p^c$  for all codimension excess  $c \geq 1$  (Corollary 30), via the error-locator normal vector identity (Lemma 27) and pairwise independence for non-overlapping supports (Theorem 28). At FRI-relevant codimension ( $c = \Theta(n)$ ),  $\mathbb{E}[M]$  is exponentially small.

2. *Worst-case landscape by codimension.* At  $c = 2$ , a Möbius identity (Lemma 39) yields a universal  $M_{\text{true}} \leq \min(p, 2\binom{n}{w-1})$  (Theorem 40) and a Bernstein-style sub-Poisson tail on bin sizes with high probability over  $(s_1, s_2)$  (Theorem 41); the empirical worst-case fits  $c_0 c_1^n$  with  $c_0 \approx 0.66$ ,  $c_1 \approx 1.36$  in the tested range  $n \leq 24$  (Remark 42; descriptive, not asymptotic). At  $c \geq 3$ , an Open-Set Rank Lemma (Conjecture 43) predicts a linear bound  $M_{\text{true}} \leq \lfloor (2D-1)/c \rfloor$ , empirically verified through  $n = 40$  (§7.6). The large- $c$  regime  $c = \Theta(n)$  falls in this third phase, where the conjecture predicts  $M_{\text{true}} = O(1)$  at the Johnson radius.
3. *OP2 obstruction.*  $M_{\text{max}} \geq 2$  at FRI parameters ( $n \geq 32$ , rate  $1/2$ ) for all  $p$  (Proposition 35). The strongest  $M = 0$  conjecture is false.
4. *Practical  $M = 0$ .* At deployment parameters ( $p \approx 2^{31}$ ,  $c = \Theta(n)$ ):  $\mathbb{E}[M] = \binom{n}{w}/p^c \ll 1$  (Corollary 30);  $M = 0$  in practice even though worst-case  $M = 0$  fails.

Additional structural results (fiber bounds, Bézout bounds, codimension bounds, incidence bounds, irreducibility of the remainder hypersurface) are developed in a companion note [24].

**Reader’s guide.** A reader interested only in the FRI soundness theorem may read §2–§6 and skip §7 entirely.

## 2 Definitions

Let  $F$  be a finite field and  $L \subset F$  with  $|L| = n$ . The FRI folding requires a *fixed-point-free involution* on  $L$ : a bijection  $\iota : L \rightarrow L$  with  $\iota^2 = \text{id}$  and  $\iota(x) \neq x$  for all  $x \in L$ . In the main body (§3–§6), we use  $\iota(x) = -x$  (i.e.,  $\text{char}(F) \neq 2$  and  $L$  closed under negation). Appendix B uses  $\iota(x) = x + \beta$  (i.e.,  $\text{char}(F) = 2$  and  $L$  an  $\mathbb{F}_2$ -linear subspace containing  $\beta \neq 0$ ).  $\text{RS}_k = \text{RS}[F, L, k]$  is the Reed–Solomon code of evaluations of polynomials of degree  $< k$  on  $L$ , with rate  $\rho = k/n$ . We write  $\delta_J = 1 - \sqrt{\rho}$  for the *Johnson radius* (the classical relative list-decoding radius with polynomial list size [15, 16]) and use this term consistently throughout.

**Open zone.** Throughout, the *open zone* refers to the strict interval  $\delta \in (\delta_J, 1 - \rho)$  — distances strictly above the Johnson radius and strictly below capacity. All end-to-end soundness theorems in this paper (Theorems 20, 64, 72, 48; Conjecture 54; the non-interactive variants of Appendix D) operate in this regime. The CA, equal-threshold, proximity-gap, and list-size theorems hold for any  $\delta$  within the technical hypotheses of each statement; the open-zone window enters only when these are composed into an end-to-end FRI/STIR/WHIR soundness statement. At the boundary  $\delta = \delta_J$  the proved zero-loss CA of BCHKS [6] applies and our reduced-threshold statement is strictly weaker.

**Remark 1** (Half-threshold lies inside BCIKS’s zero-loss regime). *The single algebraic identity that drives every soundness theorem in this paper is, for  $\delta < 1 - \rho$ ,*

$$\frac{\delta}{2} < \frac{1 - \rho}{2} \leq 1 - \sqrt{\rho} = \delta_J, \quad (2)$$

where the first inequality is  $\delta < 1 - \rho$  divided by two, and the second is the elementary identity  $(1 - \sqrt{\rho})^2 \geq 0 \Leftrightarrow 1 - \rho \leq 2(1 - \sqrt{\rho})$ . Hence after one halving step the threshold  $\delta/2$  lies strictly below the Johnson radius  $\delta_J$ , where the BCIKS zero-loss proximity gap [7, Theorem 1.2 / Eq. (1.1)] is established. We refer to this as “BCIKS’s zero-loss regime” (sometimes “the Johnson regime”); we deliberately avoid the term “unique-decoding regime” because the unique-decoding radius  $(1 - \rho)/2$  is strictly smaller than  $\delta_J$  for every  $\rho < 1$ , so (2) is a stronger statement than  $\delta/2$  being below half the minimum distance.

**Remark 2** (Base vs. extension fields). *The evaluation domain  $L$  typically sits inside a prime subfield or a small extension (e.g.,  $L \subset \mathbb{F}_p$  with  $|L| = n \mid (p-1)$ ). The soundness bounds in this paper depend on  $|F|$ , not on the subfield containing  $L$ . In deployed systems, FRI randomness  $\alpha_i$  is sampled from an extension field  $\mathbb{F}_{p^d}$  to achieve sufficient Schwartz–Zippel probability:  $|F| = p^d$ . For example, BabyBear ( $p = 2^{31} - 1$ ) with degree-4 extension gives  $|F| = p^4 \approx 2^{124}$ , yielding  $\log_2(|F|/(nR)) \approx 100$  bits of commit-phase security (at  $n = 2^{20}$ ,  $R = 20$ ). The  $|F|$  in Theorems 6, 20, and 48 refers to the field from which random challenges are drawn, which is the extension field in practice.*

**Definition 3** (Correlated Agreement [4]).  $\varepsilon_{\text{ca}}(C, \delta_{\text{hd}}, \delta_{\text{nt}}) = \max_{f_1, f_2: L \rightarrow F} \Pr_{\gamma \leftarrow F} [\Delta(f_1 + \gamma f_2, C) \leq \delta_{\text{hd}} \text{ and } \Delta_{\text{joint}}((f_1, f_2), C^{=2}) > \delta_{\text{nt}}]$ , where  $\gamma$  is drawn uniformly at random from  $F$ , the max is over all functions  $f_1, f_2 : L \rightarrow F$  (not necessarily codewords), and  $\Delta_{\text{joint}}((f_1, f_2), C^{=2}) = \min_{g_1, g_2 \in C} \frac{1}{|D|} |\{x \in D : f_1(x) \neq g_1(x) \text{ or } f_2(x) \neq g_2(x)\}|$  is the joint Hamming distance to the pair-domain code  $C^{=2}$  (equivalently, to the set product  $C \times C$ ) [7], where  $D$  is the common domain of  $f_1, f_2$  (always the domain of  $C$ ; when the CA bound is applied to  $(f_{\text{even}}, f_{\text{odd}})$  on  $L'$ ,  $D = L'$  and  $|D| = n/2$ ). We call  $\delta_{\text{hd}}$  the conclusion threshold (distance of the random combination to  $C$ ) and  $\delta_{\text{nt}}$  the premise threshold (joint distance from  $C^{=2}$ ). Standard zero-loss CA [4] corresponds to  $\delta_{\text{hd}} = \delta_{\text{nt}} = \delta$ .

**Definition 4** (FRI Folding).  $L' = \{x^2 : x \in L\}$  with  $|L'| = n/2$ . For each  $y \in L'$ , fix a square root  $\sqrt{y} \in L$  (the expressions below are invariant under the choice  $\sqrt{y} \mapsto -\sqrt{y}$ ). For  $f : L \rightarrow F$ :  $f_{\text{even}}(y) = (f(\sqrt{y}) + f(-\sqrt{y}))/2$ ,  $f_{\text{odd}}(y) = (f(\sqrt{y}) - f(-\sqrt{y}))/(2\sqrt{y})$ . FRI fold:  $f'(y) = f_{\text{even}}(y) + \alpha \cdot f_{\text{odd}}(y)$ .

**Remark 5** (Even/odd isomorphism for even  $k$ ). For  $k$  even:  $g \mapsto (g_{\text{even}}, g_{\text{odd}})$  is a linear isomorphism  $\text{RS}_k \xrightarrow{\sim} \text{RS}_{k/2}^{=2}$  (both have dimension  $k$ ; the map permutes the  $k$  coefficients). Consequently:  $f \notin \text{RS}_k$  iff  $(f_{\text{even}}, f_{\text{odd}}) \notin \text{RS}_{k/2}^{=2}$ .

For  $k$  odd:  $\text{RS}_{\lceil k/2 \rceil}^{=2}$  has dimension  $k+1 > k$ , and the isomorphism fails. The monomial  $x^k$  (for  $k \geq 3$  odd) satisfies  $\Delta(x^k, \text{RS}_k) = 1 - \rho$  (since  $x^k - g$  has at most  $k$  roots for any  $g$  of degree  $< k$ ), yet  $f_{\text{even}} = 0$  and  $f_{\text{odd}}(y) = y^{(k-1)/2} \in \text{RS}_{\lceil k/2 \rceil}$  (since  $(-1)^k = -1$ ), so the fold  $f' = \alpha \cdot y^{(k-1)/2} \in \text{RS}_{\lceil k/2 \rceil}$  for all  $\alpha$  and FRI always accepts. The degenerate case  $k = 1$  is vacuous:  $\text{RS}_1 = \{\text{constants}\}$  and FRI is not invoked ( $R = 0$ ). This is why we require  $k = 2^m$  with  $m \geq 1$ .

### 3 Half-Threshold CA Bound

**Theorem 6** (RVW13 [31]; BCHKS [6], Table 1, row 1). For any linear code  $C$  over  $F$  with  $|L| = n$ :  $\varepsilon_{\text{ca}}(C, \delta/2, \delta) \leq 1/|F|$ .

Equivalently: for any  $f_1, f_2 : L \rightarrow F$  with  $\Delta_{\text{joint}}((f_1, f_2), C \times C) > \delta$ , at most one  $\gamma \in F$  satisfies  $\Delta(f_1 + \gamma f_2, C) \leq \delta/2$ .

**Proof (due to [31]).** Two bad  $\gamma$ 's give two linear equations in  $(f_1, f_2)$  on overlapping agreement sets. Solving for both components simultaneously yields a pair  $(g_1, g_2) \in C \times C$  matching  $(f_1, f_2)$  on  $\geq (1 - \delta)n$  points, so joint distance  $\leq \delta$  — contradicting the premise.

*Proof.* Fix  $f_1, f_2$  with  $\Delta_{\text{joint}} > \delta$ . Call  $\gamma$  “bad” if  $\exists h_\gamma \in C$  with  $\Delta(f_1 + \gamma f_2, h_\gamma) \leq \delta/2$ ; let  $S_\gamma = \{x \in L : f_1(x) + \gamma f_2(x) = h_\gamma(x)\}$ , so  $|S_\gamma| \geq (1 - \delta/2)n$ .

Suppose  $\gamma_1 \neq \gamma_2$  are both bad. By inclusion-exclusion:

$$|S_{\gamma_1} \cap S_{\gamma_2}| \geq 2(1 - \delta/2)n - n = (1 - \delta)n.$$

On  $S_{\gamma_1} \cap S_{\gamma_2}$ , the system  $f_1 + \gamma_i f_2 = h_i$  ( $i = 1, 2$ ) has the unique solution

$$g_1 := \frac{\gamma_2 h_1 - \gamma_1 h_2}{\gamma_2 - \gamma_1} \in C, \quad g_2 := \frac{h_1 - h_2}{\gamma_1 - \gamma_2} \in C$$

(using linearity of  $C$ ). So  $(f_1, f_2) = (g_1, g_2)$  on  $S_{\gamma_1} \cap S_{\gamma_2}$ , giving:

$$\Delta_{\text{joint}}((f_1, f_2), C \times C) \leq \frac{|L \setminus (S_{\gamma_1} \cap S_{\gamma_2})|}{n} \leq \delta.$$

This contradicts  $\Delta_{\text{joint}} > \delta$ . Hence at most 1 bad  $\gamma$ .  $\square$

**Remark 7** (Attribution and scope). *Theorem 6 is the classical 1:2 proximity-gap inequality of Rothblum–Vadhan–Wigderson [31] (STOC 2013), formalized in BCHKS [6], Table 1, as the row [RVW13]:  $\gamma$  arbitrary,  $a \geq 2$ ,  $\varepsilon^* = \gamma$ . The proof uses only linearity of  $C$  (closure under  $F$ -linear combinations): no polynomial structure, no multiplicative structure on  $L$ , no restriction on  $k$  or  $\delta$ . It applies to any linear code over any finite field.*

What is new in this paper is not the CA bound, but (a) its application above the Johnson radius to obtain an unconditional FRI soundness theorem in the open zone  $(\delta_J, 1-\rho)$ , where no such theorem was previously available; (b) the unconditional analytic equal-threshold CA upper bound  $\varepsilon_{\text{ca}}(C, \delta, \delta) \leq \binom{n}{w}/|F|$  at  $w = n - k - 1$  (Theorem 8, with the matching lower bound holding in the asymptotic-in- $|F|$ -at-fixed- $n$  regime; see Remark 12 and Remark 14); and (c) the tightness argument showing the 1:2 ratio is optimal at FRI scale (Remark 14). The structural observation enabling (a) is the chain  $\delta/2 < (1-\rho)/2 \leq \delta_J$  (Remark 1), so after one round the protocol enters BCIKS’s zero-loss regime where BCIKS [7] supplies the round-by-round gap.

The prior proximity-gap results of AHIV17 [32] ( $\delta < \delta_V/4$ ), RZ18 [33] ( $\delta < \delta_V/3$ ), and BKS18 [34] ( $\delta$  above unique decoding) each proved successively stronger gap statements. The BKS18 FRI soundness proof [34] (Theorem 16, the one used by BCIKS [7]) splits on  $\Delta(f_2, C)$  and bounds bad  $\alpha$ ’s in each case separately (Case 1/2), yielding a 1:3 ratio ( $\delta_{\text{hd}} = \delta/3$ ), tight for that framework because the case bookkeeping gives  $2\delta_{\text{hd}} + \delta_{\text{hd}} = 3\delta_{\text{hd}} \leq \delta$  only when  $\delta_{\text{hd}} \leq \delta/3$ . The direct proof above bypasses Case 1/2 entirely: it solves for both  $f_1$  and  $f_2$  simultaneously from two bad  $\gamma$ ’s, obtaining the joint distance bound  $\leq 2\delta_{\text{hd}}$  in one step. This is  $\leq \delta$  when  $\delta_{\text{hd}} \leq \delta/2$ , improving the ratio from 1:3 to 1:2.

We state  $nR/|F|$  (not  $R/|F|$ ) in Theorem 20 to maintain a unified presentation with STIR (Theorem 48), where the per-round count cannot be tightened below  $n$  in general.

## 4 Equal-Threshold CA: Exact Bound and Tightness of the 1:2 Ratio

This section establishes the exact form of the equal-threshold CA bound (Theorem 8) and shows that it is tight up to lower-order terms (Proposition 10). The conclusion (Remark 14) is that the  $\binom{n}{w}$  constant is exponential in the code length, vacuous at FRI scale; this is the formal statement behind the “OP1 obstruction” framing of §1.9.

### 4.1 Upper bound at equal threshold

**Theorem 8** (Equal-threshold CA upper bound). *Let  $C \subseteq F^L$  be any linear code with  $|L| = n$ , and let  $w$  be any integer with  $0 \leq w \leq n - 1$ . Set  $\delta := w/n$ . For any  $f_1, f_2 : L \rightarrow F$  satisfying*

$\Delta_{\text{joint}}((f_1, f_2), C \times C) > \delta$ :

$$|\{\gamma \in F : \Delta(f_1 + \gamma f_2, C) \leq \delta\}| \leq \binom{n}{w}.$$

Consequently  $\varepsilon_{\text{ca}}(C, \delta, \delta) \leq \binom{n}{w}/|F|$ . For  $C = \text{RS}[F, L, k]$  at the boundary  $w = n - k - 1$ , the bound is attained for  $|F| > \binom{n}{w}^2$  (Proposition 10).

*Proof.* Let  $\Gamma = \{\gamma \in F : \Delta(f_1 + \gamma f_2, C) \leq \delta\}$ . For each  $\gamma \in \Gamma$  there exists a witness codeword  $c_\gamma \in C$  with  $d(f_1 + \gamma f_2, c_\gamma) \leq w$ , so the agreement set has size  $\geq n - w$ . Choose any  $A_\gamma \subseteq L$  with  $|A_\gamma| = n - w$  on which  $f_1 + \gamma f_2 = c_\gamma$ , and define the assignment

$$\varphi : \Gamma \rightarrow \binom{L}{n-w}, \quad \varphi(\gamma) = A_\gamma.$$

*Claim:*  $\varphi$  is injective. Suppose  $\gamma \neq \gamma'$  in  $\Gamma$  with  $\varphi(\gamma) = \varphi(\gamma') = A$  ( $|A| = n - w \geq 1$ ). On  $A$ :

$$f_1 + \gamma f_2 = c_\gamma, \quad f_1 + \gamma' f_2 = c_{\gamma'}.$$

Subtracting and dividing by  $\gamma - \gamma' \neq 0$  (allowed since  $C$  is closed under  $F$ -linear combinations):

$$f_2|_A = \frac{c_\gamma - c_{\gamma'}}{\gamma - \gamma'}|_A, \quad g_2 := \frac{c_\gamma - c_{\gamma'}}{\gamma - \gamma'} \in C.$$

Likewise  $f_1|_A = c_\gamma|_A - \gamma g_2|_A$ , so setting  $g_1 := c_\gamma - \gamma g_2 \in C$  gives  $(f_1, f_2) = (g_1, g_2)$  on  $A$ . Hence the joint disagreement is at most  $|L \setminus A| = n - (n - w) = w$ , so

$$\Delta_{\text{joint}}((f_1, f_2), C \times C) \leq \frac{w}{n} = \delta,$$

contradicting the hypothesis  $\Delta_{\text{joint}} > \delta$ . Therefore  $\varphi$  is injective and  $|\Gamma| \leq \binom{n}{n-w} = \binom{n}{w}$ .  $\square$

**Remark 9** (Scope of Theorem 8: convention  $w = n - k - 1$ ). The bound holds (i) for all  $(f_1, f_2)$  with  $\Delta_{\text{joint}} > \delta$  — not just generic ones; (ii) over any field  $F$ , including small fields where  $|F| \leq \binom{n}{w}$  (in which case the bound  $\binom{n}{w}/|F| \geq 1$  is trivial); (iii) at the specific weight  $w = n - k - 1$  (the largest  $w$  with  $\delta n < n - k$ , i.e., one below the covering radius).

*Smaller decoding balls.* For a strictly smaller ball  $w' < w$  corresponding to a smaller threshold  $\delta' = w'/n$ : the  $(k+1)$ -subset injection of Theorem 8 still applies (the agreement set has size  $\geq n - w' \geq k + 1$ ), and yields the loose bound  $|\Gamma| \leq \binom{n}{k+1} = \binom{n}{w'}$ . This is not tight: the right quantity at threshold  $w'$  is  $\binom{n}{w'}$  (not  $\binom{n}{w}$ ), obtained by injecting  $\Gamma$  into  $\binom{L}{n-w'}$ , the family of  $(n-w')$ -subsets on which a witness codeword agrees. The tight statement is therefore  $|\Gamma| \leq \binom{n}{w'} = \binom{n}{n-w'}$  at threshold  $\delta' = w'/n$ , with the prize-mapping table and Remark 14 below using the convention  $w = \lfloor \delta n \rfloor$  for the asymptotic-in- $n$  statements (Stirling:  $\binom{n}{\lfloor \delta n \rfloor} \sim 2^{nH(\delta)}$ ).

*Convention.* Throughout, we fix  $w = n - k - 1$  for Theorem 8 (the equality-case statement) and  $w = \lfloor \delta n \rfloor$  for the asymptotic vacuity statements at general  $\delta < (n - k - 1)/n$  (Remark 14). Mixing the two conventions on the same statement is avoided.

## 4.2 Tightness

**Proposition 10** (Equal-threshold CA lower bound, large field, RS code). Let  $C = \text{RS}[F, L, k]$  with  $|L| = n$ ,  $w = n - k - 1$ ,  $\delta = w/n$ . For every field  $F$  with  $|F| > \binom{n}{w}^2$  there exist  $f_1, f_2 : L \rightarrow F$  with  $\Delta_{\text{joint}}((f_1, f_2), C \times C) > \delta$  and

$$|\{\gamma \in F : \Delta(f_1 + \gamma f_2, C) \leq \delta\}| = \binom{n}{w}.$$

*Proof.* For each  $(k+1)$ -subset  $A \subset L$ , encode  $A$ -consistency of  $f_1 + \gamma f_2$  as a single linear equation in  $\gamma$ . Fix  $A$  and let  $\pi_A : F^L \rightarrow F$  be the linear functional defined by:  $\pi_A(f) := (\text{value of the unique degree-}<k \text{ Lagrange interpolant of } f|_{A_0} \text{ at the remaining point of } A) - f \text{ at that point, where } A_0 \text{ is any fixed } k\text{-subset of } A \text{ (the choice of } A_0 \text{ does not affect the kernel). Then } f \text{ agrees with some codeword of } C \text{ on all of } A \text{ iff } \pi_A(f) = 0$ . Set

$$a_A := \pi_A(f_1) \in F, \quad b_A := \pi_A(f_2) \in F.$$

By linearity of  $\pi_A$ :  $\pi_A(f_1 + \gamma f_2) = a_A + \gamma \cdot b_A$ . Hence  $f_1 + \gamma f_2$  is  $A$ -consistent iff

$$a_A + \gamma \cdot b_A = 0. \tag{3}$$

For  $b_A \neq 0$ , equation (3) has the unique solution  $\gamma_A = -a_A/b_A \in F$ . As  $A$  ranges over the  $\binom{n}{k+1} = \binom{n}{w}$  choices, the values  $\{\gamma_A\}_A$  exhaust all candidate decodable  $\gamma$ 's by Theorem 8's injection (when each agreement set is exactly  $(k+1)$ -sized, the injection becomes a bijection onto its image).

*Construction.* Fix  $f_2 \in F^L$  such that  $b_A(f_2) \neq 0$  for every  $A \in \binom{L}{k+1}$ ; the bad set  $\{f : \pi_A(f) = 0\}$  is a single proper hyperplane in  $F^L$  for each  $A$ , contributing  $|F|^{n-1}/|F|^n = 1/|F|$  measure, so the union of  $\binom{n}{w}$  such hyperplanes leaves at least  $|F|^n(1 - \binom{n}{w}/|F|) > 0$  valid  $f_2$  when  $|F| > \binom{n}{w}$ . Pick  $f_1 \in F^L$  uniformly at random. For  $A \neq A'$ : the equation  $\gamma_A = \gamma_{A'}$  becomes

$$a_A b_{A'} - a_{A'} b_A = 0 \iff b_{A'} \cdot \pi_A(f_1) - b_A \cdot \pi_{A'}(f_1) = 0,$$

a single linear constraint on  $f_1$ . Since  $\pi_A$  and  $\pi_{A'}$  are non-proportional functionals (their kernels differ — they encode interpolation residuals on different  $(k+1)$ -subsets), the linear combination  $b_{A'}\pi_A - b_A\pi_{A'}$  is a nonzero functional on  $F^L$ , hence vanishes on a hyperplane of measure  $1/|F|$ . By union bound over  $\binom{\binom{n}{w}}{2} \leq \binom{n}{w}^2/2$  pairs:

$$\Pr_{f_1}[\text{some } \gamma_A = \gamma_{A'}] \leq \frac{\binom{n}{w}^2}{2|F|} < 1$$

provided  $|F| > \binom{n}{w}^2/2$ . The hypothesis  $|F| > \binom{n}{w}^2$  suffices, giving an  $f_1$  with all  $\binom{n}{w}$  values  $\gamma_A$  pairwise distinct in  $F$ . By construction each  $\gamma_A$  is decodable (witnessed by codeword agreement on the  $(k+1)$ -subset  $A$ ), so  $|\Gamma| \geq \binom{n}{w}$ . Combined with the upper bound (Theorem 8),  $|\Gamma| = \binom{n}{w}$ .

*Joint-distance check.* We must verify  $\Delta_{\text{joint}}((f_1, f_2), C \times C) > \delta$ . If  $\Delta_{\text{joint}} \leq \delta$ , then there exists a codeword pair  $(g_1, g_2) \in C \times C$  with  $|\{x : (f_1(x), f_2(x)) \neq (g_1(x), g_2(x))\}| \leq \delta n$ . For any  $\gamma \in F$ :  $f_1 + \gamma f_2$  agrees with  $g_1 + \gamma g_2 \in C$  on  $\geq (1 - \delta)n$  points, so every  $\gamma \in F$  would be decodable, giving  $|\Gamma| = |F|$ . But the construction gives  $|\Gamma| = \binom{n}{w} < |F|$  (since  $|F| > \binom{n}{w}^2 \geq \binom{n}{w}$ ). Therefore  $\Delta_{\text{joint}} > \delta$ .  $\square$

**Corollary 11** (Asymptotic equality at  $w = n - k - 1$ , large field at fixed  $n$ , RS code). *For  $C = \text{RS}[F, L, k]$  with  $w = n - k - 1$ ,  $\delta = w/n$ , and every field  $F$  with  $|F| > \binom{n}{w}^2$ :*

$$\varepsilon_{\text{ca}}(C, \delta, \delta) = \frac{\binom{n}{w}}{|F|}.$$

**Remark 12** (Scope of Corollary 11: asymptotic regime). *The matching lower bound (Proposition 10) requires  $|F| > \binom{n}{w}^2$ , so Corollary 11 is an asymptotic equality in  $|F|$  at fixed  $n$ , not an asymptotic equality in  $n$  at fixed  $|F|/n$ . At FRI scale ( $n \sim 2^{20}$ ,  $|F| \sim 2^{124}$ ),  $\binom{n}{w}^2 \gg |F|$ , so the lower-bound construction does not apply and Corollary 11 cannot be invoked at deployment parameters. What does apply at deployment is the upper bound (Theorem 8), which is unconditional in  $|F|$ , and which already gives  $\binom{n}{w}/|F| \gg 1$  at deployment scale (Remark 14); the upper bound being vacuous at deployment is a strictly stronger statement than “the equal-threshold ratio is intrinsically tight”, and does not require the matching-tightness construction.*

### 4.3 Computational verification

The exact bound is verified for small parameters where exhaustive enumeration is feasible.

**Remark 13** (Empirical confirmation). *For RS[6, 3] over  $\mathbb{F}_q$ ,  $q \in \{7, 13, 19, \dots, 163\}$ : the empirical maximum (over all tested  $(f_1, f_2)$ ) of the number of  $\delta$ -close  $\gamma$ 's equals  $\min(q, \binom{6}{2}) = \min(q, 15)$  exactly. For RS[8, 4] over  $q \in \{17, \dots, 281\}$ :  $\min(q, \binom{8}{3}) = \min(q, 56)$ , also exactly. See Note [25]; scripts `op1_scaling.py` and `op1_scaling_n8.py` in [39].*

### 4.4 Why the 1:2 ratio is optimal at FRI scale

**Remark 14** (Why the 1:2 ratio is tight). *Algebraic ceiling. The proof of Theorem 6 shows: two bad  $\gamma$ 's at conclusion threshold  $\delta_{\text{hd}}$  produce a joint pair  $(g_1, g_2) \in C \times C$  matching  $(f_1, f_2)$  on  $\geq (1 - 2\delta_{\text{hd}})n$  points, giving  $\Delta_{\text{joint}} \leq 2\delta_{\text{hd}}$ . For this to contradict  $\Delta_{\text{joint}} > \delta$ : need  $2\delta_{\text{hd}} \leq \delta$ , i.e.,  $\delta_{\text{hd}} \leq \delta/2$ . The bound  $\delta_{\text{hd}} = \delta/2$  is therefore optimal for any argument that extracts a joint codeword pair from two bad random-linear-combination values.*

*Why pushing to equal threshold is vacuous (within this CA route). The upper bound of Theorem 8 is  $\binom{n}{w}/|F|$  unconditionally in  $|F|$ . The constant  $\binom{n}{w}$  grows as  $2^{nH(\delta)}$  (Stirling). At FRI-deployment code lengths and fields — e.g.  $n = 2^{14}$  already gives  $\binom{n}{\lfloor \delta n \rfloor} > 2^{1000}$  at  $\delta \geq 0.1$ , vastly exceeding any extension field used in production (BabyBear<sub>4</sub> at  $|F| \approx 2^{124}$ , KoalaBear sextic at  $|F| \approx 2^{186}$ ) — the bound  $\binom{n}{w}/|F|$  is exponentially larger than 1, so this route is presently vacuous against any meaningful security target like  $2^{-100}$ . Within the equal-threshold CA route, removing the  $2\times$  query overhead would therefore require either a tighter equal-threshold upper bound (none in the literature) or a non-CA proximity argument; whether the algebraic ceiling above can be improved by tracking  $\geq 3$  bad  $\gamma$ 's remains open (§8).*

*Half-threshold is tight. At  $\delta_{\text{hd}} = \delta/2$ ,  $\max_{\gamma} \leq 1$  in all tested cases [39], matching the  $\leq 1$  bound of Theorem 6. The constant 2 in the half-threshold bound is independent of  $n$ , in sharp contrast to the  $\binom{n}{w}$  constant at equal threshold.*

*Comparison with the BKS18 1:3 ratio. The BKS18 [34] Case 1/2 approach (split on  $\Delta(f_2, C)$ , bound each case separately) achieves ratio 1:3 with zero proximity loss at  $\delta_{\text{hd}} < \delta/3$ . The RVW13 [31]  $2\times 2$ -solve argument of Theorem 6 trades zero loss for a better ratio (1:2 with loss  $\delta_{\text{hd}}$ ); the 1:2 ratio has been folklore in the proximity-gap community since 2013.*

## 5 FRI Coupling and Proximity Gap

**Lemma 15** (FRI Coupling, even  $k$ ). *For even  $k$  and any  $f : L \rightarrow F$ :  $\Delta(f, \text{RS}_k) \leq \Delta((f_{\text{even}}, f_{\text{odd}}), \text{RS}_{k/2}^2)$ .*

*Proof.* By Remark 5:  $g \mapsto (g_{\text{even}}, g_{\text{odd}})$  is an isomorphism  $\text{RS}_k \rightarrow \text{RS}_{k/2}^2$ . For any  $g \in \text{RS}_k$  and  $y \in L'$  let  $c_y := |\{x \in \{\sqrt{y}, -\sqrt{y}\} : f(x) \neq g(x)\}| \in \{0, 1, 2\}$  denote the number of base-level errors in the fibre over  $y$ . If  $f_{\text{even}}(y) = g_{\text{even}}(y)$  and  $f_{\text{odd}}(y) = g_{\text{odd}}(y)$ , then  $f(\pm\sqrt{y}) = g(\pm\sqrt{y})$ , so every joint-error position  $y$  contributes  $c_y \geq 1$  error on  $L$  and at most 2. Let  $E = |\{y \in L' : c_y \geq 1\}|$  denote the number of joint-error positions. Then  $\Delta(f, g) = \sum_y c_y/n \leq 2E/n = E/|L'| = \Delta_{\text{joint}}$ , where the cancellation uses  $n = 2|L'|$ . Taking  $\min_g$  (using the isomorphism, the min over  $\text{RS}_{k/2}^2$  equals the min over  $\text{RS}_k$ ):  $\Delta(f, \text{RS}_k) \leq \Delta_{\text{joint}}$ .  $\square$

**Theorem 16** (Half-Threshold Proximity Gap). *For even  $k$  and  $\Delta(f, \text{RS}_k) > \delta$ :*

$$|\{\alpha \in F : \Delta(f_{\text{even}} + \alpha f_{\text{odd}}, \text{RS}_{k/2}) \leq \delta/2\}| \leq 1.$$

(The Johnson/capacity window  $\delta_J < \delta < 1-\rho$  is only needed when this theorem is composed downstream into the full FRI soundness theorem; see Theorem 20.)

*Proof.* By Lemma 15:  $\Delta((f_{\text{even}}, f_{\text{odd}}), \text{RS}_{k/2}^2) \geq \Delta(f, \text{RS}_k) > \delta$ . The isomorphism  $\text{RS}_k \xrightarrow{\sim} \text{RS}_{k/2}^2$  (Remark 5) gives  $\text{RS}_{k/2}^2 = \text{RS}_{k/2} \times \text{RS}_{k/2}$ , so  $\Delta_{\text{joint}}((f_{\text{even}}, f_{\text{odd}}), \text{RS}_{k/2} \times \text{RS}_{k/2}) > \delta$ . Apply Theorem 6 to  $(f_1, f_2) = (f_{\text{even}}, f_{\text{odd}})$  on  $C = \text{RS}_{k/2}$  over  $L'$ , with  $\alpha$  drawn from the challenge field  $F$  (typically an extension of the field containing  $L'$ ; see Remark 2): at most 1 value of  $\alpha \in F$  has  $\Delta(f_{\text{even}} + \alpha f_{\text{odd}}, \text{RS}_{k/2}) \leq \delta/2$ .  $\square$

## 6 FRI Soundness

**Definition 17** (Standard FRI query model). *The verifier samples  $q$  base-level points  $z_1, \dots, z_q \in L$  i.i.d. uniform with replacement; for each  $z_j$  it queries the folding-tree path  $z_j \rightarrow \pi_1(z_j) \rightarrow \pi_2(z_j) \rightarrow \dots \rightarrow \pi_R(z_j)$ , where  $\pi_i : L \rightarrow L^{(i)}$  is the iterated squaring map, and at every round  $i = 1, \dots, R$  checks the consistency relation  $f^{(i)}(\pi_i(z_j)) = f_{\text{even}}^{(i-1)}(\pi_i(z_j)) + \alpha_i \cdot f_{\text{odd}}^{(i-1)}(\pi_i(z_j))$ . (Equivalent statements for additive folding (App. B) and circle folding (App. C) replace  $\pi_i$  with the corresponding round- $i$  projection.)*

**Remark 18** (Implementation models: with-replacement vs. without-replacement). *This paper's bounds are stated for the i.i.d. with-replacement model (Definition 17), which is standard in the FRI/STIR/WHIR analyses [20, 7, 10, 11]. Deployed implementations typically draw  $q$  distinct base-level indices via Fiat–Shamir hashing (without replacement, or with domain separation that makes collisions cryptographically negligible). The without-replacement miss probability  $\binom{|L| - |\pi_i^{-1}(D_i)|}{q} / \binom{|L|}{q}$  is bounded above by  $(1 - \tau)^q$  for the same  $\tau = |D_i|/|L^{(i)}|$  by negative association of without-replacement indicators (Joag-Dev–Proschan [30]; see also Dubhashi–Ranjan [29]), so the bounds in Theorems 20, 48 and Conjecture 54 apply unchanged to deployed protocols.*

**Lemma 19** (Catch probability under i.i.d. base sampling). *Under Definition 17, fix any round  $i$  and any deviation set  $D_i \subseteq L^{(i)}$  with  $|D_i|/|L^{(i)}| \geq \tau$ . Then*

$$\Pr_{z_1, \dots, z_q \stackrel{\text{i.i.d.}}{\sim} L} [\pi_i(z_j) \notin D_i \text{ for all } j = 1, \dots, q] \leq (1 - \tau)^q.$$

*Proof.* Let  $X_j := \mathbf{1}\{\pi_i(z_j) \in D_i\}$ . The map  $\pi_i : L \rightarrow L^{(i)}$  is  $2^i$ -to-1 (a composition of  $i$  squarings, each 2-to-1 on the relevant subgroup), so  $|\pi_i^{-1}(D_i)| = 2^i \cdot |D_i|$  and

$$\Pr_{z_j \sim L} [X_j = 1] = \frac{2^i \cdot |D_i|}{2^i \cdot |L^{(i)}|} = \frac{|D_i|}{|L^{(i)}|} \geq \tau.$$

The variables  $X_1, \dots, X_q$  are i.i.d. Bernoulli with parameter  $\geq \tau$  because  $z_1, \dots, z_q$  are i.i.d. and each  $X_j = h(z_j)$  for the deterministic function  $h(z) = \mathbf{1}\{\pi_i(z) \in D_i\}$ . What matters is independence of the *base-level* samples  $z_j$ , not of the *induced* samples  $\pi_i(z_j)$ : the latter can collide (and are guaranteed to do so by pigeonhole in late rounds, when  $|L^{(i)}| \leq q$ ), but the Bernoulli parameters  $\Pr[X_j = 1]$  are identical across  $j$ , and the joint distribution factors as a product of Bernoullis regardless of induced collisions. Therefore

$$\Pr[X_1 = \dots = X_q = 0] = \prod_{j=1}^q \Pr[X_j = 0] \leq (1 - \tau)^q. \quad \square$$



**Theorem 20** (FRI Soundness Above Johnson). *For  $\text{RS}[F, L, k]$  with  $\text{char}(F) \neq 2$ ,  $k = 2^m$ ,  $L$  closed under negation,  $\delta_J < \delta < 1 - \rho$ , and  $\Delta(f^{(0)}, \text{RS}_k) > \delta$ : FRI with  $R = m$  rounds and  $q$  base-level queries (each opening one full path) satisfies*

$$\Pr[\text{FRI accepts}] \leq \frac{nR}{|F|} + (1 - \delta/2)^q.$$

*Proof.* We partition all adaptive prover strategies into two cases according to whether the committed oracles match the honest fold. The prover sees  $\alpha_i$  before committing  $f^{(i)}$ ; the analysis below covers every such adaptive strategy.

**Strategy A: the prover commits the honest fold at every round.** In this case,  $f^{(i)}$  is the function  $f_{\text{even}}^{(i-1)} + \alpha_i \cdot f_{\text{odd}}^{(i-1)}$  (determined by  $f^{(0)}$  and  $\alpha_1, \dots, \alpha_i$ ). No adaptive freedom remains: the prover’s only choice was to fold honestly, and the resulting  $f^{(R)}$  is a *fixed* function of  $f^{(0)}$  and  $(\alpha_1, \dots, \alpha_R)$ .

For any fixed evaluation point  $y \in L^{(R)}$ ,  $f^{(R)}(y)$  is *multilinear* in  $(\alpha_1, \dots, \alpha_R)$  — individual degree  $\leq 1$  in each  $\alpha_i$  (one round introduces each variable, each via the linear even/odd combination  $f_{\text{even}}^{(i-1)} + \alpha_i \cdot f_{\text{odd}}^{(i-1)}$ ). Consequently, every syndrome  $s_j(f^{(R)}) = \sum_y c_y \cdot f^{(R)}(y)$  is multilinear in  $(\alpha_1, \dots, \alpha_R)$ .

*Some syndrome is nonzero as a polynomial.* Suppose all syndromes vanished identically (as polynomials in  $\alpha_1, \dots, \alpha_R$ ). Then  $f^{(R)} \in \text{RS}_{k/2^R}$  for *every* choice of  $(\alpha_1, \dots, \alpha_R) \in F^R$ . Equivalently — since  $\text{RS}_{k/2^R} \subset F^{L^{(R)}}$  is an  $F$ -linear subspace and the assignment  $(\alpha_1, \dots, \alpha_R) \mapsto f_{\alpha_1, \dots, \alpha_R}^{(R)}$  is multilinear — every multilinear coefficient of  $f^{(R)}$  in  $(\alpha_1, \dots, \alpha_R)$  also lies in  $\text{RS}_{k/2^R}$ . We extract  $f^{(0)} \in \text{RS}_k$  from this universal hypothesis as a polynomial-identity argument (*not* a protocol-level reconstruction; FRI does not expose intermediate  $f_{\text{even}}^{(i-1)}, f_{\text{odd}}^{(i-1)}$ , and an attacker cannot reconstruct  $f^{(0)}$  from  $f^{(R)}$  at runtime). For each round  $i = R, R-1, \dots, 1$ : by the multilinearity in  $\alpha_i$ , evaluating at  $\alpha_i = 0$  yields  $f_{\text{even}}^{(i-1)} \in \text{RS}_{k/2^i}$ , and the difference at  $\alpha_i = 1$  minus  $\alpha_i = 0$  yields  $f_{\text{odd}}^{(i-1)} \in \text{RS}_{k/2^i}$  (each as a polynomial-valued function of  $\alpha_1, \dots, \alpha_{i-1}$ , with  $|F| \geq 2$  ensuring two distinct evaluation points; the same multilinear-coefficient argument runs coefficient-wise in the remaining variables, so each multilinear coefficient of  $f_{\text{even}}^{(i-1)}, f_{\text{odd}}^{(i-1)}$  in  $\alpha_1, \dots, \alpha_{i-1}$  lies in  $\text{RS}_{k/2^i}$ ). The isomorphism of Remark 5 — which requires  $k/2^{i-1}$  even, guaranteed by  $k = 2^m$  — then gives  $f^{(i-1)} \in \text{RS}_{k/2^{i-1}}$ . Iterating to  $i = 1$  yields  $f^{(0)} \in \text{RS}_k$ , contradicting the hypothesis. Hence there exists a syndrome  $s_j$  that is a nonzero multilinear polynomial.

Since  $s_j$  is multilinear in  $\alpha_1, \dots, \alpha_R$  (individual degree  $\leq 1$  in each variable), the DeMillo–Lipton–Schwartz–Zippel lemma in the individual-degree form (a nonzero polynomial with individual degree  $\leq d_i$  in variable  $i$  vanishes on a uniformly random point with probability  $\leq \sum_i d_i/|F|$ ) gives

$$\Pr_{\alpha_1, \dots, \alpha_R \stackrel{\text{i.i.d.}}{\sim} F} [s_j(f^{(R)}) = 0] \leq R/|F|.$$

A fortiori,  $\Pr[f^{(R)} \in \text{RS}_{k/2^R}] \leq R/|F|$ .

**Strategy B: any prover strategy with at least one deviation.** We give a single round-by-round argument that applies uniformly to every adaptive prover, regardless of how many rounds it deviates at and whether its commits are codewords or not. This replaces an earlier exposition that case-split on single-vs-multi deviation and on a “lucky/unlucky” challenge dichotomy; the case split conflated the prover’s commit chain with the honest-fold chain. The argument below tracks the honest-fold chain explicitly, which is the standard round-by-round (RBR) FRI soundness framework [20, 7, 37] and is also the form taken by Theorem 75 in §D.1.

*Honest-fold reference sequence.* Define

$$\hat{f}^{(0)} := f^{(0)}, \quad \hat{f}^{(i)}(y) := \hat{f}_{\text{even}}^{(i-1)}(y) + \alpha_i \cdot \hat{f}_{\text{odd}}^{(i-1)}(y) \quad (y \in L^{(i)}).$$

This is a deterministic function of  $f^{(0)}$  and the verifier's challenges  $(\alpha_1, \dots, \alpha_R)$ ; in particular  $\hat{f}^{(i)}$  does *not* depend on the prover's subsequent commits  $f^{(1)}, \dots, f^{(R)}$ .

*Step 1 — Bad-challenge bound:*  $\leq nR/|F|$ . By induction on  $i$ , we bound the probability that some  $\hat{f}^{(i)}$  becomes  $(\delta/2)$ -close to  $\text{RS}_{k/2^i}$ :

- *Round 1.*  $\Delta(\hat{f}^{(0)}, \text{RS}_k) = \Delta(f^{(0)}, \text{RS}_k) > \delta$ . By Theorem 16, at most 1 value of  $\alpha_1$  makes  $\Delta(\hat{f}^{(1)}, \text{RS}_{k/2}) \leq \delta/2$ . Probability  $\leq 1/|F|$ .
- *Rounds  $i = 2, \dots, R$ .* Inductively assume  $\Delta(\hat{f}^{(i-1)}, \text{RS}_{k/2^{i-1}}) > \delta/2$ . By Remark 1,  $\delta/2 < (1-\rho)/2 \leq \delta_J$ , so the threshold  $\gamma = \delta/2$  lies inside BCIKS's zero-loss regime. By Lemma 15,  $\Delta_{\text{joint}}((\hat{f}_{\text{even}}^{(i-1)}, \hat{f}_{\text{odd}}^{(i-1)}), \text{RS}_{k/2^i}^2) > \delta/2$ , and the BCIKS proximity gap [7] (Theorem 1.2, Eq. (1.1)) at threshold  $\gamma = \delta/2$  gives: at most  $n$  values of  $\alpha_i$  make  $\Delta(\hat{f}^{(i)}, \text{RS}_{k/2^i}) \leq \delta/2$ . Probability  $\leq n/|F|$ .

Union over  $R$  rounds:  $\Pr[\exists i : \Delta(\hat{f}^{(i)}, \text{RS}_{k/2^i}) \leq \delta/2] \leq 1/|F| + (R-1)n/|F| \leq nR/|F|$ .

*Step 2 — Catch-probability bound:*  $\leq (1 - \delta/2)^q$ . Condition on the good event from Step 1:  $\Delta(\hat{f}^{(i)}, \text{RS}_{k/2^i}) > \delta/2$  for every  $i$ .

If the prover commits  $f^{(i)} = \text{fold}(f^{(i-1)}, \alpha_i)$  at every round, then  $f^{(i)} = \hat{f}^{(i)}$  at every round (induction on  $i$  from the common base  $f^{(0)} = \hat{f}^{(0)}$ ), so  $\Delta(f^{(R)}, \text{RS}_{k/2^R}) > \delta/2 > 0$ ,  $f^{(R)} \notin \text{RS}_{k/2^R}$ , and the final check rejects. The verifier therefore accepts only if the prover commits at least one  $f^{(i)}$  that disagrees with the fold of its own previous commit (the verifier-checkable notion of deviation), in such a way that the resulting chain reaches  $f^{(R)} \in \text{RS}_{k/2^R}$ .

Define the round- $i$  *deviation set* (relative to the prover's own previous commit  $f^{(i-1)}$ , since this is what the verifier's consistency check examines):

$$B_i := \{y \in L^{(i)} : f^{(i)}(y) \neq f_{\text{even}}^{(i-1)}(y) + \alpha_i \cdot f_{\text{odd}}^{(i-1)}(y)\}, \quad \beta_i := |B_i|/|L^{(i)}|.$$

A query  $z_j \in L$  “catches” the prover at round  $i$  iff  $\pi_i(z_j) \in B_i$ , so the verifier rejects  $z_j$  iff  $z_j \in T := \bigcup_{i=1}^R \pi_i^{-1}(B_i) \subseteq L$ .

*Path-consistency lemma.* We invoke the standard FRI catch lemma in the form recorded by [37, Lemma 5.5 / honest-fold reference variant], which reformulates the original BBHR/BCIKS catch lemma in exactly the honest-fold-reference framework used here: under the conditions  $f^{(R)} \in \text{RS}_{k/2^R}$  and  $\Delta(\hat{f}^{(R)}, \text{RS}_{k/2^R}) > \delta/2$ , the merged catching set  $T = \bigcup_{i=1}^R \pi_i^{-1}(B_i) \subseteq L$  has relative measure at the base level satisfying  $|T|/|L| \geq \delta/2$ . The original statements [20, Lemma 4.1 / §6], [7, Lemma 8.2] bound the catching set relative to the prover's commit chain folded against *itself*; the conversion to the honest-fold reference  $\hat{f}^{(j)}$  used here is the standard induction-on-rounds argument carried out in [37] and is what we cite. The underlying path-coverage step is: by induction on  $j$ ,  $B_j$  contains every  $y \in L^{(j)}$  where  $f^{(j)}(y) \neq \hat{f}^{(j)}(y)$  except possibly via ancestor-collisions in  $B_1, \dots, B_{j-1}$ ; hence  $T = \bigcup_j \pi_j^{-1}(B_j)$  at the base level covers every disagreement position between  $f^{(R)}$  and  $\hat{f}^{(R)}$ , and counting paths through level- $R$  disagreements gives the  $\delta/2$  lower bound on  $|T|/|L|$ . We invoke the catch lemma as a black box; see the cited references for the detailed coverage argument.

By Lemma 19 applied to  $T$  with  $\tau = \delta/2$ : the probability that all  $q$  i.i.d. base-level queries miss  $T$  is at most  $(1 - \delta/2)^q$ .

*Combining Steps 1 and 2.*

$$\Pr[\text{Strategy B succeeds}] \leq \frac{nR}{|F|} + (1 - \delta/2)^q.$$

This bound holds simultaneously for any number of deviation rounds (single or multi) and any commit pattern (codeword or non-codeword), since the honest-fold reference and the merged catching set are defined uniformly across all prover strategies.

**Combining Strategies A and B.** Every adaptive prover either folds honestly at every round (Strategy A) or deviates at some round (Strategy B):

$$\Pr[\text{accept}] \leq \max\left(\frac{R}{|F|}, \frac{nR}{|F|} + (1-\delta/2)^q\right) \leq \frac{nR}{|F|} + (1-\delta/2)^q. \quad \square$$

**Remark 21** (Why  $\delta/2$  does not compound across rounds). *A common concern: the reduced threshold  $\delta/2$  could degrade further at each round (to  $\delta/4$ ,  $\delta/8$ , ...), eventually making the consistency check trivial. This does not happen. After round 1 reduces the effective distance from  $\delta$  to  $\delta/2$ , this threshold sits strictly below the Johnson radius  $\delta_J = 1 - \sqrt{\rho}$  via the chain  $\delta/2 < (1-\rho)/2 \leq \delta_J$  (Remark 1; note that the Johnson radius  $\delta_J$  exceeds the unique-decoding radius  $(1-\rho)/2$  for every  $\rho < 1$ , so  $\delta/2$  is in fact inside the strictly larger BCIKS zero-loss regime). The BCIKS proximity gap [7] (Theorem 1.2, Eq. (1.1)) at threshold  $\gamma = \delta/2 < \delta_J$  bounds the number of bad challenges to  $\leq n$  per round. The distance remains  $\delta/2$  (not  $\delta/4$ ) because we invoke the equal-threshold proximity gap rather than re-applying the two-threshold Theorem 6 iteratively. In short: Theorem 6 crosses into BCIKS's zero-loss regime at round 1; the BCIKS proximity gap [7] then locks the distance at  $\delta/2$  for all subsequent rounds.*

## 7 List-Size Landscape and OP2

The results in this section are independent of the FRI soundness theorem (Theorem 20), which uses only the half-threshold CA (Theorem 6) and classical list decoding. They address Open Problem 2 (§8): can one show  $M = 0$  on all RS-compatible flats? We give the first  $p$ -dependent moment bound on the list size (Theorem 25 for  $c = 1$ ; Corollary 30 for general  $c$ ), prove the  $M = 0$  conjecture is false at deployment code lengths (Proposition 35), and reduce the remaining gap to a worst-case-from-average problem.

**Remark 22** ( $p$ -independent list-size bounds). *Classical bounds give  $M \leq n^2/(t(t-1))$  for  $k = 2$  (Fisher double-counting; see Sudan [15], Guruswami–Sudan [16] for general list decoding) and  $M \leq \sum_{i=0}^{k-2c} \binom{n}{i}$  for general  $k$  (Frankl–Wilson [14] linear algebra method), where  $c = n-k-w$  is the codimension excess. Both are independent of  $p$  and cannot capture the  $p^{-\alpha}$  decay observed computationally [39]: for RS[16, 8] at  $c = 1$ ,  $M_{\max}(p) \approx 4200 \cdot p^{-3/4}$  across 19 primes from  $p = 17$  to 131.*

### 7.1 Syndrome hyperplane structure and second moment

Classical  $p$ -independent bounds (Remark 22) cannot capture the  $p^{-\alpha}$  decay observed computationally. The following results explain this decay via the hyperplane geometry of syndrome space and provide the first  $p$ -dependent moment bounds on the list size. The results hold for *all*  $k$ .

Fix RS $[n, k]$  over  $\mathbb{F}_p$  with  $p > n$ , evaluation domain  $L = \{\alpha_1, \dots, \alpha_n\}$ , and codimension excess  $c = 1$  (error weight  $w = n - k - 1$ , syndrome dimension  $D = n - k$ ). For each  $w$ -subset  $E \subset [n]$ , the syndrome equations  $\sum_{i \in E} e_i \alpha_i^j = s_j$  ( $j = k, \dots, n-1$ ) form a  $D \times w$  Vandermonde system. With  $D = w + 1$ , the compatibility condition is a single hyperplane  $H_E = \{s \in \mathbb{F}_p^D : \langle n_E, s \rangle = 0\}$ , and  $M(s) = |\{E \in \binom{[n]}{w} : s \in H_E\}|$ .

**Lemma 23** (Error-locator normal). *The normal vector of  $H_E$  is the coefficient vector of the error-locator polynomial:*

$$n_E = (\Lambda_{E,0}, \dots, \Lambda_{E,D-1}), \quad \Lambda_E(x) = \prod_{e \in E} (x - \alpha_e).$$

*Proof.* Since  $\alpha_e$  is a root of  $\Lambda_E$ :  $\sum_{\ell=0}^w \Lambda_{E,\ell} \alpha_e^{k+\ell} = \alpha_e^k \Lambda_E(\alpha_e) = 0$  for each  $e \in E$ . So the vector  $(\Lambda_{E,0}, \dots, \Lambda_{E,w})$  annihilates every column of the  $D \times w$  Vandermonde submatrix  $V_E$ . Since  $V_E$  has rank  $w$  (distinct evaluation points,  $p > n$ ), its left kernel is 1-dimensional and spanned by  $n_E$ .  $\square$

**Lemma 24** (Pairwise independence). *For distinct  $w$ -subsets  $E_1 \neq E_2$ , the normals  $n_{E_1}$  and  $n_{E_2}$  are linearly independent over  $\mathbb{F}_p$ .*

*Proof.* Both  $\Lambda_{E_1}$  and  $\Lambda_{E_2}$  are monic of degree  $w$ . Monic polynomials of the same degree are proportional only if equal, requiring identical root sets:  $E_1 = E_2$ .  $\square$

**Theorem 25** (Exact second moment of the list size). *For RS $[n, k]$  over  $\mathbb{F}_p$  with  $p > n$  and  $c = 1$ , the list size  $M$  over uniform  $s \in \mathbb{F}_p^D$  satisfies:*

$$\mathbb{E}[M] = \frac{\binom{n}{w}}{p}, \quad \text{Var}[M] = \mathbb{E}[M] \left(1 - \frac{1}{p}\right).$$

*Proof.* Write  $M = \sum_E \mathbf{1}_{s \in H_E}$ . Each indicator has  $\Pr[s \in H_E] = 1/p$ . For  $E_1 \neq E_2$ : by Lemma 24,  $H_{E_1} \cap H_{E_2}$  has codimension 2 in  $\mathbb{F}_p^D$ , so  $\Pr[s \in H_{E_1} \cap H_{E_2}] = 1/p^2 = \Pr[s \in H_{E_1}] \cdot \Pr[s \in H_{E_2}]$ . The indicators are pairwise uncorrelated, giving  $\text{Var}[M] = \binom{n}{w} \cdot \frac{1}{p} (1 - \frac{1}{p})$ .  $\square$

**Remark 26** (Implications for the list-size landscape).

(i) *Near-Poisson concentration.* For large  $p$ :  $\text{Var}[M] \approx \mathbb{E}[M]$ , matching the Poisson dispersion relation. This justifies the heuristic that  $M$  concentrates around  $\binom{n}{w}/p$  and is typically 0 once  $p > \binom{n}{w}$ . For RS[16, 8] at  $c = 1$ :  $\binom{16}{7} = 11440$ , so  $\mathbb{E}[M] < 1$  for  $p \geq 11441$ . The power-law decay of  $M_{\max}$  (see Remark 22) is substantially sharper than the first-moment prediction, suggesting algebraic structure beyond pairwise independence governs the worst case.

(ii) *The  $c = 1$  threshold and  $c \geq 2$  generalization.* At  $c = 1$  (near capacity,  $w = n - k - 1$ ) with rate  $\rho = 1/2$ :  $\binom{n}{w} = \binom{n}{n/2-1} \sim 2^n / \sqrt{n}$ , so  $\mathbb{E}[M] < 1$  requires  $p > 2^n / \sqrt{n}$  — far beyond any practical field once  $n$  exceeds a few dozen. At the FRI-relevant distance  $\delta > \delta_J$  with codimension  $c = n/2 - \lfloor \delta n \rfloor \gg 1$  (for  $\rho = 1/2$ , just above Johnson gives  $c \approx 0.2n$ ): Theorem 28 and Corollary 30 establish  $\mathbb{E}[M] = \binom{n}{w}/p^c$  with pairwise independence for all pairs with  $|E_1 \cap E_2| \leq w - c$ , giving  $\text{Var}[M] \approx \mathbb{E}[M]$  (Poisson dispersion). At current code lengths,  $\mathbb{E}[M] = \binom{n}{w}/p^c \leq 2^{-5.3n} \approx 0$ .

(iii) *Average vs. worst case.* By Markov:  $\Pr_s[M \geq 1] \leq \binom{n}{w}/p$ . However, §8, item 2, asks for  $M_{\max} = 0$  over all syndromes (or all RS-compatible flats). The second-moment bound  $\text{Var}[M] \leq \mathbb{E}[M]$  does not improve on Markov for this worst-case question. Bridging the average-to-worst-case gap requires exploiting the algebraic structure of the error-locator hyperplane arrangement beyond pairwise independence — for example, showing that the  $\binom{n}{w}$  hyperplanes, despite their number, avoid all RS-compatible flats when  $p$  is sufficiently large relative to  $n$ .

## 7.2 Generalization to codimension excess $c \geq 2$

Theorem 25 treats codimension excess  $c = 1$ , where each  $w$ -subset  $E$  defines a single hyperplane  $H_E$  in syndrome space. For general  $c \geq 2$  (error weight  $w = n - k - c$ , syndrome dimension  $D = n - k$ ), the compatibility condition is  $s \in W_E := \text{Im}(V_E)$ , the column span of the  $D \times w$

Vandermonde submatrix  $V_E$ . Since  $V_E$  has rank  $w$  (MDS,  $p > n$ ),  $W_E$  is a  $w$ -dimensional subspace of  $\mathbb{F}_p^D$ , equivalently a codimension- $c$  subspace. Its orthogonal complement  $W_E^\perp$  has dimension  $c$  and is the left null space of  $V_E$ . The following results extend the error-locator normal structure and pairwise independence to this regime. The case  $c \geq 2$  is the *FRI-relevant* regime: at rate  $\rho = 1/2$  and distance  $\delta \approx 0.3$ , the codimension excess is  $c = n/2 - \lfloor 0.3n \rfloor \approx 0.2n$ .

**Lemma 27** (Generalized normal space). *For codimension excess  $c \geq 1$ , the left null space of the  $D \times w$  Vandermonde submatrix  $V_E$  has dimension  $c$  and is spanned by the coefficient vectors of*

$$\Lambda_E(x) \cdot x^r, \quad r = 0, 1, \dots, c-1,$$

where  $\Lambda_E(x) = \prod_{e \in E} (x - \alpha_e)$  is the error-locator polynomial.

*Proof.* Since  $\alpha_e$  is a root of  $\Lambda_E$ , the polynomial  $\Lambda_E(x) \cdot x^r$  vanishes at every  $\alpha_e$  for  $e \in E$ . Its coefficient vector (as a polynomial of degree  $w + r < D = w + c$ , embedded in  $\mathbb{F}_p^D$ ) therefore annihilates every column of  $V_E$ . The  $c$  polynomials  $\{\Lambda_E \cdot x^r\}_{r=0}^{c-1}$  have degrees  $w, w+1, \dots, w+c-1$ ; since  $\Lambda_E$  is monic of degree  $w$ , these have leading terms  $x^w, x^{w+1}, \dots, x^{w+c-1}$ , so they are linearly independent. The left null space of  $V_E$  has dimension  $D - w = c$  (distinct evaluation points,  $p > n$ ), and  $c$  linearly independent vectors span it.  $\square$

**Theorem 28** (Pairwise independence at  $c \geq 2$ ). *For distinct  $w$ -subsets  $E_1 \neq E_2$  with  $j = |E_1 \cap E_2| \leq w - c$ : the  $2c$  normal vectors from Lemma 27 (i.e., the coefficient vectors of  $\{\Lambda_{E_1} x^r, \Lambda_{E_2} x^r : 0 \leq r < c\}$ ) are linearly independent over  $\mathbb{F}_p$ . Consequently:*

$$W_{E_1} \cap W_{E_2} \text{ has codimension } 2c, \quad \Pr_s[s \in W_{E_1} \cap W_{E_2}] = \frac{1}{p^{2c}} = \Pr_s[s \in W_{E_1}] \cdot \Pr_s[s \in W_{E_2}].$$

The membership indicators  $\mathbf{1}_{s \in W_{E_1}}$  and  $\mathbf{1}_{s \in W_{E_2}}$  are pairwise uncorrelated for uniform  $s$ .

*Proof.* Suppose  $\sum_{r=0}^{c-1} (a_r \Lambda_{E_1}(x) x^r + b_r \Lambda_{E_2}(x) x^r) = 0$  as polynomials of degree  $< D = w + c$ . Set  $P(x) = \sum a_r x^r$  and  $Q(x) = \sum b_r x^r$  (both of degree  $< c$ ). Then  $\Lambda_{E_1} P + \Lambda_{E_2} Q = 0$ .

Let  $G = \gcd(\Lambda_{E_1}, \Lambda_{E_2}) = \Lambda_{E_1 \cap E_2}$  of degree  $j$ , and write  $\Lambda_{E_i} = G \cdot A_i$  where  $\gcd(A_1, A_2) = 1$  and  $\deg A_i = w - j$ . Dividing by  $G$ :  $A_1 P = -A_2 Q$ . Since  $\gcd(A_1, A_2) = 1$ , we need  $A_1 \mid Q$ . But  $\deg A_1 = w - j \geq w - (w - c) = c > \deg Q$ , so  $Q = 0$ , hence  $P = 0$ .

The  $2c$  normals are linearly independent, so  $W_{E_1}^\perp + W_{E_2}^\perp$  has dimension  $2c$ . Since  $\dim(W_{E_1}^\perp) = \dim(W_{E_2}^\perp) = c$ , the intersection  $W_{E_1} \cap W_{E_2}$  has codimension  $2c$ , giving exact independence.  $\square$

**Remark 29** (Sharpness of the threshold  $j \leq w - c$ ). *The threshold is sharp. For  $j > w - c$ , set  $\ell = j - (w - c) \geq 1$ . In the notation of the proof above,  $\deg A_i = w - j = c - \ell$ . For any polynomial  $R$  of degree  $< \ell$ , setting  $Q = A_1 R$  (degree  $< c$ ) and  $P = -A_2 R$  (degree  $< c$ ) gives  $\Lambda_{E_1} P + \Lambda_{E_2} Q = G(-A_1 A_2 R + A_2 A_1 R) = 0$ : a nontrivial relation among the  $2c$  normals. The space of such relations has dimension  $\ell = j - (w - c)$ , so the rank drops from  $2c$  to  $2c - \ell$ .*

*In particular, at  $j = w - c + 1$  ( $\ell = 1$ ): rank =  $2c - 1$ , verified computationally for all tested pairs [39].*

**Corollary 30** (Exact moments of the list size at  $c \geq 2$ ). *For  $\text{RS}[n, k]$  over  $\mathbb{F}_p$  with  $p > n$ , codimension excess  $c \geq 1$ , and  $w = n - k - c$ , define  $M = |\{E \in \binom{[n]}{w} : s \in W_E\}|$  for uniform  $s \in \mathbb{F}_p^D$ . Then:*

$$\mathbb{E}[M] = \frac{\binom{n}{w}}{p^c}.$$

For the variance, decomposing by intersection size  $j = |E_1 \cap E_2|$ :

$$\text{Var}[M] = \underbrace{\sum_{j=0}^{w-c} \binom{n}{w} \binom{w}{j} \binom{n-w}{w-j} \cdot 0}_{\text{pairwise independent: } j \leq w-c} + \underbrace{\sum_{j=w-c+1}^w \binom{n}{w} \binom{w}{j} \binom{n-w}{w-j} \left( \frac{1}{p^{2c-\ell(j)}} - \frac{1}{p^{2c}} \right)}_{\text{correlated: } j > w-c}, \quad (4)$$

where  $\ell(j) = j - (w - c) \in \{1, \dots, c\}$  is the rank deficiency from Remark 29.

*Proof.* Write  $M = \sum_E \mathbf{1}_{s \in W_E}$ . Each indicator has  $\Pr[s \in W_E] = 1/p^c$  since  $W_E$  has codimension  $c$ . For  $\mathbb{E}[M]$ : linearity of expectation gives  $\binom{n}{w}/p^c$ .

For the second moment:  $\mathbb{E}[M^2] = \sum_{E_1, E_2} \Pr[s \in W_{E_1} \cap W_{E_2}]$ . By Theorem 28, pairs with  $j \leq w - c$  contribute  $1/p^{2c}$  each — exactly the product of marginals, so their contribution to  $\text{Var}[M]$  is zero. Pairs with  $j > w - c$  have  $\ell(j) = j - (w - c)$  dependent normal vectors (Remark 29), so  $W_{E_1} \cap W_{E_2}$  has codimension  $2c - \ell(j)$ , giving the stated excess probability  $1/p^{2c-\ell(j)} - 1/p^{2c}$ .  $\square$

**Corollary 31** (Expected true list size). *Under the hypotheses of Corollary 30, the expected true list size (counting only supports whose Vandermonde error values are all nonzero) satisfies*

$$\mathbb{E}[M_{\text{true}}] = \frac{\binom{n}{w} (p-1)^w}{p^{w+c}} = \frac{\binom{n}{w}}{p^c} \left(1 - \frac{1}{p}\right)^w.$$

*Proof.* Write  $M_{\text{true}} = \sum_E \mathbf{1}_{s \in W_E, \text{ all } v_i \neq 0}$ . For fixed  $E$ , the map  $v \mapsto V_E v$  is a bijection  $\mathbb{F}_p^w \rightarrow W_E$  (the  $D \times w$  Vandermonde matrix  $V_E$  has rank  $w$  since  $p > n$ , so it has a  $w$ -dimensional image). Under uniform  $s \in \mathbb{F}_p^D$ :  $\Pr[s \in W_E] = p^w/p^D = 1/p^c$ . Conditional on  $s \in W_E$ , the unique preimage  $v$  with  $V_E v = s$  is uniform in  $\mathbb{F}_p^w$ , so  $\Pr[\text{all } v_i \neq 0 \mid s \in W_E] = ((p-1)/p)^w$ . By linearity:

$$\mathbb{E}[M_{\text{true}}] = \binom{n}{w} \cdot \frac{1}{p^c} \cdot \left(\frac{p-1}{p}\right)^w = \frac{\binom{n}{w} (p-1)^w}{p^{w+c}}. \quad \square$$

**Remark 32** (Poisson dispersion at general  $c$ ). *The variance decomposition (4) has a clean structure:*  
(i) *Diagonal dominance.* The  $j = w$  term (same subset,  $\ell = c$ ) contributes  $\binom{n}{w} (1/p^c - 1/p^{2c}) \approx \mathbb{E}[M]$  to  $\text{Var}[M]$ . All other correlated terms ( $w - c < j < w$ ) involve  $\binom{w}{j} \binom{n-w}{w-j}$  which is polynomially bounded in  $n$ , while the excess probability  $1/p^{2c-\ell} \leq 1/p^{c+1}$  decays exponentially in  $c$ . For  $c \geq 2$  and  $p \gg n$ :  $\text{Var}[M] \approx \mathbb{E}[M]$ , matching the Poisson dispersion relation exactly as in the  $c = 1$  case.  
(ii) *Average vs. worst case.* The Poisson dispersion means that for random syndromes,  $M$  concentrates sharply around  $\binom{n}{w}/p^c$ . At deployment ( $c \approx 0.2n$ ,  $p = 2^{31}$ ):  $\mathbb{E}[M] \leq 2^{-5.3n} \approx 0$ , so  $M = 0$  for overwhelmingly most syndromes. The worst case is governed by evaluation syndromes and the incidence bound; see Remark 33 for the precise structure.

**Remark 33** (Syndrome-compatible supports vs. true list size). *The error-locator normal conditions (the key equations of Berlekamp) count supports  $E$  where  $\langle \Lambda_E \cdot x^r, s \rangle = 0$  for  $0 \leq r < c$ , but do not verify that the error values (the Vandermonde solution  $V_E^{-1} s$ ) are all non-zero. Define:*

$$M_{\text{compat}}(s) = |\{E \in \binom{L}{w} : \langle \Lambda_E \cdot x^r, s \rangle = 0 \forall r < c\}|,$$

$$M_{\text{true}}(s) = |\{E \in \binom{L}{w} : V_E^{-1} s_{0:w} \text{ has all entries } \neq 0 \text{ and extension conditions hold}\}|.$$

The list-decoding list size is  $M_{\text{true}}$ ; the count  $M_{\text{compat}} \geq M_{\text{true}}$  overcounts by including supports where some error values vanish (actual error weight  $< w$ ).

*Evaluation syndromes:*  $M_{\text{compat}} = \binom{n-1}{w-1}$  but  $M_{\text{true}} = 0$ . Fix  $\alpha \in L$  and set  $s_\alpha := (1, \alpha, \alpha^2, \dots, \alpha^{D-1})$ . For every  $w$ -subset  $E$  and  $0 \leq r < c$ :  $\langle \Lambda_E \cdot x^r, s_\alpha \rangle = \alpha^r \Lambda_E(\alpha)$ , which

vanishes iff  $\alpha \in E$ . Hence  $M_{\text{compat}}(s_\alpha) = \binom{n-1}{w-1}$ . However, for any  $E \ni \alpha$ : the Vandermonde system  $V_E \cdot v = (1, \alpha, \dots, \alpha^{w-1})^T$  has the unique solution  $v_\alpha = 1$ ,  $v_\beta = 0$  for  $\beta \in E \setminus \{\alpha\}$  — a weight-1 error, not weight- $w$ . Therefore  $M_{\text{true}}(s_\alpha) = 0$  for  $w \geq 2$ . The  $\binom{n-1}{w-1}$  “compatible” supports are all false positives: the error-locator conditions are satisfied, but the actual error has weight 1.

True list size: bounded by incidence,  $p$ -dependent for large  $n$ . Exhaustive computation over all  $s \neq 0$  yields  $\max_{s \neq 0} M_{\text{true}}(s)$ :

| $n$ | $w$ | $c$ | $d$ | $\max M_{\text{compat}}$ | $\max M_{\text{true}}$         | $p$ tested               |
|-----|-----|-----|-----|--------------------------|--------------------------------|--------------------------|
| 5   | 2   | 1   | 1   | 4                        | 2                              | 5, 7, 11, 13, 31         |
| 6   | 3   | 1   | 2   | 10                       | 4                              | 7, 11, 13                |
| 7   | 3   | 1   | 2   | 15                       | 5                              | 7, 11, 13                |
| 8   | 3   | 1   | 2   | 21                       | 6                              | 11, 13                   |
| 6   | 3   | 2   | 1   | 10                       | 2                              | 7, 11                    |
| 7   | 3   | 2   | 1   | 15                       | 2                              | 7, 11                    |
| 8   | 3   | 2   | 1   | 21                       | 2                              | 11, 13                   |
| 9   | 3   | 2   | 1   | $\binom{8}{2}$           | $3 (p \leq 23); 2 (p \geq 29)$ | 11–97                    |
| 12  | 3   | 2   | 1   | $\binom{11}{2}$          | $4 (p \leq 17); 3 (p \geq 19)$ | 13–31                    |
| 15  | 3   | 2   | 1   | $\binom{14}{2}$          | $5 (p = 17); M_\infty = 3$     | 17; rank $_{\mathbb{Q}}$ |
| 7   | 3   | 3   | 0   | 15                       | 1                              | 7, 11                    |

Here  $d := w - c$  and  $\max M_{\text{true}} \leq \binom{n}{d} / \binom{w}{d}$  (the incidence bound [24]) in every case. The table is generated by exhaustive enumeration over  $\mathbb{F}_p^D$ ; scripts and outputs are in [39].

$p$ -dependence. For  $n \leq 8$ ,  $\max M_{\text{true}}$  is identical across all tested  $p > n$ . At  $n \geq 9$ ,  $M_{\text{true}}$  becomes  $p$ -dependent: the  $\mathbb{Q}$ -rank of the combined normal matrix (Proposition 36) determines whether a configuration persists for all large  $p$ , or only for finitely many primes. The asymptotic value  $M_\infty$  (Definition 37) is computed in Remark 38.

Provable bounds.

- (a) Unique decoding ( $c \geq w$ ):  $M_{\text{true}} \leq 1$ , since  $\min \Delta(\text{RS}[n, n-D]) = D+1 = w+c+1 > 2w$ .
- (b) Incidence bound ( $c < w$ ):  $M_{\text{true}} \leq \binom{n}{d} / \binom{w}{d}$  with  $d = w - c$  [24], tight at  $d = 1$  for small  $p$ .
- (c) Average:  $\mathbb{E}[M_{\text{true}}] = \binom{n}{w} (1 - 1/p)^w / p^c$  (Corollary 31), exponentially small at deployment.

Consequence for the error-locator normal approach. The error-locator normal machinery (Theorems 25 and 28) gives tight average bounds ( $\mathbb{E}[M] = \binom{n}{w} / p^c$ ) but operates on  $M_{\text{compat}}$ , which overcounts  $M_{\text{true}}$  by including false-positive supports. The worst-case  $M_{\text{true}}$  is bounded by the incidence bound (polynomial in  $n$ ) and can depend on  $p$ , but this does not affect the FRI soundness theorem (Theorem 20), which uses only the half-threshold CA (Theorem 6).

**Remark 34** (Coset extraction on multiplicative subgroups). When  $L = \langle \omega \rangle$  and  $k = 2$ , the Vieta conditions on an agreement set  $S$  force the characteristic polynomial  $\prod (x - \alpha_i)$  to be lacunary (only leading, sub-leading, and constant terms nonzero). For  $n = 2^K$  and even  $t$ :  $\gcd(t-1, n) = 1$ , so the lacunary polynomial’s root set on  $L$  is a single coset of the subgroup of order  $\gcd(t-1, n)$  (Lidl–Niederreiter [13]; Golomb–Gong [18] Ch. 3–4), giving  $M = O(1)$ . The power-of-2 domain structure that FRI uses eliminates the coset-union complication entirely.

### 7.3 OP2 obstruction and the average-to-worst-case gap

The strongest form of OP2 asks: does  $M_{\text{max}} = 0$  for all centers  $u$  when  $p$  is sufficiently large? The following elementary construction shows this is *false* at FRI-relevant code lengths.

**Proposition 35** (List-size obstruction). *For  $\text{RS}[n, k]$  over  $\mathbb{F}_p$  with  $p \geq 3$ , integer Johnson radius  $w_J := \lfloor n(1 - \sqrt{\rho}) \rfloor$  (where  $\rho = k/n$ ), and integer codimension excess  $c_J = n - k - w_J$ : if  $c_J \leq \lfloor (n-k-1)/2 \rfloor$ , then  $M_{\max} \geq 2$  for all  $p$ .*

*Proof.* Let  $c_1, c_2 \in \text{RS}_k$  with  $d(c_1, c_2) = d_{\min} = n-k+1$ , and let  $S = \{i : c_1(i) \neq c_2(i)\}$  with  $|S| = d_{\min}$ . The hypothesis gives  $2w_J \geq d_{\min}$ , so we can choose  $B_1, B_2 \subset S$  with  $|B_1| = |B_2| = w_J$  and  $B_1 \cup B_2 = S$ . Define  $u$  by:  $u = c_1 = c_2$  on  $L \setminus S$ ;  $u = c_2$  on  $B_1 \setminus B_2$ ;  $u = c_1$  on  $B_2 \setminus B_1$ ;  $u \notin \{c_1, c_2\}$  on  $B_1 \cap B_2$  (feasible since  $p \geq 3$ ). Then  $d(u, c_1) = |B_1| = w_J$  and  $d(u, c_2) = |B_2| = w_J$ , so  $M(u) \geq 2$ .  $\square$

For rate  $\rho = 1/2$ :  $c_J \approx 0.207n$  and  $\lfloor (n-k-1)/2 \rfloor \approx n/4$ . At  $\rho = 1/2$  with  $n$  a power of 2 (the FRI setting), the obstruction applies for  $n \geq 2^5 = 32$  (verified computationally up to  $n = 64$  [39]; fails at  $n = 16$  due to ceiling effects).

*What this means.* The strongest form of OP2 ( $M_{\max} = 0$  for all  $p$ ) is false at deployment-scale FRI. The average-case bounds ( $\mathbb{E}[M] = \binom{n}{w}/p^c$ , Corollary 30) and practical  $M = 0$  at deployment parameters are unaffected. See §8, item 2, for the precise remaining questions.

## 7.4 Asymptotic stabilization of $M_{\text{true}}$

The following result explains the  $p$ -dependent behavior observed in Remark 33: for large  $p$ , the worst-case list size stabilizes at a value determined by the rational rank of a certain normal matrix.

**Proposition 36** ( $\mathbb{Q}$ -rank characterization). *Fix  $n, w, c$  with  $1 \leq c < w \leq n$  and  $D = w + c$ . Let  $E_1, \dots, E_k$  be pairwise-compatible  $w$ -subsets of  $L = \{0, 1, \dots, n-1\}$  (i.e.,  $|E_i \cap E_j| \leq w - c$  for  $i \neq j$ ). Form the combined normal matrix*

$$\mathcal{N} = \begin{pmatrix} \Lambda_{E_1} \cdot x^0 \\ \vdots \\ \Lambda_{E_1} \cdot x^{c-1} \\ \Lambda_{E_2} \cdot x^0 \\ \vdots \\ \Lambda_{E_k} \cdot x^{c-1} \end{pmatrix} \in \mathbb{Z}^{kc \times D},$$

where row  $\Lambda_{E_i} \cdot x^r$  is the coefficient vector of  $\Lambda_{E_i}(x) \cdot x^r$  in the basis  $\{1, x, \dots, x^{D-1}\}$ .

- (a) *Upper bound.* If  $\text{rank}_{\mathbb{Q}}(\mathcal{N}) = D$ , then for all primes  $p$  not dividing any  $D \times D$  minor of  $\mathcal{N}$ , the intersection  $\bigcap_{i=1}^k W_{E_i} = \{0\}$  over  $\mathbb{F}_p$ . In particular, no nonzero syndrome can simultaneously have all  $k$  supports in its list, so  $M_{\text{true}}(s) < k$  for every  $s \neq 0$ .
- (b) *Lower bound.* If  $\text{rank}_{\mathbb{Q}}(\mathcal{N}) < D$  and the configuration is nondegenerate (see proof), then for all primes  $p > \max(n, kw)$ , there exists  $s \in \mathbb{F}_p^D \setminus \{0\}$  such that  $s \in W_{E_i}$  and the Vandermonde solution  $V_{E_i}^{-1}s$  has all entries nonzero, for every  $i = 1, \dots, k$ . Hence  $M_{\text{true}}(s) \geq k$ .

*Proof.* (a) If  $\text{rank}_{\mathbb{Q}}(\mathcal{N}) = D$ , then  $\mathcal{N}$  has a  $D \times D$  submatrix with nonzero determinant  $\Delta \in \mathbb{Z}$ . For any prime  $p \nmid \Delta$ :  $\text{rank}_{\mathbb{F}_p}(\mathcal{N}) = D$ , so the only  $s \in \mathbb{F}_p^D$  satisfying  $\mathcal{N}s = 0$  is  $s = 0$ . Since  $s \in W_{E_i}$  iff the  $c$  normals of  $E_i$  annihilate  $s$ , the intersection  $\bigcap W_{E_i} = \ker(\mathcal{N}) = \{0\}$ .

(b) If  $\text{rank}_{\mathbb{Q}}(\mathcal{N}) = D - \delta$  with  $\delta \geq 1$ , the rational kernel  $K = \ker_{\mathbb{Q}}(\mathcal{N})$  has dimension  $\delta$ . Clearing denominators gives  $\delta$  linearly independent integer vectors in  $\ker(\mathcal{N})$ . For any prime  $p$  larger than the largest prime factor of the denominators (which depends only on  $n$ ), these reduce to  $\delta$  independent vectors in  $\ker_{\mathbb{F}_p}(\mathcal{N})$ , so  $\dim_{\mathbb{F}_p}(\bigcap W_{E_i}) \geq \delta$ .



It remains to show that some  $s$  in the kernel gives all-nonzero Vandermonde solutions. For each support  $E_i$ , the map  $s \mapsto V_{E_i}^{-1}s$  restricted to  $\bigcap W_{E_i}$  is a linear map to  $\mathbb{F}_p^w$ . Call the configuration *nondegenerate* if for every  $(i, j)$ , the linear functional  $s \mapsto (V_{E_i}^{-1}s)_j$  does not vanish identically on  $\bigcap W_{E_i}$ . When nondegenerate: each of the  $kw$  “bad” sets  $\{s : (V_{E_i}^{-1}s)_j = 0\}$  is a proper hyperplane in the  $p^\delta$ -element kernel, and the union of  $kw$  proper hyperplanes omits at least  $p^\delta - kw \cdot p^{\delta-1} = p^{\delta-1}(p - kw) > 0$  points when  $p > kw$ . Nondegeneracy is verified computationally for all configurations in Remark 38; see [39].  $\square$

**Definition 37** (Asymptotic worst-case list size).

$$M_\infty(n, w, c) := \max\{k : \exists k \text{ pairwise-compatible } w\text{-subsets of } [n] \text{ with } \text{rank}_{\mathbb{Q}}(\mathcal{N}) < D\}.$$

By Proposition 36,  $M_{\text{true}}(n, w, c, p) = M_\infty(n, w, c)$  for all sufficiently large primes  $p$ , and  $M_{\text{true}} \leq M_\infty + 1$  for all  $p > n$ .

**Remark 38** (Computed values of  $M_\infty$ ). For  $w = 3, c = 2$  ( $d = 1$ ):

| $n$      | $M_\infty$ | $\lfloor n/3 \rfloor$ | rank-deficient triples | verification                               |
|----------|------------|-----------------------|------------------------|--------------------------------------------|
| $\leq 8$ | 2          | $\leq 2$              | 0                      | exhaustive, nondegenerate                  |
| 9–10     | 2          | 3                     | 0                      | exhaustive; 280 ( $n=9$ ), 2800 ( $n=10$ ) |
| 11       | 3          | 3                     | 1 of 15400             | nondeg.; no rank-def. 4-tuple              |
| 12       | 3          | 4                     | 2 of 61600             | nondeg.; no rank-def. 4-tuple              |
| 13       | 3          | 4                     | 4 of 200200            | no rank-def. 4-tuple                       |
| 14       | 3          | 4                     | 6 of 560560            | no rank-def. 4-tuple                       |
| 15       | 3          | 5                     | 18 of 1401400          | no rank-def. 4-tuple                       |

Rank-deficient triples include both translates of

$$\mathcal{C}_a = \{(a, a+2, a+7), (a+3, a+8, a+10), (a+4, a+5, a+6)\}$$

and non-translate configurations (e.g.,  $\{(0, 7, 8), (2, 6, 10), (4, 5, 12)\}$  at  $n = 13$ ; 13 of 18 configurations at  $n = 15$  are non-translates).

*Algebraic obstruction to  $M_\infty \geq 4$ .* For any rank-deficient triple, the kernel of the  $6 \times 5$  normal matrix is one-dimensional. A fourth disjoint triple  $\{d, e, f\}$  preserves rank 4 iff all three roots satisfy a compatibility cubic. Polynomial division reduces integer compatibility to a divisibility condition  $(Q(u)) \mid (R(u))$  after centering  $u = d - \bar{d}$ . For the translate family:  $(u^2 - 5) \mid 4u$  (independent of  $a$ ), yielding exactly the 9 used points as the only integer solutions. Verified computationally [39]: for every rank-deficient triple at  $n \leq 15$  (including all non-translate configurations), the analogous divisibility condition has solutions only among the 9 used points.

No rank-deficient 4-tuple was found in exhaustive search up to  $n = 15$  or in 10,000 random trials at  $n \in \{18, 21, 30\}$  [39]. Conjecture:  $M_\infty(n, 3, 2) = 3$  for all  $n \geq 11$  and  $M_\infty = 2$  for  $6 \leq n \leq 10$ .

Additional structural results — fiber bounds, pinned-flat uniqueness, Bézout bounds, codimension bounds ( $\dim V(r_0-1, r_1, r_2) \leq w-3$ ), incidence bounds ( $M \leq \binom{n}{d} / \binom{w}{d}$ ), and irreducibility of the remainder hypersurface — are developed in a companion note [24] and provide the correct order-of-magnitude picture of  $M$  as a function of field size and code parameters.

## 7.5 Möbius structure and worst-case bound at $c = 2$

The FRI-relevant case is codimension excess  $c = 2$  at rate  $\rho = 1/2$  (so  $w = n/2 - 2$ ). Here a clean algebraic structure controls how  $M_{\text{true}}$  depends on the syndrome line. Fix  $s_1, s_2 \in \mathbb{F}_p^D$  defining a line  $s(\gamma) = s_1 + \gamma s_2$ , and let

$$M_{\text{true}}(s_1, s_2) := \#\{\gamma \in \mathbb{F}_p : \exists E \in \binom{L}{w} \text{ s.t. } V_E^{-1}(s(\gamma)) \text{ has all entries } \neq 0\}.$$

**Lemma 39** (Möbius structure on cores). *Let  $C \subset L$  with  $|C| = w - 1$  (a core) and  $E = C \cup \{x\}$  for  $x \in L \setminus C$ . As  $x$  varies, the unique  $\gamma_E$  for which the first compatibility equation  $\langle \Lambda_E, s(\gamma) \rangle = 0$  holds is a degree-1 rational function (Möbius transform) of the evaluation point  $\alpha_x$ :*

$$\gamma_E = \frac{A_C \cdot \alpha_x + B_C}{C_C \cdot \alpha_x + D_C}$$

for constants  $A_C, B_C, C_C, D_C \in \mathbb{F}_p$  depending on  $(s_1, s_2)$  and  $C$  but not on  $x$ .

*Proof.* Factor  $\Lambda_E(z) = \Lambda_C(z) \cdot (z - \alpha_x)$ . The coefficient of  $z^j$  in  $\Lambda_E$  is  $\varepsilon_{j-1}(C) - \alpha_x \cdot \varepsilon_j(C)$ , where  $\varepsilon_j(C)$  are the coefficients of  $\Lambda_C$  (with  $\varepsilon_{-1}(C) := 0$ ). Hence the coefficient vector of  $\Lambda_E$  is linear in  $\alpha_x$ . For any  $s \in \mathbb{F}_p^D$ ,  $\langle \Lambda_E, s \rangle = \sum_j \varepsilon_{j-1}(C) s_j - \alpha_x \sum_j \varepsilon_j(C) s_j$  is linear in  $\alpha_x$ . Setting  $\langle \Lambda_E, s_1 + \gamma s_2 \rangle = 0$  and solving for  $\gamma$  gives a ratio of two linear functions of  $\alpha_x$ .  $\square$

**Theorem 40** (Worst-case bound at  $c = 2$ ). *At codimension excess  $c = 2$  over any  $\mathbb{F}_p$  with  $p > n$ ,*

$$M_{\text{true}}(s_1, s_2) \leq M_{\text{compat}}(s_1, s_2) \leq \min\left(p, 2 \binom{n}{w-1}\right)$$

for every line  $(s_1, s_2)$ .

*Proof.* A pair  $(E, \gamma)$  contributes to  $M_{\text{compat}}$  iff  $\langle \Lambda_E, s(\gamma) \rangle = 0$  and  $\langle \Lambda_E \cdot x, s(\gamma) \rangle = 0$  (the two normals at  $c = 2$ ). Decompose  $E = C \cup \{x\}$  with  $|C| = w - 1$  and  $x \in L \setminus C$ . Both inner products are bilinear in  $(\alpha_x, \gamma)$ : linear in  $\alpha_x$  by the shift identity used in Lemma 39, linear in  $\gamma$  by the line parameterization. Eliminating  $\gamma$  between the two bilinear equations yields a polynomial in  $\alpha_x$  of degree  $\leq 2$ .

Two cases. If this polynomial is not identically zero, it has at most two roots, so at most two values of  $\alpha_x$  (and hence at most two pairs  $(\alpha_x, \gamma_E)$ ) arise from this core. Summing over  $\binom{n}{w-1}$  cores:  $M_{\text{compat}} \leq 2 \binom{n}{w-1}$ . If the polynomial vanishes identically, the line  $\ell(\gamma) = s_1 + \gamma s_2$  lies inside the union  $\bigcup_{x \in L \setminus C} W_{C \cup \{x\}}$ , but  $M_{\text{compat}} \leq p = |\mathbb{F}_p|$  holds trivially. Either way the stated bound holds. Since  $M_{\text{true}} \leq M_{\text{compat}}$ , the bound passes to  $M_{\text{true}}$ .  $\square$

The bound complements the incidence bound:  $M_{\text{true}} \leq \min(p, 2 \binom{n}{w-1}, \binom{n}{d} / \binom{w}{d})$  with  $d = w - c$ . The  $2 \binom{n}{w-1}$  factor is loose at small  $n$  (e.g.  $n = 15$ ,  $w = 6$ ,  $c = 2$ : bound 210 vs. empirical  $\max M_{\text{true}} = 5$  [39]) but it is the only  $p$ -independent upper bound currently known at the FRI-relevant codimension  $c = 2$  that is universal in  $(s_1, s_2)$ .

### Sub-Poisson concentration on bin sizes

Theorem 40 bounds the Möbius-image bin size  $M_\gamma(s_1, s_2) := |\{E \in \binom{L}{w} : \gamma_E(s_1, s_2) = \gamma\}|$  by  $2 \binom{n}{w-1} / w$  universally, but this ignores the negative correlation induced by Lemma 39. A sharper *high-probability* bound (over uniform  $(s_1, s_2) \in \mathbb{F}_p^D \times \mathbb{F}_p^D$ ) recovers the birthday rate.

**Theorem 41** (Sub-Poisson tail on  $M_\gamma$  at  $c = 2$ ). *Fix  $\rho = k/n$ ,  $w = n - k - 2$ ,  $D = n - k$ . For any prime  $p > \max(n^2, 2(n - w + 1))$ , with probability  $\geq 1 - 1/p$  over uniform  $(s_1, s_2) \in \mathbb{F}_p^D \times \mathbb{F}_p^D$ :*

$$\max_{\gamma \in \mathbb{F}_p} M_\gamma(s_1, s_2) \leq \frac{\binom{n}{w}}{p} + 4\sqrt{2 \log(p) \cdot \frac{\binom{n}{w}}{p}}.$$

Combined with  $M_\gamma \leq p$  trivially, in the same probability event:

$$\max_{\gamma \in \mathbb{F}_p} M_\gamma(s_1, s_2) \leq C_0 \cdot \sqrt{\binom{n}{w} \cdot \log p}, \quad C_0 \leq 16.$$

*Proof sketch.* Set  $X_E := \mathbf{1}[\gamma_E(s_1, s_2) = \gamma]$ , with marginal  $\Pr_s[X_E = 1] = 1/p$  (Lemma 39 gives a Möbius bijection on the leaf  $\alpha_x$ , so for fixed core  $C$  and uniform  $(s_1, s_2)$ , the value  $\gamma_E$  is uniform on  $\mathbb{F}_p$ ). Pairs  $(E_1, E_2)$  classify by overlap  $r$ : at  $r \leq w - 2$ , the four normals are linearly independent (Theorem 28), so  $\text{Cov}_s(X_{E_1}, X_{E_2}) = 0$ ; at  $r = w - 1$ , the Möbius bijectivity excludes equal images from distinct leaves except on a degenerate-Möbius locus of measure  $\leq 1/p$ , giving  $\text{Cov}_s(X_{E_1}, X_{E_2}) \leq 0$ . Hence  $\{X_E\}$  are pairwise non-positively correlated indicators with  $\text{Var}[M_\gamma] \leq \mathbb{E}[M_\gamma] = \binom{n}{w}/p =: \mu$ .

The standard sub-Poisson tail bound for indicators with non-positive pairwise covariance (Janson [28], Dubhashi–Ranjan [29]; the Bernstein argument requires only the non-positive correlation hypothesis) gives

$$\Pr_s[M_\gamma \geq \mu + t] \leq \exp\left(-\frac{t^2}{2(\mu + t/3)}\right).$$

Setting  $t = 4\sqrt{2 \log(p) \cdot \mu}$  and applying a union bound over  $\gamma \in \mathbb{F}_p$ :

$$\Pr_s\left[\max_\gamma M_\gamma \geq \mu + t\right] \leq p \cdot \exp\left(-\frac{16 \cdot 2 \log p \cdot \mu}{2(\mu + (4/3)\sqrt{2 \log p \cdot \mu})}\right) \leq \frac{1}{p}$$

for  $\mu \geq \log p$ , which holds for  $p > p_{\min}(n)$  as stated. The combined bound with  $M_\gamma \leq p$  follows from the AM–GM step  $\min(\mu, \sqrt{\binom{n}{w} \log p}) \leq \sqrt{\binom{n}{w} \log p}$ . Full proof in companion note [26].  $\square$

The high-probability bound matches the empirical  $0.66 \cdot 1.36^n$  envelope (Remark 42) up to  $\sqrt{\log p}$  factors. A *worst-case* extension over all  $(s_1, s_2)$  is open (§8, the birthday-bound conjecture); the natural higher-moment Markov approach is structurally blocked by configurations where three supports share a common  $(w - c)$ -element core (companion note [26], script in [39]).

**Remark 42** (Empirical exponential growth at the worst case). *At  $c = 2$  and  $w = n/2 - 2$  (the FRI-relevant regime), exhaustive search over  $(s_1, s_2)$  pairs, prime fields, and supports up to  $n = 24$  [39] gives an empirical fit*

$$\max_{s_1, s_2, p} M_{\text{true}}(s_1, s_2; p) \approx c_0 \cdot c_1^n$$

with constants  $c_0 \approx 0.66$ ,  $c_1 \approx 1.36$  in the tested range  $n \in \{8, 10, \dots, 24\}$  (OLS regression of  $\log_2 M_{\text{true}}$  on  $n$  over 9 data points,  $R^2 > 0.99$ ). We report these as point estimates descriptive of the tested range only — with 9 data points spanning 16 in the exponent, the 95% confidence interval on the slope  $\log_2 c_1$  is non-trivial and we make no claim about population constants. The peak prime tracks  $p_{\text{peak}} \approx \sqrt{\binom{n}{w}}$ , the regime where bin sizes (the number of supports projecting to each  $\mathbb{F}_p$  value of the linear part of  $\gamma$ ) become  $\Theta(\sqrt{\binom{n}{w}})$  on average and concentrate within a constant factor of

this average. The data are consistent with  $M_{\text{true}}$  at  $c = 2$  growing super-polynomially in  $n$ , but a rigorous super-polynomial worst-case lower bound from the data alone would require either a longer scaling range or a complementary analytic argument; this remark records the scaling shape on the tested range, not an asymptotic claim. The corresponding ratio  $M_{\text{true}}/p$ , the FRI-relevant quantity, is observed to decay in the tested range ( $M_{\text{true}}/p \approx 0.676^n$  at  $p = p_{\text{peak}}$ ).

## 7.6 Phase diagram across codimension excess $c$

Theorem 40 and Remark 42 characterize the worst-case list size  $M_{\text{true}}$  at  $c = 2$ . The picture across all  $c$  is a clean three-phase diagram:

- $c = 1$  (covering radius):  $\mathbb{E}[M] = \binom{n}{w}/p$  with Poisson dispersion (Corollary 30); the trivial counting bound  $M_{\text{true}} \leq \binom{n}{w}$  holds. The worst-case behaviour at large  $n$  is empirically observed to saturate but a closed-form upper bound is open at  $c = 1$ .
- $c = 2$ : empirical fit  $M_{\text{true}} \approx c_0 c_1^n$  with  $c_0 \approx 0.66$ ,  $c_1 \approx 1.36$  in the tested range  $n \leq 24$  (Remark 42; descriptive, not an asymptotic claim); governed by the Möbius coupling on  $(w-1)$ -cores (Lemma 39).
- $c \geq 3$ : empirically linear in  $n$  at sufficiently large  $p$ ,  $M_{\text{true}} \leq \lfloor (2D-1)/c \rfloor$  where  $D = n - k$ .

The deployment-relevant regime for FRI is  $c = \Theta(n)$  (specifically  $c \approx 0.2n$  above the Johnson radius at rate  $1/2$ ), so the third phase is the one that controls FRI soundness analysis.

**Conjecture 43** (Open-Set Rank Lemma at  $c \geq 3$ ). *For supports  $E_1, \dots, E_m \subset L$  each of size  $w = D - c$  with  $c \geq 3$ , and distinct  $\gamma_1, \dots, \gamma_m \in \mathbb{F}_p^*$ , let  $A \in \mathbb{F}_p^{mc \times 2D}$  be the constraint matrix with row blocks  $[N_{E_i} \mid \gamma_i N_{E_i}]$ , where  $N_{E_i}$  collects the  $c$  error-locator normals from  $E_i$  (Lemma 27). For all  $p$  above an effective threshold  $p_0(n, k, c)$  (polynomial in  $n$ ):*

$$\text{either } \text{rank}(A) = \min(mc, 2D), \quad \text{or } \exists i \text{ with } \langle n_0(E_i), s_2 \rangle = 0 \text{ for all } (s_1, s_2) \in \ker A.$$

Equivalently, every  $(s_1, s_2)$  with  $M_{\text{true}}(s_1, s_2) = m$  satisfies  $m \leq \lfloor (2D-1)/c \rfloor$ .

**Remark 44** (Empirical evidence and small- $p$  counterexamples). *Conjecture 43 has been verified at large  $p$  for  $(n, k)$  up to  $n = 40$  with  $c$  up to 9 via random-syndrome search [27, 39]; selected data points (rate  $1/2$ ,  $\geq 14000$  random configs each):*

| $n$ | $k$ | $c$ | bound $\lfloor (2D-1)/c \rfloor$ | observed max $M_{\text{true}}$ | $p$ tested |
|-----|-----|-----|----------------------------------|--------------------------------|------------|
| 20  | 10  | 5   | 3                                | $\leq 2$                       | large      |
| 24  | 12  | 5   | 4                                | $\leq 2$                       | large      |
| 28  | 14  | 6   | 4                                | $\leq 4$                       | large      |
| 32  | 16  | 7   | 4                                | $\leq 4$                       | large      |
| 36  | 18  | 8   | 4                                | $\leq 4$                       | large      |
| 40  | 20  | 9   | 4                                | $\leq 4$                       | large      |

The observed  $M_{\text{true}}$  is constant ( $\leq 4$ ) at the Johnson-radius codimension for rate  $1/2$ , well within the  $O(n/c)$  envelope.

At small primes, Conjecture 43 can fail due to mod- $p$  algebraic coincidences. Two structural counterexamples are known:

- *Triangle obstruction at  $c = 2$ . For  $n = 8$ ,  $w = 2$ ,  $p = 113$ : the four supports  $\{(3, 6), (5, 6), (3, 5), (0, 1)\}$  realize  $M_{\text{true}} = 4$ , exceeding the would-be bound  $\lfloor (2D-1)/c \rfloor = 3$ . The first three supports form a  $K_3$  on vertices  $\{3, 5, 6\}$ , providing a row dependency that the fourth disjoint support extends. Consistent with the  $c = 2$  phase being exponential, not linear.*

- *Tetrahedron obstruction at  $c = 3$ , small  $p$ .* For  $n = 12$ ,  $w = 3$ ,  $p = 61$ : the four supports  $\{(1, 4, 5), (1, 4, 8), (4, 5, 8), (1, 5, 8)\}$  are exactly the size-3 subsets of  $\{1, 4, 5, 8\}$ , forming a  $K_4$  tetrahedron, and realize  $M_{\text{true}} = 4 > \lfloor (2D - 1)/c \rfloor = 3$ . The lemma fails at this small  $p$  but holds for  $p > p_0(n, k, c)$ .

The general obstruction at general  $c$  is the  $(w+1)$ -clique — all size- $w$  subsets of a  $(w+1)$ -vertex set. These configurations are realizable only at primes below an explicit threshold; the conjecture asserts an effective Schwartz–Zippel-style bound on this threshold.

The connection to FRI: the FRI verifier accepts a low-distance codeword on the evaluation domain only when the syndrome line hits one of  $M_{\text{true}}$  many error patterns. At FRI scale,  $c \approx 0.2n$  exceeds 3, putting us in the third phase of the diagram. At fixed rate,  $\lfloor (2D - 1)/c \rfloor$  is a *constant* in  $n$  (equal to 4 at rate  $1/2$  near the Johnson radius); Conjecture 43 therefore predicts  $M_{\text{true}} = O(1)$  uniformly in code length at current parameters, consistent with the empirical row of the table.

## 8 Open Problems

We list open problems in order of deployment relevance.

### 1. Beating the 1:2 ratio via $\geq 3$ bad $\gamma$ 's or non-CA arguments.

*Status of the standard two-bad- $\gamma$  CA framework* (settled in this paper). Theorem 6 proves  $\varepsilon_{\text{ca}}(C, \delta/2, \delta) \leq 1/|F|$  for any linear code, and the 1:2 ratio is optimal for the standard two-bad- $\gamma$  argument (Remark 14). At *equal* threshold ( $\delta_{\text{hd}} = \delta$ ), Theorem 8 gives the unconditional upper bound  $\varepsilon_{\text{ca}}(C, \delta, \delta) \leq \binom{n}{w}/|F|$ , matched at the equality case  $w = n - k - 1$  for  $|F| > \binom{n}{w}^2$  by Proposition 10 (RS code). Since  $\binom{n}{w} \sim 2^{nH(\delta)}$  grows exponentially in the code length, this upper bound is vacuous against any meaningful security target at FRI scale.

*The open question* is whether a strictly different argument can beat 1:2. Two natural candidates:

- $\geq 3$  bad  $\gamma$  in the CA framework. Three bad  $\gamma$ 's give an overdetermined system (three equations in two unknowns  $(f_1, f_2)$ ), yielding polynomial constraints on  $h_3$  given  $h_1, h_2$ . The additional constraint does not tighten the standard inclusion-exclusion bound beyond the two- $\gamma$  result, but tighter analyses tracking the algebraic relations among  $h_1, h_2, h_3$  are not ruled out.
- *Non-CA proximity proofs* bypassing the CA framework entirely, pursued in concurrent companion work in preparation [2, 3]. Specific quantitative claims await those papers' completion and are not asserted here.

### 2. Worst-case $M_{\text{true}}$ : structure and bounds.

Proposition 35 shows  $M_{\text{max}} \geq 2$  at FRI parameters ( $n \geq 32$ , rate  $1/2$ ) for all  $p$ ; the strongest form of the  $M = 0$  conjecture is false. The average-case bounds  $\mathbb{E}[M] = \binom{n}{w}/p$  (Theorem 25) for  $c = 1$  and  $\mathbb{E}[M] = \binom{n}{w}/p^c$  (Corollary 30) for general  $c$ , both with Poisson dispersion  $\text{Var}[M] \approx \mathbb{E}[M]$  (via pairwise independence of error-locator normals, Theorem 28), show  $M = 0$  for a random syndrome once  $p^c > \binom{n}{w}$ . At deployment ( $c \approx 0.2n$ ,  $p = 2^{31}$ ):  $\mathbb{E}[M] \leq 2^{-5.3n} \approx 0$ .

*Evaluation syndromes are false positives.* Remark 33 shows that the error-locator normal compatibility count  $M_{\text{compat}}$  reaches  $\binom{n-1}{w-1}$  at evaluation syndromes  $s_\alpha = (1, \alpha, \dots, \alpha^{D-1})$ . But these  $\binom{n-1}{w-1}$  “compatible” supports all produce weight-1 errors ( $v_\alpha = 1$ ,  $v_\beta = 0$  for  $\beta \neq \alpha$ ), so  $M_{\text{true}}(s_\alpha) = 0$  for  $w \geq 2$ . The  $\binom{n-1}{w-1}$  is a failure of the error-locator normal machinery to enforce non-zero error values, not a genuine list-size obstruction.

*Provable bounds on  $M_{\text{true}}$ .* At  $c \geq w$  (within the unique decoding radius):  $M_{\text{true}} \leq 1$ . At  $c < w$ :  $M_{\text{true}} \leq \binom{n}{d} / \binom{w}{d}$  with  $d = w - c$  (the incidence bound [24]).

*p-dependence and asymptotic stabilization.* Proposition 36 proves that  $\max M_{\text{true}}$  stabilizes at the value  $M_\infty$  (Definition 37) for all sufficiently large  $p$ , and Remark 38 tabulates computed values. Corollary 31 gives the exact formula  $\mathbb{E}[M_{\text{true}}] = \binom{n}{w}(1 - 1/p)^w/p^c$ .

What remains open: (i) a closed-form for  $M_\infty(n, w, c)$  (the data at  $w = 3$ ,  $c = 2$  suggest  $M_\infty = 3$  for all  $n \geq 11$ ; see Remark 38); (ii) whether any worst-case list-size bound can improve the proximity gap beyond the half-threshold  $\delta/2$ ; (iii) the *birthday-bound conjecture*: extending Theorem 41 from a high-probability bound (over the random parameters  $(s_1, s_2)$ ) to a uniform worst-case bound. The high-probability form

$$\max_{\gamma} M_{\gamma}(s_1, s_2) \leq \frac{\binom{n}{w}}{p} + 4\sqrt{2 \log(p) \cdot \frac{\binom{n}{w}}{p}} \quad \text{w.p.} \geq 1 - 1/p$$

is proved in Theorem 41 via a Bernstein-style sub-Poisson tail on the variance bound from Lemma 39. The conjectural worst-case version is

$$\max_{s_1, s_2} \max_{\gamma \in \mathbb{F}_p} M_{\gamma}(s_1, s_2) \leq C_1 \cdot \frac{\binom{n}{w}}{p} \quad (\text{uniform in } (s_1, s_2)).$$

If true, this combines with  $M_{\gamma} \leq p$  to yield  $\max M_{\text{true}} = O(\sqrt{\binom{n}{w}}) = O(2^{n/2})$ , matching the  $1.41^n$  envelope of Remark 42 and bridging the gap between the universal bound  $2\binom{n}{w-1}$  (Theorem 40) and the empirical  $0.66 \cdot 1.36^n$ .

*Status of the worst-case extension.* The natural higher-moment + Markov + union-bound approach for promoting Theorem 41 to all  $(s_1, s_2)$  is structurally blocked. The  $k$ -wise extension of the pairwise-independence theorem (Theorem 28) is empirically false at  $c = 2$  for  $k \geq 3$ : three supports  $E_1, E_2, E_3$  sharing a common  $(w-c)$ -element core  $C^*$  have all pairwise overlaps  $|E_i \cap E_j| = w - c$  (saturating the pairwise-independence hypothesis) but their six normals all lie in the principal ideal  $(\Lambda_{C^*}) \subset \mathbb{F}_p[x]/(x^D)$  of codimension  $w - c$ , hence dimension  $2c$ , so the rank is at most  $2c < 3c$ . Joint compatibility probability is then  $1/p^{2c}$  rather than the independent prediction  $1/p^{3c}$ , making the third moment super-Gaussian and ruling out sub-Gaussian Markov tails. Empirical verification: companion script [39] finds thousands of rank-deficient triples for  $n \in \{12, 14, 16\}$ . This is a higher-order analog of Lemma 39 (the two-leaf case at  $c = 2$ ): the multi-leaf coupling at  $c \geq 2$  requires sequence-school cross-correlation tools beyond what the error-locator normal machinery alone provides.

*The  $c \geq 3$  regime is qualitatively different (§7.6).* The same rank-deficiency that creates exponential growth at  $c = 2$  disappears at  $c \geq 3$ , and Conjecture 43 (Open-Set Rank Lemma) predicts  $M_{\text{true}} \leq \lfloor (2D - 1)/c \rfloor$  at sufficiently large  $p$ , linear in  $n$ . The large- $c$  regime  $c = \Theta(n)$  falls in this third phase, where the conjecture predicts  $M_{\text{true}} = O(1)$  at the Johnson radius and the empirical data confirms  $M_{\text{true}} \leq 4$  for rate  $1/2$  across  $n = 20$  to  $40$ . Resolving Conjecture 43 for  $c \geq 3$  would close the worst-case worry at current code lengths even without progress on the  $c = 2$  birthday-bound conjecture above.

3. **Odd  $k$ .** The  $k = 2^m$  condition excludes odd  $k$  (Remark 5). Standard FRI uses  $k = 2^m$ , so this is not a practical blocker.

**Characteristic 2.** Appendix B (Theorem 64) gives the same  $nR/|F| + (1 - \delta/2)^q$  FRI soundness bound over binary tower fields, by replacing the multiplicative  $\pm\sqrt{y}$  decomposition with the additive subspace-polynomial fold  $s(x) = x^2 + \beta x$ . The half-threshold CA bound (Theorem 6) is characteristic-independent (it uses only linearity and Hamming distances); only the coupling lemma requires the fold replacement, and the additive fold preserves the isomorphism  $\text{RS}_k \xrightarrow{\sim} \text{RS}_{k/2}^2$  for even  $k$ . This covers Binius-style settings on  $\mathbb{F}_{2^{128}}$  tower fields without additional assumptions.

## 9 Deployment Roadmap

**Scope of the deployment-applicable results.** The FRI, STIR, and WHIR protocol structures — folding maps, Merkle commitments, prover messages, and verifier checks — remain unchanged. The above-Johnson soundness coverage is: FRI (in both characteristics) and circle FRI at theorem level; STIR at theorem level in odd characteristic, with the characteristic-2 STIR analogue sketched in Remark 66; WHIR at conjecture level in odd characteristic (Conjecture 54; the proximity-substitution licensing is rigorously proved via Lemma 55, the surrounding iteration-level induction of WHIR [11, §5] is left as follow-up work), with the characteristic-2 WHIR analogue similarly sketched. The only parameter that changes is the query count  $q$ , and increasing  $q$  increases proof size, verifier hashing, calldata, and wrapper-circuit cost linearly in the query-heavy portion of the verifier. The theorem bounds interactive oracle proof of proximity (IOPP) soundness of the interactive proximity test; non-interactive knowledge-soundness (via Fiat–Shamir, DEEP-quotienting, polynomial commitment wrapping) requires standard but separate compilation arguments. We summarize what our results provide in concrete parameters.

|                                                                                | Conjectured (zero-loss CA)         | Proved (this paper)                                                                                                                     |
|--------------------------------------------------------------------------------|------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------|
| Status                                                                         | Unproven                           | Theorem (FRI/circle-FRI/STIR); Conjecture (WHIR)                                                                                        |
| Security budget ( $q = 128$ , $R = 20$ , $\delta = 0.4$ , rate $\rho = 1/2$ ): |                                    |                                                                                                                                         |
| Commit-phase (BabyBear base, $ F  = 2^{21}$ )                                  | $\sim 7$ bits                      | $\sim 7$ bits                                                                                                                           |
| Commit-phase (BabyBear <sub>4</sub> , $ F  = 2^{124}$ )                        | $\sim 100$ bits                    | $\sim 100$ bits                                                                                                                         |
| Query-phase ( $q = 128$ , $\delta$ vs $\delta/2$ )                             | $(1-\delta)^{128} \approx 2^{-94}$ | $(1-\delta/2)^{128} \approx 2^{-41}$                                                                                                    |
| Total (BabyBear base)                                                          | $\min(7, 94) = 7$                  | $\min(7, 41) = 7$                                                                                                                       |
| Total (BabyBear <sub>4</sub> , extension)                                      | $\min(100, 94) = 94$               | $\min(100, 41) = 41$                                                                                                                    |
| Query overhead                                                                 | baseline                           | $\sim 2.3\times$ baseline at $\delta = 0.4$                                                                                             |
| Protocol-structure changes                                                     | none                               | none ( $q$ recalibrated)                                                                                                                |
| Protocols covered                                                              | FRI, STIR, WHIR                    | <b>FRI, STIR (theorem), WHIR (conjecture)</b> (see matrix) <sup>†</sup>                                                                 |
| Characteristic                                                                 | odd only                           | <b>FRI: any (theorem); STIR: char<math>\neq 2</math> theorem; WHIR: char<math>\neq 2</math> conjecture; char 2 sketched<sup>†</sup></b> |
| Survives $(\delta, \delta)$ -CA exponential constant (§8)                      | no                                 | yes                                                                                                                                     |

<sup>†</sup> FRI soundness is proved as a formal theorem in both characteristics (Theorem 20 for char  $\neq 2$ ; Theorem 64 for char  $= 2$ ). STIR’s commit term is  $(2k+2n)R/|F|$  in general and simplifies to  $(2k+2)R/|F|$  at  $\rho \geq 1/2$  (Theorem 48), tightening to  $(k+n)R/|F|$  (or  $(k+2)R/|F|$  at  $\rho \geq 1/2$ ) at deployment scale; the security rows above are for FRI. The STIR/WHIR extension to characteristic 2 follows from the same argument (the additive coupling lemma replaces the multiplicative one and the per-coordinate batch-CA helper extends without change), as stated in Remark 66; we do not package it as a separate theorem or corollary.

**Proof-coverage matrix.** The precise scope of what is formally proved vs. deferred:

| Protocol / composition step                                          | Status                                                                      | Reference                |
|----------------------------------------------------------------------|-----------------------------------------------------------------------------|--------------------------|
| FRI above Johnson, $\text{char}(F) \neq 2$                           | <b>fully proved</b>                                                         | Theorem 20               |
| FRI above Johnson, $\text{char}(F) = 2$                              | <b>fully proved</b>                                                         | Theorem 64               |
| Half-threshold CA, any linear code (RVW13 [31])                      | classical                                                                   | Theorem 6                |
| Batch CA, any field                                                  | <b>fully proved</b>                                                         | Theorem 46               |
| STIR soundness, $\text{char}(F) \neq 2$                              | <b>fully proved</b>                                                         | Theorem 48               |
| STIR soundness, $\text{char}(F) = 2$                                 | <b>sketched</b>                                                             | Remark 66                |
| WHIR, linear-subcode CA core, $\text{char}(F) \neq 2$                | <b>fully proved</b>                                                         | Corollary 50             |
| WHIR soundness / affine-subcode composition, $\text{char}(F) \neq 2$ | <i>conjectural</i> (proximity-substitution proved via Lemma 55)             | Conjecture 54            |
| WHIR soundness, $\text{char}(F) = 2$                                 | <b>sketched</b>                                                             | Remark 66                |
| Circle-FRI variant (Stwo)                                            | <b>fully proved</b>                                                         | Theorem 72               |
| Fiat-Shamir / non-interactive soundness, multiplicative FRI          | <b>fully proved</b>                                                         | Theorem 76; Appendix D   |
| Fiat-Shamir compilation, char-2 / circle / STIR                      | <i>follows by the same RBR template; not packaged as a separate theorem</i> | —                        |
| DEEP-FRI knowledge soundness                                         | <b>fully proved</b>                                                         | Theorem 78; Corollary 79 |
| Composition with Plonkish lookups (logUp)                            | <b>not addressed</b>                                                        | future work              |

The “safety net” covers the full proximity-testing stack at theorem level for FRI in both characteristics, circle FRI, and STIR in odd characteristic, and at conjecture level for WHIR in odd characteristic (Theorems 20, 64, 72, 48; Conjecture 54); the characteristic-2 STIR/WHIR extension is sketched in Remark 66. Fiat–Shamir compilation to non-interactive soundness (Theorem 76) and DEEP-FRI knowledge soundness for the polynomial commitment scheme (Theorem 78, Corollary 79) are fully proved. The remaining “not addressed” item (Plonkish/logUp composition) is the arithmetization layer, which is protocol-specific and orthogonal to proximity testing.

**Adoption requires no changes to protocol structure:** in the settings formalized here, the FRI/STIR/WHIR prover algorithm, round structure, folding equations, and proof format are unchanged; only the query count  $q$  is recalibrated. For Binius-style settings over binary tower fields ( $\mathbb{F}_{2^{128}}$ ), Appendix B proves this at theorem level for additive FRI: no protocol modification beyond using subspace-polynomial folding  $s(x) = x^2 + \beta x$  in place of  $y = x^2$  is required — this is the canonical characteristic-2 analogue introduced by Ben-Sasson–Bentov–Horesh–Riabzev [20] and further developed in additive FRI and Binius [21]; our contribution is a *soundness theorem* for this fold above the Johnson radius, not the fold itself. The same additive-fold substitution yields the sketched characteristic-2 STIR/WHIR extension in Remark 66. The verifier implementation must still be parameterized by the chosen query count  $q$ , so changing  $q$  changes the number of Merkle openings and consistency checks. Remaining engineering tasks:

1. **Interactive security parameter tuning:** the commit-phase term  $nR/|F|$  matches BCHKS up to constants; on BabyBear<sub>4</sub> ( $|F| = p^4 \approx 2^{124}$ ), this gives  $\sim 100$  bits. The binding constraint is the query-phase term  $(1-\delta/2)^q$ : at  $q = 128$  this gives  $\sim 41$  bits for  $\delta = 0.4$ . Reaching 128-bit interactive query-phase security at  $\delta = 0.4$  requires  $q \approx 398$  (a  $\sim 3.1\times$  increase from typical  $q = 128$ ), but BabyBear<sub>4</sub> remains capped at  $\sim 100$  total interactive bits by the commit term; the non-interactive setting is stricter because of the Fiat–Shamir  $Q$ -loss in the next item. Multi-rate reference:

| $\rho$ | $\delta_J = 1 - \sqrt{\rho}$ | $\delta$ | $q$ (conj. zero-loss) | $q$ (ours) |
|--------|------------------------------|----------|-----------------------|------------|
| 1/2    | 0.293                        | 0.30     | $\sim 195$            | $\sim 427$ |
| 1/2    | 0.293                        | 0.40     | $\sim 136$            | $\sim 311$ |
| 1/4    | 0.500                        | 0.55     | $\sim 87$             | $\sim 216$ |
| 1/8    | 0.646                        | 0.70     | $\sim 58$             | $\sim 161$ |

All rows target  $\lambda = 100$  bits of query-phase security, with  $\delta$  chosen strictly above the rate-dependent Johnson radius  $\delta_J = 1 - \sqrt{\rho}$  so that our theorem applies (the open zone is  $\delta_J < \delta < 1 - \rho$ ). Below  $\delta_J$ , BCHKS’s zero-loss bound is the relevant baseline and our reduced-threshold result offers no improvement; the table therefore omits below-Johnson operating points.

2. **Non-interactive parameter tuning:** Fiat–Shamir compilation multiplies the interactive RBR error by the adversary’s random-oracle query budget  $Q$  and adds the



Merkle/random-oracle binding term  $3(Q^2+1)/2^\kappa$  (Theorem 76). Table 3 gives the resulting parameter budget for  $Q = 2^{64}$  and  $\delta = 0.4$ . The entries are generated by `notes/scripts/deployment_params.py`; “commit” means no finite  $q$  can reach the target because the  $Q \cdot nR/|F|$  term already exceeds the target.

| Config                  | $ F $ bits | $q$ now | Interactive | NI now | $q_{64}$ | $q_{100}$ | $q_{128}$ |
|-------------------------|------------|---------|-------------|--------|----------|-----------|-----------|
| BabyBear base           | 31         | 128     | 6.7         | −57.3  | commit   | commit    | commit    |
| BabyBear <sub>4</sub>   | 124        | 128     | 41.2        | −22.8  | commit   | commit    | commit    |
| BabyBear <sub>5</sub>   | 155        | 128     | 41.2        | −22.8  | 398      | commit    | commit    |
| BabyBear <sub>7</sub>   | 217        | 128     | 41.2        | −22.8  | 398      | 510       | 597       |
| Mersenne31 <sub>4</sub> | 124        | 128     | 41.2        | −22.8  | commit   | commit    | commit    |
| Goldilocks <sub>2</sub> | 128        | 128     | 41.2        | −22.8  | commit   | commit    | commit    |
| Binary tower $2^{128}$  | 128        | 128     | 41.2        | −22.8  | commit   | commit    | commit    |

**Table 3.** Non-interactive FRI security under the loose Theorem 76 accounting, with  $n = 2^{20}$ ,  $R = 20$ ,  $\delta = 0.4$ ,  $Q = 2^{64}$ ,  $\kappa = 256$  except BabyBear<sub>7</sub> where  $\kappa = 384$  so the hash-binding term is not the bottleneck. “Interactive” and “NI now” are proved bits at  $q = 128$ ; negative NI bits mean the theorem gives no non-trivial non-interactive bound at that query count after the  $Q$ -loss. The  $q_\lambda$  columns give the required query count for  $\lambda$ -bit non-interactive security when the commit and hash terms permit it.

3. **Verifier cost update:** increasing  $q$  is a parameter change with real cost. It leaves the IOP structure intact but increases Merkle path verifications, FRI consistency checks, proof bytes, calldata, and any recursive/wrapper verifier circuit work roughly linearly in  $q$ .
4. **Reproducibility:** the table above is generated by `notes/scripts/deployment_params.py`; the paper artifact, Lean status, proximity-gap figure, and saved computational outputs are documented in `REPRODUCING.md`. The structural companion notes [24, 25, 26, 27] cited from this paper are public in [39] at the paths recorded in their bib entries.
5. **Formal verification:** the half-threshold CA bound (Theorem 6) and the equal-threshold CA upper bound (Theorem 8) are formally verified in Lean 4 with Mathlib at the statement-faithful level (zero `sorry`, zero project-level axioms beyond Mathlib’s standard ones); the half-threshold theorem also has a paper-form  $|F|$ -probability corollary. Building blocks for the FRI coupling chain (`coupling_pointwise`, `coupling_counting`, `rs_iso_forward/surj`), the char-2 generalized pairing, and the batch-CA fixed-rest union bound are present in the Lean library. Remaining: a concrete `RSIsomorphismWitness` instance for the multiplicative-coset FRI domain, a faithful import of BCIKS ’20 Theorem 1.2, and the protocol-event wrapper that converts the rational-form soundness skeleton into the paper-form probability bound. See the companion repository’s `paper1/lean/STATUS.md` for the per-theorem state [39].

## 10 Conclusion and Outlook

We have proved an unconditional FRI soundness theorem above the Johnson radius for plain Reed–Solomon codes on the evaluation domains used in practice — in every characteristic, under the structural hypotheses ( $k = 2^m$ ,  $L$  admitting a fixed-point-free involution) that are standard for FRI. The bound  $\varepsilon_{\text{FRI}} \leq nR/|F| + (1-\delta/2)^q$  is a theorem, not a conjecture, for multiplicative FRI, additive FRI, and circle FRI; the same reduced-threshold method also yields a theorem-level odd-characteristic extension to STIR and a *conjectural* odd-characteristic extension to WHIR (Conjecture 54, with the proximity-layer substitution proved via Lemma 55); the characteristic-2 STIR/WHIR extension is sketched in Remark 66 (Section 9). The only protocol-level change is the query count  $q$ ; the prover/verifier equations remain unchanged, while verifier cost scales with the chosen  $q$ .

**The CA framework is settled.** The half-threshold CA bound is the classical 1:2 proximity-gap inequality of RVW13 [31]. Our contribution is applying it *above the Johnson radius*: the chain  $\delta/2 < (1-\rho)/2 \leq \delta_J$  (Remark 1) places the protocol into BCIKS’s zero-loss regime after round 1, where BCIKS [7] locks the distance. The 1:2 ratio is optimal at FRI scale (Remark 14), and the  $M = 0$  approach is obstructed at FRI-relevant code lengths (Proposition 35). The  $2\times$  query overhead is intrinsic to any CA-based proof at current parameters.

**Beyond the CA framework.** Every published FRI soundness proof factors through a CA lemma (§1.7). Reducing the  $2\times$  overhead requires proving FRI soundness *without* this intermediate step — for example, via a direct algebraic analysis of the fold map on multiplicative subgroups, or via character-sum techniques exploiting that  $L = \langle \omega \rangle$  is a cyclic group. Cross-correlation distributions from the Helleseth–Golomb–Gong tradition [17, 18] and Weil-type bounds on incomplete character sums [19] are natural candidates. Concurrent companion work in preparation [2, 3] pursues this direction; specific results await those papers’ completion. Whether FRI admits a non-CA soundness proof with overhead  $< 2\times$  at FRI scale is a key remaining question for FRI soundness theory.

**Practical bottom line.** Prior to this work, above the Johnson radius: zero bits of proven FRI soundness. After:  $\varepsilon_{\text{FRI}} \leq nR/|F| + (1-\delta/2)^q$ , giving  $\sim 41$  interactive query bits at  $q = 128$  or  $\sim 128$  interactive query bits at  $q \approx 398$  for  $\delta = 0.4$ . Total security is still capped by the commit term (for example  $\sim 100$  interactive bits on BabyBear<sub>4</sub> at  $n = 2^{20}$ ,  $R = 20$ ). The  $2\times$  query overhead vs. the conjectured zero-loss baseline is optimal within the CA framework; against what is actually proved, our bound is the tightest available.

**Roadmap.** Three follow-ups: (a) complete the Lean 4 formalization [39] (the half-threshold CA theorem is statement-faithful in Lean with a paper-form  $|F|$ -probability corollary; building blocks for the FRI coupling, char-2, and batch-CA chains are present, and the remaining work is the concrete pairing instances, a faithful BCIKS ’20 import, and the protocol-event wrapper for  $\text{Pr}[\text{accept}]$ ); (b) tight bounds on worst-case  $M_{\text{true}}$  and whether they can improve the proximity gap beyond  $\delta/2$  (§8, item 2); (c) non-CA proximity statements pursued in concurrent work in preparation by the same authors [2, 3], addressing the question of whether the  $< 2\times$  regime is reachable outside the CA framework. Specific quantitative claims await those papers’ completion.

**Concurrent work in preparation.** Two companion papers in preparation by the same authors [2, 3] pursue non-CA proximity statements complementary to this paper’s CA-framework boundary. Specific quantitative claims await those papers’ completion and are not asserted here.

## A Extension to STIR and WHIR

We show that the half-threshold CA bound (Theorem 6) extends to the batch setting required by STIR [10] and to the constrained RS subcodes required by WHIR [11]. This makes the safety net directly applicable to deployed STIR/WHIR-based systems.

### A.1 Batch Correlated Agreement

STIR and its variants test proximity of multiple functions simultaneously by taking a random linear combination and checking the result. This requires a *batch* CA bound.

**Definition 45** (Batch Correlated Agreement). For code  $C$  over domain  $L$ ,  $B$  functions  $f_1, \dots, f_B : L \rightarrow F$ , and independent uniform  $\gamma_1, \dots, \gamma_B \in F$ :

$$\varepsilon_{\text{ca}}^{(B)}(C, \delta_{\text{hd}}, \delta_{\text{nt}}) = \max_{f_1, \dots, f_B} \Pr_{\gamma_1, \dots, \gamma_B} \left[ \Delta \left( \sum_{i=1}^B \gamma_i f_i, C \right) \leq \delta_{\text{hd}} \text{ and } \exists i : \Delta(f_i, C) > \delta_{\text{nt}} \right].$$

**Theorem 46** (Half-Threshold Batch CA). For any integer  $B \geq 1$  and any linear code  $C$  over  $F$ :

$$\varepsilon_{\text{ca}}^{(B)}(C, \delta/2, \delta) \leq \frac{B}{|F|}.$$

*Proof.* Fix  $f_1, \dots, f_B$  with  $\Delta(f_j, C) > \delta$  for some index  $j$ . We bound the probability over  $(\gamma_1, \dots, \gamma_B)$  that  $g = \sum_i \gamma_i f_i$  is  $(\delta/2)$ -close to  $C$ .

**Step 1: Conditioning.** Fix all  $\gamma_i$  for  $i \neq j$ , so that  $h = \sum_{i \neq j} \gamma_i f_i$  is a fixed function. Then  $g = \gamma_j f_j + h$ , which is linear in the single random variable  $\gamma_j$ .

**Step 2: Pairwise argument.** Call  $\gamma_j$  “bad” if  $\exists c_\gamma \in C$  with  $|S_\gamma| \geq (1 - \delta/2)n$  where  $S_\gamma = \{x \in L : \gamma_j f_j(x) + h(x) = c_\gamma(x)\}$ .

For two distinct bad values  $\gamma_j, \gamma'_j$ : on  $S_\gamma \cap S_{\gamma'}$  (size  $\geq (1 - \delta)n$ ), we have  $(\gamma_j - \gamma'_j)f_j = c_\gamma - c_{\gamma'}$ , so  $f_j$  agrees with the codeword  $(c_\gamma - c_{\gamma'})/(\gamma_j - \gamma'_j) \in C$  on  $\geq (1 - \delta)n$  points, giving  $\Delta(f_j, C) \leq \delta$ . This contradicts  $\Delta(f_j, C) > \delta$ . Hence at most 1 bad  $\gamma_j$ .

**Step 3: Union bound.** The bound in Step 2 holds for each fixed  $h$  (i.e., each fixed choice of  $\gamma_i, i \neq j$ ). Averaging over the other  $\gamma_i$ :

$$\Pr_{\gamma_1, \dots, \gamma_B} [\Delta(g, C) \leq \delta/2 \mid f_j \text{ is } \delta\text{-far}] \leq \frac{1}{|F|}.$$

By the union bound over all indices  $j$  with  $\Delta(f_j, C) > \delta$  (at most  $B$  of them):

$$\varepsilon_{\text{ca}}^{(B)}(C, \delta/2, \delta) \leq \frac{B}{|F|}. \quad \square$$

**Remark 47** (Individual vs. joint distance premise). Definition 45 uses the individual distance condition  $\exists i : \Delta(f_i, C) > \delta$  rather than the joint distance  $\Delta_{\text{joint}}((f_1, \dots, f_B), C^{=B}) > \delta$ . The two premises are incomparable: individual distance can be small even when joint distance is large (disjoint per-coordinate error sets), and conversely. The individual form is the one tailored to the STIR use-case below.

In the STIR proof (Theorem 48, Strategy B), the relevant condition is  $\Delta(f^{(i-1)}|_{L^{(i)}}, \text{RS}_{k_i}) > \delta$ : this is the individual-distance condition applied to one specific function, which is what Theorem 46 requires. The pairwise argument in the proof of Theorem 6 is invoked directly.

## A.2 STIR Soundness Above Johnson

We show that our half-threshold CA is a drop-in replacement for the conjectured zero-loss CA in the STIR soundness proof [10].

**STIR protocol recap.** STIR [10] is an IOPP for Reed–Solomon proximity testing. Fix a degree schedule  $k_0 = k, k_1, \dots, k_R$  with  $k_i \mid k_{i-1}$  and domains  $L = L^{(0)} \supset L^{(1)} \supset \dots \supset L^{(R)}$  with  $|L^{(i)}| = n/2^i$  (the standard radix-2 schedule sets  $k_i = k_{i-1}/2$ , matching FRI). At each round  $i$ , the verifier sends a random evaluation point  $z_i \in F \setminus L^{(i)}$  and a random batching coefficient  $\gamma_i$ . The prover responds with  $v_i = f^{(i-1)}(z_i)$  and commits the quotient oracle  $q_i(x) = (f^{(i-1)}(x) - v_i)/(x - z_i)$ , then sets  $f^{(i)} = q_i + \gamma_i \cdot f^{(i-1)}|_{L^{(i)}}$  (the batched next-round function). The key soundness step is: if  $f^{(i-1)} \notin \text{RS}_{k_{i-1}}$ , then with high probability over  $\gamma_i$ , the batched  $f^{(i)}$  is far from  $\text{RS}_{k_i}$  — this is precisely the CA bound applied to the pair  $(q_i, f^{(i-1)}|_{L^{(i)}})$  with randomness  $\gamma_i$ .

**Theorem 48** (STIR Soundness Above Johnson). *For  $\text{RS}[F, L, k]$  with  $\text{char}(F) \neq 2$ ,  $k = 2^m$ ,  $L$  closed under negation,  $\delta_J < \delta < 1 - \rho$ ,  $\Delta(f^{(0)}, \text{RS}_k) > \delta$ , and the standard STIR parameter regime  $|F| \geq 2|L^{(i)}|$  at every round (in particular  $|F|/|L| = \Omega(1)$ ), satisfied at every deployment instance considered in this paper, see §9): STIR with  $R$  rounds and  $q$  base-level queries (each opening one full path) satisfies*

$$\Pr[\text{STIR accepts}] \leq \frac{(2k + 2n)R}{|F|} + \left(1 - \frac{\delta}{2}\right)^q,$$

where the  $n = |L^{(0)}|$  term comes from the BCIKS bad-challenge bound at rounds  $i \geq 2$  of Strategy B (which gives  $\leq |L^{(i)}|$  bad scalars per round) and is dominated by the OOD term  $k_{i-1}$  when  $\rho \geq 1/2$ , in which case the commit term simplifies to  $(2k + 2)R/|F|$ . At the stronger deployment-scale regime  $|F|/|L^{(i)}| \rightarrow \infty$  at fixed rate (always satisfied where  $|L^{(i)}|/|F|$  is exponentially small), the commit-phase term tightens to  $(k + n)R/|F|$  in general (or  $(k + 2)R/|F|$  when  $\rho \geq 1/2$ ).

*Proof.* The proof follows the same Strategy A / Strategy B dichotomy as Theorem 20, adapted to STIR's round structure.

**Strategy A: the prover commits honest quotients and batching at every round.** The prover computes honest quotients  $q_i(x) = (f^{(i-1)}(x) - f^{(i-1)}(z_i))/(x - z_i)$  and honest batching  $f^{(i)} = q_i + \gamma_i \cdot f^{(i-1)} \upharpoonright_{L^{(i)}}$ . No adaptive freedom remains beyond the randomness of  $(z_i, \gamma_i)$ . We bound the probability that  $f^{(R)} \in \text{RS}_{k_R}$  by a per-round union bound.

At round  $i$ , suppose  $f^{(i-1)} \notin \text{RS}_{k_{i-1}}$ . Then:

- Over  $z_i \in F \setminus L^{(i)}$ : by the STIR function-quotienting lemma [10] (Lemma 4.4), if  $f^{(i-1)} \notin \text{RS}_{k_{i-1}}$  then the quotient  $q_i(x) = (f^{(i-1)}(x) - f^{(i-1)}(z_i))/(x - z_i)$  inherits non-codeword status except on a set of  $\leq k_{i-1}$  degenerate values of  $z_i$  (roots of a nonzero degree- $(\leq k_{i-1})$  polynomial in  $z_i$  arising from the interpolation identity). Lemma 4.4 of [10] is unconditional polynomial algebra; no proximity-gap or list-decodability hypothesis is invoked.
- Over  $\gamma_i$ : for a non-degenerate  $z_i$ , the batched syndrome of  $f^{(i)} = q_i + \gamma_i \cdot f^{(i-1)} \upharpoonright_{L^{(i)}}$  is degree-1 in  $\gamma_i$  (nonzero), so it vanishes for at most 1 value of  $\gamma_i$ .

Per round (Strategy A only), under the regime  $|F| \geq 2|L^{(i)}|$  in the theorem hypothesis:  $\Pr[f^{(i)} \in \text{RS}_{k_i} \mid f^{(i-1)} \notin \text{RS}_{k_{i-1}}] \leq k_{i-1}/(|F| - |L^{(i)}|) + 1/|F| \leq 2k_{i-1}/|F| + 1/|F| \leq (2k + 1)/|F|$ , where the  $k_{i-1}/(|F| - |L^{(i)}|)$  term comes from [10, Lemma 4.4] ( $z_i \in F \setminus L^{(i)}$  landing in the at-most- $k_{i-1}$  degenerate locus), and the  $1/|F|$  from the  $\gamma_i$  Schwartz–Zippel over  $F$ . Adding the Strategy B OOD-sampling-failure term  $\binom{\ell_i}{2} \cdot k_{i-1}/(|F| - |L^{(i)}|) \leq O(k/|F|)$  from [10, Lemma 4.5] (in the Johnson zone  $\delta/2 < \delta_J$  which the post-halving threshold occupies, the classical Johnson list-size bound [10, Theorem 3.4] gives  $\ell_i = O(1/\sqrt{1 - \rho})$  uniform in  $n$ ): per-round commit-phase budget  $\leq (2k + 2)/|F|$ , matching the theorem statement.

By the union bound over  $R$  rounds:  $\Pr[f^{(R)} \in \text{RS}_{k_R}] \leq (2k + 2)R/|F|$ .

*Tightening to  $(k + 2)R/|F|$  at deployment scale.* The factor of 2 in  $2k_{i-1}/|F|$  comes from  $1/(|F| - |L^{(i)}|) \leq 2/|F|$  under  $|F| \geq 2|L|$ ; under the stronger deployment-scale regime  $|F|/|L^{(i)}| \rightarrow \infty$  at fixed rate (e.g.  $|F| \approx 2^{124}$ ,  $|L| \leq 2^{20}$ , ratio  $\geq 2^{100}$ ), the factor of 2 collapses to  $1 + o(1)$  and the commit-phase term tightens to  $(k + 2)R/|F|$  as in the alternative form of the theorem statement.

*Note on the  $k$ -factor.* The  $k_{i-1}/|F|$  term is intrinsic to STIR's out-of-domain evaluation: the quotient involves evaluating a degree- $(< k_{i-1})$  polynomial at the random point  $z_i$ , producing  $\leq k_{i-1}$  degenerate values. On extension fields (e.g., BabyBear<sub>4</sub> with  $|F| = p^4 \approx 2^{124}$ ), this term is negligible ( $k/|F| \approx 2^{-105}$ ). On the base field ( $|F| = p \approx 2^{31}$ ), it yields  $\sim 7$  bits — comparable to the BCHKS commit-phase bound. The improvement of this paper over BCHKS for STIR is therefore in the query phase:  $(1 - \delta/2)^q$  replaces the conjectured  $(1 - \delta)^q$ .

**Strategy B: the prover deviates from honest computation at some round  $i$ .** The

optimal adversarial strategy: send a codeword  $q_i \in \text{RS}_{k_i}$  as the quotient, then compute honestly from round  $i+1$  onward (all subsequent functions land in the code). Two checks threaten at round  $i$ :

1. *Out-of-domain (OOD) check*: the verifier sends  $z_i \in F \setminus L^{(i)}$  and receives the claimed value  $v_i = f^{(i-1)}(z_i)$ , then checks  $f^{(i-1)}(y) - v_i = (y - z_i)q_i(y)$  at  $q$  random points  $y \in L^{(i)}$ . In Strategy B,  $f^{(i-1)}$  is an arbitrary prover-committed oracle (not necessarily a polynomial). The soundness here decomposes into two STIR-internal pieces:

- *OOD-sampling failure* ( $z_i$  landing on a list-collision): bounded by Lemma 4.5 of [10] as  $\binom{\ell_i}{2} \cdot (k_{i-1}/(|F| - |L^{(i)}|))^s$ , where  $\ell_i$  is the list-decoding bound for  $\text{RS}[F, L^{(i)}, k_{i-1}]$  at the relevant threshold and  $s$  is the number of OOD samples (we use  $s = 1$  in the standard schedule). This is unconditional once the underlying RS code is list-decodable at that threshold, which holds in the Johnson zone by the classical Johnson bound ([10] Theorem 3.4) without any proximity-gap conjecture; our regime  $\delta/2 < (1-\rho)/2 \leq \delta_J$  post-halving sits unconditionally in this zone.
- *Shift-query failure* (the  $q$  random checks on  $L^{(i)}$  all miss the deviation set): by Lemma 4.4 of [10] (unconditional function-quotienting), an adversarial pair  $(v_i, q_i)$  with  $q_i \in \text{RS}_{k_i}$  that satisfies the OOD identity at fewer than the threshold number of points of  $L^{(i)}$  produces a deviation set of relative size  $\geq 1 - k_{i-1}/|L^{(i)}|$ , giving a catch probability  $\geq 1 - (k_{i-1}/|L^{(i)}|)^q = 1 - \rho_i^q$ , where  $\rho_i := k_{i-1}/|L^{(i)}|$ . With STIR's halving schedule ( $k_i = k/2^i$ ,  $|L^{(i)}| = n/2^i$ ),  $\rho_i = 2\rho$ . At deployment rates  $\rho \leq 1/4$ , this gives  $\rho_i^q \leq (1/2)^q$  and the shift-query failure is dominated by  $(1-\delta/2)^q$  for every  $\delta < 1$ . At  $\rho = 1/2$  ( $\rho_i = 1$ ) the shift-query bound is vacuous, and the consistency check  $(1-\delta/2)^q$  is what carries the query-phase soundness; the commit-phase OOD term is unchanged. The  $\rho = 1/2$  case is therefore covered by the same bound, with the shift-query check contributing nothing.

The shift-query bound is what enters our final budget; the OOD-sampling term contributes only  $O(k/|F|)$  via Lemma 4.5 and is absorbed into the  $(2k + 2n)R/|F|$  commit-phase term ( $(2k + 2)R/|F|$  at  $\rho \geq 1/2$ ,  $(k + n)R/|F|$  at deployment scale).

2. *Batch-CA at the deviation round*:  $f^{(i)} = q_i + \gamma_i \cdot f^{(i-1)}|_{L^{(i)}}$  must be  $(\delta/2)$ -close to  $\text{RS}_{k_i}$ . Since  $q_i \in \text{RS}_{k_i}$ , this reduces to bounding the bad  $\gamma_i$  that make  $f^{(i-1)}|_{L^{(i)}}$  fold close to  $\text{RS}_{k_i}$ . The bound depends on the round:

- Round  $i = 1$ :  $\Delta(f^{(0)}, \text{RS}_{k_0}) > \delta$ , so Theorem 6 (half-threshold CA) gives at most 1 bad  $\gamma_1$ , failure  $\leq 1/|F|$ .
- Rounds  $i \geq 2$  (assuming honest play in rounds  $1, \dots, i-1$ ): the state is at distance  $> \delta/2$  from  $\text{RS}_{k_{i-1}}$  by Remark 21, and threshold  $\gamma = \delta/2$  lies inside BCIKS's zero-loss regime. The BCIKS proximity gap [7] (Theorem 1.2, Eq. (1.1)) gives at most  $|L^{(i)}|$  bad  $\gamma_i$ , failure  $\leq |L^{(i)}|/|F|$ .

Per-round failure  $\leq |L^{(i)}|/|F|$  in the worst case (round  $i \geq 2$ ), contributing the  $2|L^{(0)}|R/|F| = 2nR/|F|$  summand of the  $(2k + 2n)R/|F|$  total in the theorem statement. (At  $\rho \geq 1/2$ ,  $|L^{(i)}| \leq k_{i-1}$ , and this term is dominated by the OOD term in item (1), simplifying to  $(2k + 2)R/|F|$ .)

As in Theorem 20 (Remark 21), the effective distance  $\delta/2$  is preserved at rounds  $2, \dots, R$ : since  $\delta/2 < (1-\rho)/2 \leq \delta_J$  (Remark 1), the BCIKS proximity gap [7] (Theorem 1.2, Eq. (1.1)) at threshold  $\gamma = \delta/2$  bounds the bad challenges to  $\leq |L^{(i)}|$  per round.

Combining items (1) and (2) per round, and taking the union over  $R$  rounds, the per-round commit-phase contribution is  $\leq (2k_{i-1} + 2|L^{(i)}|)/|F|$ . Summing over rounds and using  $k_{i-1} \leq k$  and  $|L^{(i)}| \leq |L^{(0)}| = n$ :

$$\Pr[\text{Strategy B}] \leq \frac{(2k + 2n)R}{|F|} + (1 - \delta/2)^q.$$

At rate  $\rho \geq 1/2$  (where  $|L^{(i)}| \leq k_{i-1}$  at every round) this simplifies to  $(2k+2)R/|F| + (1-\delta/2)^q$ ; at the deployment-scale regime  $|F|/|L^{(i)}| \rightarrow \infty$  the factor of 2 in  $2k_{i-1}/|F|$  collapses, giving  $(k+n)R/|F| + (1-\delta/2)^q$  in general (and  $(k+2)R/|F| + (1-\delta/2)^q$  when  $\rho \geq 1/2$ ).

**Combining.** Every adaptive prover either commits honest quotients (Strategy A) or deviates (Strategy B); both strategies yield the same commit-phase bound, and only Strategy B has the query term, so

$$\Pr[\text{accept}] \leq \frac{(2k+2n)R}{|F|} + (1-\delta/2)^q,$$

matching the theorem statement.  $\square$

**Remark 49** (Comparison with conjectured STIR bound). *The original STIR analysis [10] assumes zero-loss CA ( $\varepsilon_{\text{ca}}(C, \delta, \delta) = O(1)/|F|$ , conjectured) and achieves  $\varepsilon = O(R)/|F| + (1-\delta)^q$ . Our proved STIR bound has two differences:*

- *Commit phase:  $(2k+2)R/|F|$  unconditionally (and  $(k+2)R/|F|$  at deployment scale, see Theorem 48) instead of  $O(R)/|F|$ . The  $k$ -factor comes from the OOD polynomial evaluation at the random point  $z_i$  (degree  $\leq k_{i-1}$ , giving  $k_{i-1}$  degenerate values). This is intrinsic to STIR's quotient structure and absent in plain FRI. On extension fields ( $|F| = p^4$ ), this term is negligible; on the base field ( $|F| = p$ ), it dominates.*
- *Query phase:  $(1-\delta/2)^q$  instead of  $(1-\delta)^q$ , the same  $\sim 2\times$  overhead as for FRI.*

*For BabyBear<sub>4</sub> ( $|F| \approx 2^{124}$ ,  $k = 2^{19}$ ,  $R = 20$ ): the commit term gives  $\approx 2^{-100}$  bits (negligible), and the query term is the bottleneck as for FRI. On the base field BabyBear ( $|F| \approx 2^{31}$ ): the commit term gives only  $\sim 7$  bits, comparable to BCHKS. The improvement over BCHKS for base-field STIR is entirely in the query phase.*

### A.3 WHIR Extension

WHIR [11] achieves super-fast verification by working with *constrained* Reed–Solomon codes: linear subcodes  $C' \subseteq \text{RS}_k$  cut out by additional linear constraints, and affine constraint sets  $\text{RS}_k \cap \{f : \langle a_j, f \rangle = b_j\}$  for the inhomogeneous (OOD-evaluation, multilinear-evaluation) cases. The CA bound (Theorem 6) uses only one property of the code:

*Linearity* (closure under  $F$ -linear combinations).

This is preserved by every *linear* subcode of  $\text{RS}_k$  over the same domain  $L$ . We make the linear-subcode reduction explicit (Corollary 50), then build up to the full WHIR soundness theorem (Conjecture 54) via a per-iteration analysis that handles affine constraints through the sumcheck weight polynomial rather than through CA on affine codes.

**Corollary 50** (Linear-subcode CA). *For any integer  $B \geq 1$ , any threshold  $\delta$ , and any **linear** subcode  $C' \subseteq \text{RS}_k$  over domain  $L$ :*

$$\varepsilon_{\text{ca}}(C', \delta/2, \delta) \leq 1/|F|, \quad \varepsilon_{\text{ca}}^{(B)}(C', \delta/2, \delta) \leq B/|F|.$$

*Proof.* The proofs of Theorems 6 and 46 invoke linearity in exactly one place: constructing  $g_1, g_2 \in C$  from  $h_1, h_2 \in C$  via  $F$ -linear combinations. Linear subcodes are closed under such combinations, so the proofs go through verbatim.  $\square$

**Remark 51** (Why this does not extend to affine subcodes). *For an affine subcode  $C' = g_0 + C_0$  with  $g_0 \notin C_0$ : the construction  $g_2 = (c_\gamma - c_{\gamma'})/(\gamma - \gamma')$  in the proof of Theorem 6 produces an element of  $C_0$  (a difference, not a sum), not of  $C'$ . Hence we cannot conclude  $g_2 \in C'$ . WHIR*

addresses this by enforcing affine constraints at a different layer (the sumcheck weight polynomial), not by extending CA to affine codes. The CA bound is only invoked on the ambient  $\text{RS}_{k/2^i}$  at each folding step.

### WHIR protocol structure: what we replace

We adopt the notation of [11] (Construction 5.1). One WHIR iteration consists of:

- (i) **OOD sample.** Verifier sends  $z \in F \setminus L$ ; prover responds with a claimed evaluation  $y = f(z)$  and continues with the augmented constraint set  $\{(z_j, y_j)\}_j \cup \{(z, y)\}$ .
- (ii) **Sumcheck folding.** Verifier sends folding randomness  $\alpha \in F$ ; the constrained RS code  $C^{(i)}$  folds via the sumcheck weight-polynomial machinery (Definition 4.5 of [11]) to a constrained RS code  $C^{(i+1)}$  of half the dimension on the half-domain  $L^{(i+1)}$ .
- (iii) **Shift queries.** Verifier samples  $t$  i.i.d. points from  $L^{(i+1)}$  and checks consistency of the folded function with the prover's commitment.

After  $M$  iterations the protocol terminates with  $q$  i.i.d. final queries on the smallest folded domain.

**Where MCA enters in [11]:** only at step (ii), the sumcheck folding. WHIR's analysis (Theorem 4.20 of [11]) bounds the probability that a  $\delta$ -far function  $f$  folds to a  $\delta$ -close folded function via the MCA error of the constrained RS code at threshold  $\delta$ . Above the Johnson radius, this requires Conjecture 4.12 of [11] (zero-loss MCA in the open zone). *This is the conjecture we eliminate.* The OOD sample (i) and shift queries (iii) do not invoke MCA; their soundness is bounded directly by polynomial-identity testing and Lemma 19 respectively.

### Substitution: half-threshold CA + proved MCA in the Johnson regime

We replace WHIR's MCA at threshold  $\delta$  with a two-stage argument:

- *Iteration 1, first folding round only:* apply our half-threshold CA (Theorem 6) on the linear ambient  $\text{RS}_k$  via Corollary 50. This drops the effective distance from  $\delta$  (premise) to  $\delta/2$  (conclusion), with  $\leq 1/|F|$  probability of escape. *This is the only invocation of half-threshold CA in the entire WHIR proof.*
- *All subsequent folding rounds (iteration 1 rounds  $2, \dots, r_1$  and all folds of iterations  $2, \dots, M$ ):* the effective distance  $\delta/2$  satisfies  $\delta/2 < (1-\rho)/2 \leq \delta_J$  (Remark 1), so the protocol enters BCIKS's zero-loss regime (the proved Johnson regime; strictly larger than the unique-decoding ball). WHIR's MCA Corollary 4.11 (proved unconditionally in [11] for  $\delta < \delta_J$ ) applies, giving  $\leq n/|F|$  bad challenges per round. We deliberately do *not* re-apply half-threshold CA at later folds: doing so would compound the halving to  $\delta/4, \delta/8, \dots$

The key structural point: *WHIR's Conjecture 4.12 is invoked only above Johnson; below Johnson, MCA is proved in [11].* Our one-time half-threshold step at iteration 1's first fold brings the effective distance below  $(1-\rho)/2 \leq \delta_J$ , so all subsequent folds (in iteration 1 and all later iterations) use the proved MCA, never the conjecture.

### Affine constraints: handled by sumcheck, not by CA

WHIR's OOD constraints define an affine subcode  $C' = g_0 + C_0$  where  $C_0 \subseteq \text{RS}_k$  is linear and  $g_0$  encodes the constraint values. The CA bound (Corollary 50) applies only to the linear part, but in WHIR the constraints are not enforced via CA: they are absorbed into the sumcheck weight polynomial  $\text{wt}(X)$  (Definition 4.5 of [11]). The folding step (ii) then asks the prover to produce a

folded function  $f'$  such that the sumcheck identity

$$\sum_{x \in L^{(i)}} \text{wt}(x) \cdot f^{(i)}(x) = (\text{wt-folded sumcheck claim})$$

holds round by round. This identity is degree- $\leq 2$  in the sumcheck round variable, with soundness  $\leq 2/|F|$  per round (Schwartz–Zippel). The CA bound enters only at the proximity-preserving step that maps  $f^{(i)} \rightarrow f^{(i+1)}$ , where the underlying code is the linear  $\text{RS}_{k/2^i}$ . Affine vs. linear is therefore handled by separation of concerns: **sumcheck verifies constraint satisfaction, CA verifies proximity.**

### Per-iteration error budget

**Lemma 52** (Per-iteration commit-phase error, conditional). *Under the above setup and conditional on the cross-iteration distance preservation invariant of Conjecture 54 (which is what currently keeps Theorem 20’s analogue a conjecture for WHIR), one WHIR iteration with  $r$  folding rounds has commit-phase soundness error (failures from random alpha/sumcheck challenges, not from query sampling)*

$$\epsilon_{\text{commit}}^{\text{iter}} \leq \frac{k}{|F|} + \frac{r n}{|F|} + \frac{2r}{|F|},$$

with terms: OOD ( $k/|F|$ ), folding rounds at  $\leq n$  bad challenges each ( $rn/|F|$ ), and sumcheck rounds ( $2r/|F|$ ). Shift queries are handled separately (Remark 53).

*Proof.* OOD ( $k/|F|$ ). The verifier’s  $z \in F \setminus L$  is uniform; an adversarial prover with non-codeword  $f$  can fool the OOD check only if  $z$  is a root of a fixed degree- $< k$  polynomial (the difference between  $f$  and the prover’s claimed evaluation extended as a low-degree interpolant). At most  $k$  such roots, giving  $\leq k/|F|$ .

*Folding rounds ( $rn/|F|$ ).* Iteration 1’s first fold uses half-threshold CA: the input is  $\delta$ -far, Corollary 50 gives at most 1 bad  $\alpha$  for  $(\delta/2)$ -closeness ( $\leq 1/|F|$ , tighter than  $n/|F|$ ). All other folds (iteration 1’s rounds  $2, \dots, r$  and all folds of iterations  $\geq 2$ ) use the BCIKS proximity gap [7] at threshold  $\delta/2 < (1-\rho)/2 \leq \delta_J$  (the input is  $(\delta/2)$ -far by distance preservation across iterations, Remark 21; we deliberately do *not* re-apply the half-threshold CA at iterations  $\geq 2$ , which would compound the distance halving). BCIKS gives  $\leq n$  bad challenges per round,  $\leq rn/|F|$  per iteration. We state  $rn/|F|$  as a uniform upper bound covering all  $r$  folds (loose by  $\frac{n-1}{|F|}$  for iteration 1’s first fold).

*Sumcheck ( $2r/|F|$ ).* Per-round Schwartz–Zippel on the degree- $\leq 2$  sumcheck identity gives  $2/|F|$ ; total  $2r/|F|$ .

Union over the three commit-phase failure modes gives the stated bound.  $\square$

**Remark 53** (Shift queries and query-phase max-over-strategies). *The prover commits to a strategy before seeing verifier randomness. Adversarial strategies (mutually exclusive choices) and their query-phase failure probabilities:*

- *Strategy A: honest fold throughout. The cumulative oracle  $f^{(M)}$  is  $(\delta/2)$ -far by distance preservation. Final  $q$  queries miss with probability  $\leq (1-\delta/2)^q$  (Lemma 19).*
- *Strategy B: commit a codeword  $f^{(i+1)} \in C^{(i+1)}$  in iteration  $i$  instead of the honest fold; then continue honestly. The discrepancy between the committed codeword and the honest fold is  $(\delta/2)$ -far on  $L^{(i+1)}$ . Iteration- $i$  shift queries miss this with probability  $\leq (1-\delta/2)^t$ . (Subsequent iterations’ shift queries pass automatically: codewords fold to codewords, so consistency is preserved from iteration  $i+1$  onward; final queries also pass since  $f^{(M)}$  is a codeword.)*



- *Strategy B-multi*: deviate at multiple iterations  $\{i_1, \dots, i_t\}$ . The events “shift queries miss in iteration  $i_j$ ” are independent across iterations (different randomness), so joint miss  $\leq (1-\delta/2)^{t \cdot |\{i_j\}|} \leq (1-\delta/2)^t$ .

The prover picks ONE strategy; the verifier’s failure probability is the maximum over strategies, not the sum:

$$\Pr[\text{query-phase miss}] \leq \max((1-\delta/2)^q, (1-\delta/2)^t) = (1-\delta/2)^q,$$

using  $t \geq q$  to make the shift-query bound tighter. No factor of  $M$  arises: a single chosen strategy can only escape its single failure mode.

**Conjecture 54** (WHIR Soundness Above Johnson, conjectural). For  $\text{RS}[F, L, k]$  with  $k = 2^m$ ,  $L$  admitting a fixed-point-free involution,  $\delta_J < \delta < 1-\rho$ , an initial constrained Reed–Solomon code  $C^{(0)} \subseteq \text{RS}_k$  defined by linear and affine constraints (multilinear-evaluation and OOD), and the standard WHIR parameter regime of [11, §5] ( $|F| \geq 2|L|$  at every iteration, multilinear-extension degree  $m_i \leq \log_2 k$ , list-decoding parameter  $\ell_{i,0} \leq O(1/\sqrt{1-\rho})$  in the Johnson zone, per-iteration shift count satisfying  $t_{i-1} \cdot \ell_{i,0} = O(r_i)$  as in [11, §5]): **provided the prover-committed initial oracle**  $f^{(0)} : L \rightarrow F$  **satisfies**  $\Delta(f^{(0)}, C^{(0)}) > \delta$ , WHIR [11] (Construction 5.1) with  $M$  iterations,  $r_i$  folding rounds in iteration  $i$ ,  $t \geq q$  shift queries per iteration, and  $q$  final queries should satisfy

$$\Pr[\text{WHIR accepts}] \leq \frac{Mk + (n+2)R}{|F|} + \left(1 - \frac{\delta}{2}\right)^q,$$

where  $R = \sum_{i=1}^M r_i$  is the total number of FRI-folding rounds across all iterations.

*Status.* The MCA-orthogonality of the OOD and shift terms in WHIR Theorem 5.2’s four-term decomposition is established rigorously in Lemma 55; this licenses replacing the MCA-bearing folding terms  $\varepsilon_{0,s}^{\text{fold}} + \varepsilon_{i,s}^{\text{fold}}$  by our half-threshold-CA + BCIKS chain at the term level, and removes any dependency on WHIR Conjecture 4.12 from the proximity layer. However, completing the proof to a formal theorem requires re-running the iteration-level induction of [11, §5]’s Theorem 5.2 with the substituted folding-term bound across all  $M$  iterations, which involves checking that the case-(a)/case-(b) recursion and constraint-set accumulation in WHIR’s argument remain sound under the substitution. We leave this re-derivation to subsequent work; the conjecture above is the bound the substitution would deliver if the iteration-level analysis goes through unchanged.

*Sketch / partial proof.* The conjecture is what one would obtain by substituting our half-threshold-CA + BCIKS chain for the MCA-bearing folding terms in WHIR’s Theorem 5.2; the proximity-layer substitution is rigorous (Lemma 55); the surrounding iteration-level analysis from [11, §5] is stated here in sketch form for context but is not re-derived in full. We retain the case-split structure and Merkle-binding analysis that motivate the iteration-level argument.

**Case split on ambient vs. constrained distance.** Two cases:

*Case (a):*  $\Delta(f^{(0)}, \text{RS}_k) > \delta$ . The ambient RS distance is  $> \delta$ , so Corollary 50 applies directly: half-threshold CA on the linear ambient  $\text{RS}_k$  at iteration 1’s first fold drops the effective distance to  $\delta/2$  with  $\leq 1/|F|$  escape (subsumed by the  $rn/|F|$  uniform bound of Lemma 52).

*Case (b):*  $\Delta(f^{(0)}, \text{RS}_k) \leq \delta$  but  $\Delta(f^{(0)}, C^{(0)}) > \delta$ . Some witness  $g \in \text{RS}_k$  with  $d(f^{(0)}, g) \leq \delta n$  exists, but every such witness violates at least one constraint defining  $C^{(0)} \subsetneq \text{RS}_k$ . We must show that, even though the proximity layer cannot reject (the ambient distance is small), the sumcheck/OOD layer rejects.

*Pre-query case-determination.* The case split is made on the *committed* oracle  $f^{(0)} : L \rightarrow F$  before any verifier randomness. The Merkle commitment binds  $f^{(0)}$  as a function on all of  $L$  at commit time (see [20, 11] for the standard binding analysis); the ambient distance  $\Delta(f^{(0)}, \text{RS}_k) =$

$\min_{g \in \text{RS}_k} d(f^{(0)}, g)/n$  is then a property of the committed function, independent of which positions are subsequently queried. The prover cannot retroactively shift  $f^{(0)}$  between Case (a) and Case (b) based on query outcomes (modulo the standard collision term  $3(Q^2+1)/2^\kappa$  for  $\kappa$ -bit Merkle binding under  $Q$  random-oracle queries, absorbed into the Fiat–Shamir budget).

*Black-box reduction to WHIR’s sumcheck/OOD soundness.* Conditional on Case (b), we invoke *as black boxes* the per-iteration bounds in WHIR’s main round-by-round soundness theorem ([11, Theorem 5.2 of Construction 5.1]), under the standard WHIR parameter regime of [11, §5] as instantiated in our deployment table §9 (extension fields  $|F| = p^d$  with  $d \geq 4$ ,  $|F| \gg \text{poly}(k)$ , multilinear extension degree  $m_i \leq \log_2 k$ , Johnson-zone list-decoding parameter  $\ell_{i,0} \leq O(1/\sqrt{1-\rho})$ ):

- *Out-of-domain (OOD) bound*  $\varepsilon_i^{\text{out}} \leq 2^{m_i} \ell_{i,0}^2 / (2|F|)$ , proved in [11] via Lemma 4.25 (Schwartz–Zippel / list-decoding on a degree- $m_i$  power map). At the deployment regime above,  $\varepsilon_i^{\text{out}} = O(k/|F|)$ ; we write  $k/|F|$  for this term throughout.
- *Shift / sumcheck-consistency bound*  $\varepsilon_i^{\text{shift}} \leq (1-\delta_{i-1})^{t_{i-1}} + \ell_{i,0}(t_{i-1}+1)/|F|$ , proved via Claim 5.4 of [11] (a polynomial identity in the combination randomness  $\gamma_i$ ) plus a standard random-point miss bound; we write  $2r_i/|F|$  for the commit-phase  $\ell_{i,0}(t_{i-1}+1)/|F| = O(r_i/|F|)$  contribution at the deployment regime, and absorb the  $(1-\delta_{i-1})^{t_{i-1}}$  query-phase contribution into our  $(1-\delta/2)^q$  term (Remark 53).

The constants  $k$  and  $2r_i$  above are the WHIR-paper figures at the standard deployment regime ([11, §5]); we do not re-derive them. **Crucially, neither bound invokes WHIR Conjecture 4.12 or the MCA error  $\text{err}^*$ .** Tracing dependencies through the proof of Theorem 5.2 of [11], the round- $i$  commit-phase escape probability is decomposed into four terms  $\varepsilon_{0,s}^{\text{fold}}, \varepsilon_i^{\text{out}}, \varepsilon_i^{\text{shift}}, \varepsilon_{i,s}^{\text{fold}}$  (the equation immediately preceding the proof of [11, Theorem 5.2]; we adopt their term names here). Each term has a derivational lineage in [11] that we trace in Lemma 55 below; the conclusion is that the OOD term  $\varepsilon_i^{\text{out}}$  and the shift term  $\varepsilon_i^{\text{shift}}$  have no MCA dependency, while  $\varepsilon_{0,s}^{\text{fold}}$  and  $\varepsilon_{i,s}^{\text{fold}}$  do, and the latter are exactly what the half-threshold CA + BCIKS zero-loss chain of §5 replaces (one halving step at iteration 1’s first fold, which lands  $\delta/2 < \delta_J$  by Remark 1, then BCIKS distance-locking for all subsequent folds). The OOD and shift/sumcheck layers are imported unchanged.

**Lemma 55** (MCA-orthogonality of OOD and shift terms in WHIR Theorem 5.2). *In the proof of [11, Theorem 5.2]:*

- (i)  $\varepsilon_i^{\text{out}}$  is bounded in [11, Lemma 4.25] by Schwartz–Zippel applied to a single degree- $m_i$  polynomial in the OOD challenge  $z_i \in F$ ; its proof invokes only  $|F|$ , the multilinear-extension degree  $m_i$ , and the list-decoding parameter  $\ell_{i,0}$  via the Johnson list-size bound ([11, Theorem 3.4], classical Johnson). No MCA conjecture, no Conjecture 4.12, no  $\text{err}^*$  ever appears in the chain.
- (ii)  $\varepsilon_i^{\text{shift}}$  is bounded in [11, Claim 5.4] by a polynomial-identity argument in the batching randomness  $\gamma_i$  plus an elementary Schwartz–Zippel estimate; its derivation cites [11, Lemma 4.4 (function-quotienting)] (unconditional polynomial algebra) and [11, Theorem 3.4 (Johnson list-size)]. No MCA conjecture appears.
- (iii)  $\varepsilon_{0,s}^{\text{fold}}$  and  $\varepsilon_{i,s}^{\text{fold}}$  are the proximity-preservation / mutual-CA terms; they are bounded in [11, Theorem 4.20] by the MCA error  $\text{err}^*$ , which [11] controls in turn via Conjecture 4.12. These are the only two terms in the decomposition that invoke MCA.

Hence, replacing  $\varepsilon_{0,s}^{\text{fold}} + \varepsilon_{i,s}^{\text{fold}}$  by our half-threshold-CA + BCIKS chain leaves  $\varepsilon_i^{\text{out}}$  and  $\varepsilon_i^{\text{shift}}$  unchanged and removes every dependency on Conjecture 4.12.

*Proof of Lemma 55.* The lemma is a tracing exercise on the proof of [11, Theorem 5.2]. (i) The chain  $\varepsilon_i^{\text{out}} \leftarrow [11, \text{Lemma 4.25}] \leftarrow [11, \text{Theorem 3.4 (Johnson)}]$  has no node referencing [11, Conjecture 4.12 / Theorem 4.20]; the only inputs are  $|F|$ ,  $m_i$ ,  $\ell_{i,0}$ , and the unconditional Johnson

list-size bound. (ii) The chain  $\varepsilon_i^{\text{shift}} \leftarrow [11, \text{Claim 5.4}] \leftarrow [11, \text{Lemma 4.4, Theorem 3.4}]$  similarly has no MCA-bearing node. (iii) For the folding terms,  $\varepsilon_{0,s}^{\text{fold}}, \varepsilon_{i,s}^{\text{fold}} \leftarrow [11, \text{Theorem 4.20}] \leftarrow [11, \text{Conjecture 4.12}]$  via the proof of [11, Theorem 4.20], which invokes mutual correlated agreement at the iteration’s pre-fold distance. These three derivational chains are disjoint in [11] (no shared lemma other than the Schwartz–Zippel lemma, which is unconditional), so substituting half-threshold CA for the MCA-bearing chain in (iii) leaves (i)–(ii) intact.  $\square$

*Black-box term substitution.* Rather than re-derive WHIR’s per-iteration recursion from scratch, we apply Lemma 55 as a *term-substitution lemma* on top of WHIR’s existing Theorem 5.2 derivation in [11]. Specifically:

- (1) Take the round- $i$  commit-phase escape budget of [11, Theorem 5.2], which has the four-term form  $\varepsilon_{0,s}^{\text{fold}} + \varepsilon_i^{\text{out}} + \varepsilon_i^{\text{shift}} + \varepsilon_{i,s}^{\text{fold}}$  as recalled above.
- (2) By Lemma 55,  $\varepsilon_i^{\text{out}}$  and  $\varepsilon_i^{\text{shift}}$  have no MCA dependency and we import them unchanged from [11] (with the deployment-regime evaluation  $k/|F|$  and  $2r_i/|F|$  stated above).
- (3) The MCA-bearing folding terms  $\varepsilon_{0,s}^{\text{fold}} + \varepsilon_{i,s}^{\text{fold}}$  are bounded by [11, Theorem 4.20] via  $\text{err}^*$ , which [11] controls via Conjecture 4.12. We *replace* this MCA chain with the half-threshold-CA chain of §5: at iteration 1’s first fold, half-threshold CA (Theorem 6) drops the effective distance from  $> \delta$  to  $> \delta/2$  at cost  $\leq 1/|F|$ ; for every subsequent fold (in iteration 1 rounds  $2, \dots, r_1$  and *all* folds of iterations  $2, \dots, M$ ), the post-halving distance  $\delta/2$  satisfies  $\delta/2 < \delta_J$  (Remark 1) so the BCIKS proximity gap [7] applies in its proved zero-loss regime, contributing  $\leq n/|F|$  bad challenges per fold (this is exactly the role  $\varepsilon_{i,s}^{\text{fold}}$  would have played, but now without invoking Conjecture 4.12). Summed over folds:  $r_i n/|F|$  at iteration  $i$ .
- (4) Substituting (3) into (1) gives the per-iteration commit-phase budget  $\varepsilon_{\text{commit},i}^{\text{iter}} \leq \varepsilon_i^{\text{out}} + \varepsilon_i^{\text{shift}} + \varepsilon_{\text{fold-our},i}^{\text{fold}} \leq k/|F| + 2r_i/|F| + r_i n/|F|$ .

This per-iteration term substitution is licensed by Lemma 55: the OOD and shift/sumcheck terms in WHIR’s derivation do not depend on what bound is used for the folding terms. Whether the surrounding induction in the proof of [11, Theorem 5.2] (per-iteration accumulation across  $M$  iterations + case split between ambient-close-but-constraint-violating versus ambient-far inputs) *would* carry through unchanged with our  $\varepsilon_i^{\text{fold-our}}$  in place of the MCA-bearing  $\varepsilon_{0,s}^{\text{fold}} + \varepsilon_{i,s}^{\text{fold}}$  is the conjectural step we leave to follow-up work; this paper does not re-derive WHIR’s iteration-level analysis under the substitution. The conjecture above states the bound that the substitution *would* deliver if WHIR’s iteration-level induction is found to be substitutable in this sense.

*Total error.* Summing per-iteration commit-phase budgets across  $M$  iterations and adding the query-phase term:

$$\Pr[\text{accept}] \leq \sum_{i=1}^M \varepsilon_{\text{commit},i}^{\text{iter}} + (1-\delta/2)^q \leq \frac{Mk + (n+2)R}{|F|} + (1-\delta/2)^q,$$

where  $\sum_i (k + r_i n + 2r_i) = Mk + (n+2)R$  for the per-round terms, and the query-phase  $(1 - \delta/2)^q$  comes from Remark 53 (max over adversarial strategies, not sum; using  $t \geq q$ ).  $\square$

**Remark 56** (Result type and scope of Conjecture 54). *Conjecture 54 is a **conjectural WHIR soundness statement above Johnson**, accompanied by Lemma 55 which rigorously establishes the proximity-layer substitution that the conjecture relies on. It does not propose a new MCA statement and is intended to remove the dependency on WHIR’s Conjecture 4.12 from the folding terms; the surrounding iteration-level analysis of [11, §5] is referenced rather than re-derived, and a full proof requires verifying that this iteration-level analysis goes through under the substitution. Conceptually, the half-threshold CA acts as a one-time “crossing” from above-Johnson to below-Johnson at iteration 1’s first fold; from that point onward (all subsequent folds), the protocol stays*

inside the Johnson regime  $\delta < \delta_J$  where MCA is proved (WHIR Corollary 4.11). The  $2\times$  query overhead would be the price of this single crossing. We list the WHIR theorem-level statement as a follow-up open task (§9).

**Remark 57** (Deployment status). *The reduced-threshold safety net covers the full proximity-testing stack, with theorem-level coverage for FRI (Theorem 20, Theorem 64) and STIR in odd characteristic (Theorem 48), and conjecture-level coverage for WHIR in odd characteristic (Conjecture 54, with proximity-substitution proved); the characteristic-2 STIR/WHIR extension is recorded only as the sketch in Remark 66. This includes the odd-characteristic multiplicative-FRI settings (BabyBear, Goldilocks), circle FRI (Stwo on Mersenne31, Theorem 72), and the characteristic-2 additive-FRI setting (Binius on binary tower fields). No protocol changes are required — only the soundness analysis is updated.*

## B Extension to Characteristic 2

We show that the reduced-threshold CA bound (Theorem 6) — which is characteristic-independent — extends to the FRI folding structure over fields of characteristic 2, yielding FRI soundness above the Johnson radius for Binius-style settings on binary tower fields. The key replacement is the *additive* (subspace-polynomial) fold in place of the multiplicative  $\pm\sqrt{y}$  fold.

### B.1 Additive FRI Folding

**Definition 58** (Additive FRI Folding). *Let  $F = \mathbb{F}_{2^s}$  and  $L \subset F$  an  $\mathbb{F}_2$ -linear subspace of dimension  $d$ , so  $|L| = n = 2^d$ . Fix  $\beta \in L$ ,  $\beta \neq 0$ . The **subspace polynomial** of  $\langle\beta\rangle$  is*

$$s(x) = x^2 + \beta x.$$

*Since  $s$  is  $\mathbb{F}_2$ -linear ( $s(x+y) = s(x) + s(y)$  in characteristic 2, using  $(x+y)^2 = x^2 + y^2$ ) and  $\ker(s|_L) = \{0, \beta\}$ :*

$$L' := s(L) \subset F, \quad |L'| = n/2,$$

*and  $L'$  is again an  $\mathbb{F}_2$ -linear subspace (of dimension  $d-1$ ). Each  $y \in L'$  has exactly two preimages  $\{x, x+\beta\}$  in  $L$ .*

*For  $f : L \rightarrow F$ , the **additive even/odd decomposition** is:*

$$\begin{aligned} f_1(y) &= \frac{f(x) + f(x+\beta)}{\beta}, \\ f_0(y) &= f(x) + x \cdot f_1(y), \end{aligned}$$

*where  $x$  is either preimage of  $y$  (the expressions are invariant under  $x \mapsto x+\beta$ : swapping adds  $\beta \cdot f_1(y)$  to the  $f_0$  formula's second term and adds  $\beta \cdot f_1(y)$  to  $f(x)$ , which cancel). The **additive FRI fold** is:*

$$f'(y) = f_0(y) + \alpha \cdot f_1(y).$$

**Remark 59** (Reconstruction).  *$f(x) = f_0(s(x)) + x \cdot f_1(s(x))$  for all  $x \in L$ : at  $x$ , the right side is  $f_0(y) + x \cdot f_1(y) = (f(x) + x \cdot f_1(y)) + x \cdot f_1(y) = f(x)$ . This is the characteristic-2 analogue of the odd-characteristic reconstruction  $f(x) = f_{\text{even}}(x^2) + x \cdot f_{\text{odd}}(x^2)$ .*

## B.2 Even/Odd Isomorphism in Characteristic 2

**Lemma 60** (Additive Even/Odd Isomorphism). *For  $k$  even:  $g \mapsto (g_0, g_1)$  is a linear isomorphism  $\text{RS}[F, L, k] \xrightarrow{\sim} \text{RS}[F, L', k/2]^2$ .*

*Proof. Forward direction* ( $g \in \text{RS}_k \Rightarrow g_0, g_1 \in \text{RS}_{k/2}$ ). Let  $g(x) = \sum_{i=0}^{k-1} a_i x^i$ . In the polynomial ring  $F[x]$ , the relation  $x^2 = s + \beta x$  (from  $s(x) = x^2 + \beta x$ ) allows recursive reduction: every monomial  $x^i$  reduces to  $p_i(s) + x \cdot q_i(s)$  where  $\deg p_i, \deg q_i \leq \lfloor i/2 \rfloor$ . Collecting terms:  $g(x) = g_0(s) + x \cdot g_1(s)$  with

$$\deg g_0 \leq \lfloor (k-1)/2 \rfloor \leq k/2 - 1, \quad \deg g_1 \leq \lfloor (k-2)/2 \rfloor \leq k/2 - 1.$$

Hence  $g_0, g_1 \in \text{RS}_{k/2}(L')$ .

**Dimension count.** The map  $\varphi : \text{RS}_k(L) \rightarrow \text{RS}_{k/2}(L')^2$  is  $F$ -linear and maps a  $k$ -dimensional space into a  $(k/2 + k/2) = k$ -dimensional space.

**Injectivity.** If  $(g_0, g_1) = (0, 0)$ , then  $g(x) = g_0(s(x)) + x \cdot g_1(s(x)) = 0$  for all  $x \in L$ . Since  $g$  has degree  $< k \leq |L|$ , this implies  $g = 0$ . So  $\ker(\varphi) = \{0\}$ .

A linear injection between spaces of equal dimension is an isomorphism.  $\square$

**Remark 61** (Why  $k$  even is needed). *For  $k$  odd in characteristic 2: the monomial  $f(x) = x^k$  satisfies  $\Delta(x^k, \text{RS}_k) = 1 - \rho$  (since  $x^k - g$  has  $\leq k$  roots for any  $g$  of degree  $< k$ ). Its additive decomposition gives  $f_0, f_1$  of degree  $\leq \lceil k/2 \rceil - 1$ , so both belong to  $\text{RS}_{\lceil k/2 \rceil}$  and the fold  $f' = f_0 + \alpha \cdot f_1 \in \text{RS}_{\lceil k/2 \rceil}$  for every  $\alpha$ . The isomorphism  $\text{RS}_k \rightarrow \text{RS}_{k/2}^2$  fails because  $\text{RS}_{\lceil k/2 \rceil}^2$  has dimension  $2\lceil k/2 \rceil = k+1 > k = \dim \text{RS}_k$ : the map is injective but not surjective, so the coupling lemma breaks (not every pair  $(g_0, g_1) \in \text{RS}_{\lceil k/2 \rceil}^2$  corresponds to a codeword in  $\text{RS}_k$ ). The fold of a far-from-code function lands in the larger code  $\text{RS}_{\lceil k/2 \rceil}$  for all  $\alpha$ , and FRI is unsound. This is the same failure mode as in odd characteristic (Remark 5), confirming that  $k = 2^m$  is a genuine structural requirement.*

## B.3 Coupling Lemma and Proximity Gap

**Lemma 62** (Additive FRI Coupling). *For even  $k$  and any  $f : L \rightarrow F$  (with  $\text{char}(F) = 2$ ,  $L$  an  $\mathbb{F}_2$ -linear subspace,  $\beta \in L \setminus \{0\}$ ):*

$$\Delta(f, \text{RS}_k(L)) \leq \Delta((f_0, f_1), \text{RS}_{k/2}(L')^2).$$

*Proof.* By Lemma 60:  $g \mapsto (g_0, g_1)$  is an isomorphism  $\text{RS}_k \rightarrow \text{RS}_{k/2}^2$ . For any  $g \in \text{RS}_k$  and  $y \in L'$ : if  $f_0(y) = g_0(y)$  and  $f_1(y) = g_1(y)$ , then by reconstruction,  $f(x) = g(x)$  for both  $x$  in the fiber  $\{x, x+\beta\}$  of  $y$ . So each joint-error position  $y$  (where  $f_0(y) \neq g_0(y)$  or  $f_1(y) \neq g_1(y)$ ) contributes  $\geq 1$  error to  $f - g$  on  $L$ , and each  $y$  contributes at most 2. Let  $E = |\{y \in L' : (f_0(y), f_1(y)) \neq (g_0(y), g_1(y))\}|$ . Then  $\Delta(f, g) \leq 2E/n = E/|L'| = \Delta_{\text{joint}}$ . Taking  $\min_g$ :

$$\Delta(f, \text{RS}_k) \leq \Delta_{\text{joint}}. \quad \square$$

**Theorem 63** (Half-Threshold Proximity Gap, Characteristic 2). *(The Johnson-window hypothesis below is included only for downstream composition into Theorem 64; the proof itself uses only additive coupling plus Theorem 6.) For  $\text{char}(F) = 2$ , even  $k$ ,  $\delta_J < \delta < 1 - \rho$ , and  $\Delta(f, \text{RS}_k(L)) > \delta$ :*

$$|\{\alpha \in F : \Delta(f_0 + \alpha f_1, \text{RS}_{k/2}(L')) \leq \delta/2\}| \leq 1.$$

*Proof.* By Lemma 62:  $\Delta((f_0, f_1), \text{RS}_{k/2}^2) \geq \Delta(f, \text{RS}_k) > \delta$ . Instantiate Theorem 6 with the pair  $(f_0, f_1)$  on  $\text{RS}_{k/2}$  over  $L'$  (premise threshold  $\delta$ , conclusion threshold  $\delta/2$ ). Theorem 6 uses only linearity and Hamming distances, with no characteristic hypothesis. The proof is identical to Theorem 16.  $\square$

## B.4 FRI Soundness in Characteristic 2

**Theorem 64** (FRI Soundness Above Johnson, Characteristic 2). *For  $\text{RS}[F, L, k]$  with  $\text{char}(F) = 2$ ,  $k = 2^m$ ,  $L \subset F$  an  $\mathbb{F}_2$ -linear subspace of dimension  $d \geq m$ , and  $\delta_J < \delta < 1 - \rho$ ,  $\Delta(f^{(0)}, \text{RS}_k) > \delta$ : FRI with additive folding (Definition 58),  $R = m$  rounds, and  $q$  base-level queries (each opening one full path) satisfies*

$$\Pr[\text{FRI accepts}] \leq \frac{nR}{|F|} + (1 - \delta/2)^q.$$

*Proof.* The proof follows the same Strategy A / Strategy B structure as Theorem 20, with the multiplicative fold replaced by the additive fold. We verify the three components and note where the argument differs from odd characteristic:

**Iterability.** At each round  $i$ , the domain  $L^{(i)} = s_i(L^{(i-1)})$  is an  $\mathbb{F}_2$ -linear subspace of dimension  $d-i$ , using  $\beta_i \in L^{(i-1)} \setminus \{0\}$  as the folding direction. The code dimension halves:  $k_i = k/2^i$ . Since  $k = 2^m$ , we have  $k_i$  even for all  $i < m$  (guaranteeing the isomorphism), and  $k_m = 1$  (so  $\text{RS}_1$  is the space of constants, and the final check verifies  $f^{(R)}$  is constant).

**Strategy A.** Each round's fold  $f^{(i)}(y) = f_0^{(i-1)}(y) + \alpha_i \cdot f_1^{(i-1)}(y)$  is linear in  $\alpha_i$ , where  $f_0^{(i-1)}, f_1^{(i-1)}$  are determined by  $f^{(0)}$  and  $(\alpha_1, \dots, \alpha_{i-1})$  via the additive decomposition on the domain chain

$$L^{(0)} := L, \quad L^{(i)} := s_i(L^{(i-1)}) \quad (i \geq 1),$$

with each  $L^{(i)}$  an  $\mathbb{F}_2$ -subspace of dimension  $d-i$  (so the chain is well-defined for  $R \leq d$ ). It follows that  $f^{(R)}(y)$  is multilinear in  $(\alpha_1, \dots, \alpha_R)$  with total degree  $\leq R$ . The *nonvanishing*: if all syndromes vanished identically, then  $f^{(R)} \in \text{RS}_1$  for every  $(\alpha_1, \dots, \alpha_R)$ . As in the proof of Theorem 20, this is a polynomial-identity argument (not a runtime reconstruction); evaluating at  $\alpha_i = 0$  and  $\alpha_i = 1$  recovers  $f_0^{(i-1)}, f_1^{(i-1)} \in \text{RS}_{k_i}$ , and Lemma 60 (requiring  $k_{i-1}$  even, guaranteed by  $k = 2^m$ ) gives  $f^{(i-1)} \in \text{RS}_{k_{i-1}}$ ; iterating yields  $f^{(0)} \in \text{RS}_k$ , a contradiction. Hence some syndrome is a nonzero multilinear polynomial, and  $\Pr[f^{(R)} \in \text{RS}_1] \leq R/|F|$  by Schwartz–Zippel.

**Strategy B.** The distance-preservation argument uses Theorem 63: at round 1, all but  $\leq 1$  value of  $\alpha_1$  gives  $\Delta(f^{(1)}, \text{RS}_{k/2}) > \delta/2$ . At subsequent rounds,  $\delta/2 < (1-\rho)/2 \leq \delta_J$  (Remark 1), and the BCIKS proximity gap [7] (Theorem 1.2, Eq. (1.1)) at threshold  $\gamma = \delta/2$  gives  $\leq n$  bad  $\alpha$  per round; distance is preserved at  $\delta/2$ . The consistency check catches a  $(\delta/2)$ -far deviation with probability  $\geq 1 - (1 - \delta/2)^q$  by Lemma 19 (the additive query model is the analogue of Definition 17 with  $\pi_i$  replaced by the iterated subspace polynomial  $s_i \circ \dots \circ s_1$ , which is  $2^i$ -to-1 on its image; the i.i.d. argument transfers verbatim). The multi-deviation argument proceeds identically to Theorem 20.

Combining:  $\Pr[\text{accept}] \leq nR/|F| + (1 - \delta/2)^q$ .  $\square$

**Remark 65** (Domain structure). *In the multiplicative (odd-characteristic) setting, FRI requires  $L$  closed under negation. In the additive (characteristic-2) setting, the analogous requirement is that  $L$  is an  $\mathbb{F}_2$ -linear subspace, which is the standard domain choice in additive-FRI systems. The folding direction  $\beta_i$  at round  $i$  must satisfy  $\beta_i \in L^{(i-1)} \setminus \{0\}$ ; it is chosen dynamically from the current domain  $L^{(i-1)} = s_{i-1}(\dots(s_1(L)))$ , which is an  $\mathbb{F}_2$ -subspace of dimension  $d-i+1$ . A concrete prescription: fix an  $\mathbb{F}_2$ -basis  $\{e_1, \dots, e_d\}$  of  $L$ ; at round  $i$ , use  $\beta_i =$  any nonzero element of  $L^{(i-1)}$  (e.g., the image  $s_{i-1}(\dots(s_1(e_i)))$  if nonzero, or any other convenient choice).*

**Remark 66** (STIR/WHIR in characteristic 2). *The batch-CA theorem (Theorem 46) is characteristic-independent (it uses only the pairwise argument from Theorem 6). Combining with Lemma 62 and Theorem 63 gives the same argument for characteristic 2, so Theorem 48 and*

Corollary 50 extend with the same bounds. We record this only as a sketch here, not as a separate theorem or corollary.

## C Extension to Circle FRI (Stwo)

We extend the soundness theorem to circle FRI [22], the variant used by Stwo [23] on Mersenne31 fields ( $p = 2^{31} - 1$ ). The half-threshold CA (Theorem 6) is code-agnostic; only the coupling lemma requires adaptation to the circle evaluation domain. The first FRI round (bivariate  $\rightarrow$  univariate) uses a circle-specific coupling lemma proved below; all subsequent rounds reduce to standard univariate FRI.

### C.1 Circle Codes and Folding

**Definition 67** (Circle Group and Circle RS Code). *Let  $p$  be an odd prime. The **circle group** is  $\mathcal{C}(\mathbb{F}_p) = \{(x, y) \in \mathbb{F}_p^2 : x^2 + y^2 = 1\}$  with group law*

$$(x_0, y_0) \oplus (x_1, y_1) = (x_0x_1 - y_0y_1, x_0y_1 + y_0x_1),$$

*identity  $(1, 0)$ , and inversion  $(x, y)^{-1} = (x, -y)$ . The order is  $|\mathcal{C}(\mathbb{F}_p)| = p+1$  when  $-1$  is a quadratic non-residue ( $p \equiv 3 \pmod{4}$ ) and  $p-1$  otherwise; for Mersenne31 ( $p = 2^{31} - 1 \equiv 3 \pmod{4}$ ), the order is  $2^{31}$ .*

*Let  $D \subset \mathcal{C}(\mathbb{F}_p)$  be an evaluation domain of even size  $N$  on which the **involution**  $\iota(x, y) = (x, -y)$  is fixed-point-free (equivalently:  $y \neq 0$  for all  $(x, y) \in D$ , i.e.,  $(\pm 1, 0) \notin D$ ). Let  $D_x = \{x : (x, y) \in D\}$  denote the  $x$ -projection, with  $|D_x| = N/2$ . The involution partitions  $D$  into  $N/2$  pairs  $\{(x, y), (x, -y)\}$ , one for each  $x \in D_x$ .*

*The **circle RS code**<sup>4</sup> is*

$$\text{CRS}_N(D) = \{(x, y) \mapsto a(x) + y \cdot b(x) : \deg a < N/2, \deg b < N/2\} \upharpoonright_D,$$

*a linear code of dimension  $N$  on  $D$  (the full evaluation space, with  $\dim = |D|$ ). In the FRI setting, the code is evaluated on the larger blowup domain  $D$  and the polynomial degree bound is  $k < N$ ; we write  $\text{CRS}_k(D)$  for the subcode with  $\deg a < k/2$  and  $\deg b < k/2$ , having dimension  $k$  and rate  $\rho = k/N$  (cf. Stwo’s “blowup factor”). All rate references in the circle theorems below use this  $\rho = k/N$  convention.*

**Definition 68** (Circle FRI Folding). *For  $f : D \rightarrow F$ , the **circle even/odd decomposition** is:*

$$f_0(x) = \frac{f(x, y) + f(x, -y)}{2}, \quad f_1(x) = \frac{f(x, y) - f(x, -y)}{2y},$$

*where  $(x, y) \in D$  is either representative of the fiber over  $x$  (both expressions are invariant under  $y \mapsto -y$ ). The **circle FRI fold** is  $f'(x) = f_0(x) + \alpha \cdot f_1(x)$ , a function on  $D_x$ .*

### C.2 Coupling Lemma and Soundness

**Lemma 69** (Circle Even/Odd Isomorphism). *For even  $k$ : the map  $g \mapsto (g_0, g_1)$  is a linear isomorphism  $\text{CRS}_k(D) \xrightarrow{\sim} \text{RS}_{k/2}(D_x)^{=2}$ .*

---

<sup>4</sup>In the notation of [22], this is the FFT subspace  $L'_N(\mathbb{F}_p)$ . The full space  $L_N$  has dimension  $N+1$ ; the extra monomial  $x^{N/2}$  is handled by the prover sending its coefficient before the first fold (Remark 73).

*Proof.* For  $g(x, y) = a(x) + y \cdot b(x) \in \text{CRS}_k$ :  $g_0 = a$  and  $g_1 = b$ , both of degree  $< k/2$ , hence both in  $\text{RS}_{k/2}(D_x)$ . The map is  $F$ -linear, and the source and target both have dimension  $k$ . *Injectivity:* if  $(a, b) = (0, 0)$  on all of  $D_x$ , then  $\deg a, \deg b < k/2 \leq |D_x|$  forces  $a = b = 0$  as polynomials, so  $g = 0$  on  $D$ .  $\square$

**Lemma 70** (Circle FRI Coupling). *For even  $k$  and any  $f : D \rightarrow F$ :*

$$\Delta(f, \text{CRS}_k(D)) \leq \Delta((f_0, f_1), \text{RS}_{k/2}(D_x)^{=2}).$$

*Proof.* By Lemma 69, each  $g \in \text{CRS}_k$  corresponds to  $(g_0, g_1) \in \text{RS}_{k/2}^{=2}$ . If  $f_0(x) = g_0(x)$  and  $f_1(x) = g_1(x)$  at some  $x \in D_x$ , then  $f(x, \pm y) = a(x) \pm y \cdot b(x) = g(x, \pm y)$  (by reconstruction:  $f(x, y) = f_0(x) + y \cdot f_1(x)$ ), so *both* points in the fiber over  $x$  are error-free. Equivalently, each joint-error position  $x$  (where  $f_0(x) \neq g_0(x)$  or  $f_1(x) \neq g_1(x)$ ) contributes  $\geq 1$  base-level error on  $D$  and at most 2. With  $E = |\{x \in D_x : (f_0(x), f_1(x)) \neq (g_0(x), g_1(x))\}|$ :

$$\Delta(f, g) = \frac{|\{P \in D : f(P) \neq g(P)\}|}{N} \leq \frac{2E}{N} = \frac{E}{|D_x|} = \Delta_{\text{joint}}.$$

Taking  $\min_g$  (using the isomorphism, the minimum over  $\text{RS}_{k/2}^{=2}$  equals the minimum over  $\text{CRS}_k$ ):  $\Delta(f, \text{CRS}_k) \leq \Delta_{\text{joint}}$ .  $\square$

**Theorem 71** (Circle FRI Proximity Gap). *(The Johnson-window hypothesis below is included only for downstream composition into Theorem 72; the proof itself uses only the circle coupling lemma plus Theorem 6.) For even  $k$ ,  $\delta_J < \delta < 1 - \rho$ , and  $\Delta(f, \text{CRS}_k(D)) > \delta$ :*

$$|\{\alpha \in F : \Delta(f_0 + \alpha f_1, \text{RS}_{k/2}(D_x)) \leq \delta/2\}| \leq 1.$$

*Proof.* By Lemma 70:  $\Delta((f_0, f_1), \text{RS}_{k/2}^{=2}) \geq \Delta(f, \text{CRS}_k) > \delta$ . Apply Theorem 6 to  $(f_0, f_1)$  on  $\text{RS}_{k/2}$  over  $D_x$  (premise threshold  $\delta$ , conclusion threshold  $\delta/2$ ): at most 1 value of  $\alpha$  has  $\Delta(f_0 + \alpha f_1, \text{RS}_{k/2}) \leq \delta/2$ .  $\square$

**Theorem 72** (Circle FRI Soundness Above Johnson). *For  $\text{CRS}_k(D)$  over  $F$  with even  $k = 2^m$ ,  $D \subset \mathcal{C}(\mathbb{F}_p)$  of size  $N$  with  $\iota$  fixed-point-free and  $0 \notin D_x$  (equivalently, the 4-torsion points  $(0, \pm 1) \notin D$ ),  $\delta_J < \delta < 1 - \rho$ , and  $\Delta(f^{(0)}, \text{CRS}_k) > \delta$ : circle FRI with  $R = m$  rounds and  $q$  base-level queries (each opening one full path) satisfies*

$$\Pr[\text{circle FRI accepts}] \leq \frac{NR}{|F|} + \left(1 - \frac{\delta}{2}\right)^q.$$

*Proof.* The proof follows the Strategy A / Strategy B structure of Theorem 20.

**Round 1 (circle  $\rightarrow$  univariate).** By Theorem 71: for all but  $\leq 1$  value of  $\alpha_1$ ,  $\Delta(f^{(1)}, \text{RS}_{k/2}(D_x)) > \delta/2$ .

**Rounds 2,  $\dots$ ,  $R$  (univariate FRI on  $D_x$ ).** After round 1, the function  $f^{(1)} : D_x \rightarrow F$  lives on a univariate evaluation domain. Subsequent circle-FRI rounds use the **doubling map**  $\pi(x) = 2x^2 - 1$  (the  $x$ -coordinate of the point-doubling endomorphism on  $\mathcal{C}$ ) with involution  $x \mapsto -x$ . This involution is fixed-point-free on  $D_x$  by the standing hypothesis  $0 \notin D_x$  (the 4-torsion points  $(0, \pm 1) \in \mathcal{C}$  have order 4, and the hypothesis excludes them from  $D$ ). Stwo's domains satisfy this automatically: with  $|D| \geq 8$ , the chosen cosets avoid the 4-torsion. The even/odd decomposition

$$f_{\text{even}}(\pi(x)) = \frac{f(x) + f(-x)}{2}, \quad f_{\text{odd}}(\pi(x)) = \frac{f(x) - f(-x)}{2x}$$



uses the relation  $x^2 = (\pi(x) + 1)/2$  to reduce: for  $g \in \text{RS}_{k/2}(D_x)$ , the monomials  $x^{2j} = ((\pi + 1)/2)^j$  and  $x^{2j+1} = x \cdot ((\pi + 1)/2)^j$  give  $g_{\text{even}}, g_{\text{odd}} \in \text{RS}_{k/4}(\pi(D_x))$ , establishing the isomorphism  $\text{RS}_{k/2}(D_x) \xrightarrow{\sim} \text{RS}_{k/4}(\pi(D_x))^2$ . From this point onward, the standard univariate coupling lemma (Lemma 15) applies verbatim.

Since  $\delta/2 < (1-\rho)/2 \leq \delta_J$  (the inequality  $(1-\rho)/2 \leq 1-\sqrt{\rho}$  holds for all  $\rho \in (0,1)$ ), the effective distance  $\delta/2$  is strictly below the Johnson radius. At each round  $j \geq 2$ : Lemma 15 gives  $\Delta_{\text{joint}} > \delta/2$ , and the BCIKS proximity gap [7] at threshold  $\delta/2 \leq \delta_J$  bounds the number of bad challenges to  $\leq N/2^{j-1}$  at round  $j$ . Distance is *preserved* at  $\delta/2$  (not further halved), as in Remark 21. The BCIKS result applies because it requires only the RS code structure and a valid coupling lemma, both of which hold for the  $\pi$ -based folding.

**Strategy A.** Each round introduces  $\alpha_i$  with degree 1, so  $f^{(R)}(y)$  is multilinear in  $(\alpha_1, \dots, \alpha_R)$ . If all syndromes vanished identically, then  $f^{(R)} \in \text{RS}_1$  for every  $(\alpha_1, \dots, \alpha_R) \in F^R$ . As in the proof of Theorem 20, we extract  $f^{(0)} \in \text{CRS}_k$  via a polynomial-identity argument (not a runtime reconstruction): evaluating at  $\alpha_i = 0$  and  $\alpha_i = 1$  at each round recovers the round- $(i-1)$  even/odd projections in the appropriate codes, and Lemma 69 (round 1) and the round- $\geq 2$  isomorphisms iterate this back to  $f^{(0)} \in \text{CRS}_k$ , a contradiction. By Schwartz–Zippel:  $\Pr[f^{(R)} \in \text{RS}_1] \leq R/|F|$ .

**Strategy B.** Combining the round-1 proximity gap ( $\leq 1$  bad  $\alpha$ ), the round- $\geq 2$  distance preservation ( $\leq N/2^{j-1}$  bad  $\alpha$  at round  $j$ , hence  $\leq N$  uniformly), and the consistency check ( $\geq 1-(1-\delta/2)^q$  catch probability via Lemma 19 applied with the round- $i$  projection  $\pi_i = (\text{involution at round 1}) \circ (\text{doubling map iterated at rounds } \geq 2)$ , each  $2^i$ -to-1 from  $D$  to its round- $i$  image):  $\Pr[\text{accept}] \leq NR/|F| + (1 - \delta/2)^q$ .  $\square$

**Remark 73** (The dimension gap). *The full circle polynomial space  $L_N$  of [22] has dimension  $N+1$ ; the FFT subspace  $\text{CRS}_N$  used above has dimension  $N$ . The extra direction is the degree- $(N/2)$  monomial  $x^{N/2}$ , corresponding to the vanishing polynomial of the half-domain. In Stwo’s protocol [23], the prover sends this coefficient before the first fold, reducing the proximity test to  $\text{CRS}_N$ . This adds at most  $1/|F|$  to the commit-phase error (absorbed into the  $NR/|F|$  term).*

**Remark 74** (STIR and WHIR on circle domains). *The per-coordinate batch-CA helper (Theorem 46) is code-agnostic. Combined with Lemma 70, the STIR analysis of Theorem 48 extends to circle-FRI-based STIR with the same rate-uniform bound  $(2k+2N)R/|F| + (1-\delta/2)^q$  (simplifying to  $(2k+2)R/|F|$  at  $\rho \geq 1/2$  and tightening to  $(k+N)R/|F|$  or  $(k+2)R/|F|$  at deployment scale), and Corollary 50 extends to circle-code linear subcodes. The above-Johnson safety net now covers all deployed proximity-testing protocols: multiplicative FRI, additive FRI (characteristic 2), and circle FRI.*

## D Fiat–Shamir Compilation and Knowledge Soundness

The theorems in this paper (Theorems 20, 64, 72, 48) establish *information-theoretic* soundness of the corresponding interactive proximity-testing protocols; Conjecture 54 states the proposed WHIR analogue conditional on its iteration-level induction (proximity-layer substitution rigorous via Lemma 55). Deployed STARK systems use the non-interactive versions obtained via the Fiat–Shamir transform (replacing verifier randomness with hash evaluations). This appendix states the non-interactive soundness and knowledge-soundness theorems that follow from our bounds via standard compilation.

## D.1 Round-by-Round Soundness of FRI

The Fiat–Shamir security of FRI requires not just end-to-end soundness but *round-by-round* (RBR) *soundness* [36, 37]: the property that at each protocol round, the prover’s probability of escaping a “doomed” state is bounded, regardless of the transcript so far.

**Theorem 75** (RBR Soundness of Half-Threshold FRI, multiplicative case). *For  $\text{RS}[F, L, k]$  with  $\text{char}(F) \neq 2$ ,  $k = 2^m$ ,  $L$  closed under negation,  $\delta_J < \delta < 1 - \rho$ , and  $q$  base queries in the final consistency round: the FRI protocol with  $R$  rounds has round-by-round soundness with per-round error*

$$\varepsilon_{\text{rbr}}^{(i)} \leq \begin{cases} 1/|F| & i = 1 \text{ (commit round: } \leq 1 \text{ bad } \alpha \text{ via Theorem 16),} \\ n/|F| & 2 \leq i \leq R \text{ (commit rounds: BCIKS at } \delta/2 \leq \delta_J), \\ (1 - \delta/2)^q & i = R + 1 \text{ (query round).} \end{cases}$$

*Proof.* The doomed predicate at round  $i$  is defined on the deterministic *honest-fold reference*  $\hat{f}^{(i)}$  (computed from  $f^{(0)}$  and  $\alpha_1, \dots, \alpha_i$  alone, independent of the prover’s commits; see Step 1 of the proof of Theorem 20): the partial transcript is doomed at round  $i$  iff  $\Delta(\hat{f}^{(i)}, \text{RS}_{k/2^i}) > \delta/2$ . Initially ( $i = 0$ ):  $\Delta(\hat{f}^{(0)}, \text{RS}_k) = \Delta(f^{(0)}, \text{RS}_k) > \delta > \delta/2$ , so every false statement starts doomed.

*Round 1 (commit).* By the proximity gap (Theorem 16): at most 1 value of  $\alpha_1$  makes the honest fold  $\hat{f}^{(1)} = \text{fold}(\hat{f}^{(0)}, \alpha_1)$  become  $(\delta/2)$ -close to  $\text{RS}_{k/2}$ . Escape probability:  $\leq 1/|F|$  (sharper than the round- $\geq 2$  bound of  $n/|F|$  stated in the theorem).

*Rounds 2,  $\dots$ ,  $R$  (commit).* Under the inductive doomed hypothesis at round  $i-1$ :  $\Delta(\hat{f}^{(i-1)}, \text{RS}_{k/2^{i-1}}) > \delta/2$  together with  $\delta/2 < (1-\rho)/2 \leq \delta_J$  (Remark 1) places the threshold inside BCIKS’s zero-loss regime. By Lemma 15, the joint distance  $\Delta_{\text{joint}}((\hat{f}_{\text{even}}^{(i-1)}, \hat{f}_{\text{odd}}^{(i-1)}), \text{RS}_{k/2^i}^2) > \delta/2$ , and the BCIKS proximity gap [7] (Theorem 1.2, Eq. (1.1)) at threshold  $\gamma = \delta/2$  gives: at most  $n$  values of  $\alpha_i$  make  $\Delta(\hat{f}^{(i)}, \text{RS}_{k/2^i}) \leq \delta/2$ . Escape probability:  $\leq n/|F|$ .

*Round  $R+1$  (query).* Conditional on doomed at round  $R$  ( $\Delta(\hat{f}^{(R)}, \text{RS}_{k/2^R}) > \delta/2$ ), the FRI tree-consistency lemma [20, §6], [7, Lemma 8.2] gives that the merged catching set  $T := \bigcup_i \pi_i^{-1}(B_i)$  (where  $B_i$  is the prover’s round- $i$  deviation set, as in Step 2 of the proof of Theorem 20) has relative measure  $|T|/|L| \geq \delta/2$  whenever  $f^{(R)} \in \text{RS}_{k/2^R}$ . By Lemma 19, the probability that all  $q$  i.i.d. base-level queries miss  $T$  is  $\leq (1 - \delta/2)^q$ .

The doomed predicate at commit rounds is defined on the deterministic honest-fold reference  $\hat{f}^{(i)}$ , so the per-round commit bounds depend on the verifier’s challenges only and are independent of the prover’s commits. The query-round set  $T$  does depend on the prover’s commits, but the lower bound  $|T|/|L| \geq \delta/2$  is uniform over all prover commit patterns (any prover whose final commit is a codeword  $f^{(R)} \in \text{RS}_{k/2^R}$  under doomed  $\hat{f}^{(R)}$ ). Hence the protocol is RBR-sound with the stated errors.  $\square$

## D.2 Non-Interactive FRI via Fiat–Shamir

The BCS transformation [35] compiles a public-coin IOP with RBR soundness into a non-interactive argument in the random oracle model. The security loss is *linear* in the number of random-oracle queries (not exponential in rounds), precisely because RBR soundness enables straight-line simulation.

**Theorem 76** (Non-Interactive FRI Soundness). *Under the hypotheses of Theorem 75, let  $H : \{0, 1\}^* \rightarrow \{0, 1\}^\kappa$  be a random oracle. For any adversary making at most  $Q$  queries to  $H$ , the*

*Fiat-Shamir-compiled FRI protocol satisfies:*

$$\varepsilon_{\text{FRI}}^{\text{NI}} \leq Q \cdot \left( \frac{nR}{|F|} + \left(1 - \frac{\delta}{2}\right)^q \right) + \frac{3(Q^2 + 1)}{2^\kappa}.$$

*Proof.* By the BCS compilation theorem [35, 37]: for a public-coin IOP with RBR soundness error  $\varepsilon_{\text{rbr}}$ , the Fiat-Shamir-compiled protocol has soundness error at most  $Q \cdot \varepsilon_{\text{rbr}} + 3(Q^2 + 1)/2^\kappa$ , where the second term is the Merkle-tree binding error. Substituting  $\varepsilon_{\text{rbr}} \leq nR/|F| + (1 - \delta/2)^q$  from Theorem 75 (summing the per-round errors:  $R \cdot n/|F| + (1 - \delta/2)^q$ ) gives the stated bound.  $\square$

**Remark 77** (Practical parameters). *For BabyBear<sub>4</sub> ( $|F| = p^4 \approx 2^{124}$ ),  $n = 2^{20}$ ,  $R = 20$ ,  $\kappa = 256$  (SHA-256), and  $Q = 2^{64}$  (a generous query bound):*

- *Interactive FRI error:*  $nR/|F| + (1 - \delta/2)^q \approx 2^{-100} + (1 - \delta/2)^q$ .
- *Fiat-Shamir overhead:*  $Q \cdot \varepsilon \approx 2^{64} \cdot \varepsilon$ , reducing security by  $\sim 64$  bits.
- *Merkle binding:*  $3Q^2/2^\kappa \approx 2^{129}/2^{256} \approx 2^{-127}$  (negligible).
- *Net:* the interactive bound loses  $\sim 64$  bits in the non-interactive setting. At  $q \approx 398$  (128-bit interactive query-phase security for  $\delta = 0.4$ ), the non-interactive query contribution is only  $\sim 64$  bits against  $Q = 2^{64}$  adversaries, and the total is further capped by the commit term unless the challenge field is enlarged.

*This matches the standard parameter regime: deployed systems set  $\kappa \geq 2\lambda$  and target  $\lambda$ -bit security against  $Q \leq 2^\lambda$  adversaries.*

*Tighter bound.* The BCS compilation [37] applies  $Q$ -loss to the per-round max of the RBR error, not the total. Since each commit round has escape probability  $\leq n/|F|$  (Theorem 75), a tighter bound is  $\varepsilon_{\text{FRI}}^{\text{NI}} \leq Q \cdot \max(n/|F|, (1 - \delta/2)^q) + 3(Q^2 + 1)/2^\kappa$ , saving a factor of  $R$  in the commit term. We state the looser bound above for consistency with the interactive Theorem 20.

### D.3 Knowledge Soundness and DEEP Polynomial Commitments

FRI is deployed not as a standalone proximity test but as the core of a *polynomial commitment scheme (PCS)* via the DEEP technique [38]: to verify a claimed evaluation  $\hat{f}(\zeta) = v$  at a random point  $\zeta \in F$ , the verifier checks that the quotient  $q(x) = (f(x) - v)/(x - \zeta)$  is close to  $\text{RS}_{k-1}$ .

**Theorem 78** (DEEP-FRI Knowledge Soundness, multiplicative case). *For  $\text{RS}[F, L, k]$  with  $\text{char}(F) \neq 2$ ,  $k = 2^m$ ,  $L$  closed under negation,  $\delta_J < \delta < 1 - \rho$ , and the underlying multiplicative FRI subprotocol of Theorem 20 with  $R = m$  rounds and  $q$  queries: the DEEP-FRI polynomial commitment scheme [38] has knowledge soundness*

$$\varepsilon_{\text{PCS}} \leq \frac{k}{|F|} + \frac{nR}{|F|} + \left(1 - \frac{\delta}{2}\right)^q,$$

*where the  $k/|F|$  term is the DEEP quotient error (probability that the random evaluation point  $\zeta$  is degenerate).*

*Proof.* The DEEP technique [38] reduces polynomial commitment to proximity testing: given oracle access to  $f$  and a claimed evaluation  $f(\zeta) = v$  at a random  $\zeta \in F \setminus L$ :

1. If  $f \in \text{RS}_k$  and  $f(\zeta) = v$ : the quotient  $q(x) = (f(x) - v)/(x - \zeta)$  satisfies  $q \in \text{RS}_{k-1}$ , and FRI accepts.
2. If  $f \in \text{RS}_k$  but  $f(\zeta) \neq v$ : the polynomial  $f(x) - v$  is not divisible by  $(x - \zeta)$  (since  $f(\zeta) - v \neq 0$ ), so the function  $q(x) = (f(x) - v)/(x - \zeta)$  on  $L$  (well-defined since  $\zeta \notin L$ ) is not the evaluation of any degree- $<(k-1)$  polynomial. Hence  $\Delta(q, \text{RS}_{k-1}) \geq 1 - (k-1)/n = 1 - \rho + 1/n > \delta$ , and FRI rejects with high probability.

3. If  $f \notin \text{RS}_k$ : the verifier sends random  $\zeta \in F \setminus L$  and the prover provides a claimed value  $v$ . By the DEEP worst-case-to-average-case reduction [38]: for all but  $\leq k$  degenerate values of  $\zeta$ , the quotient function  $q(x) = (f(x) - v)/(x - \zeta)$  on  $L$  is far from  $\text{RS}_{k-1}$ , regardless of the prover's choice of  $v$ . Probability of a degenerate  $\zeta$ :  $k/|F|$ . Conditional on non-degenerate  $\zeta$ : FRI catches the far quotient via Theorem 20.

Taking the union bound:  $\varepsilon_{\text{PCS}} \leq k/|F| + nR/|F| + (1 - \delta/2)^q$ .

*Knowledge extraction.* By the round-by-round framework [37]: the extractor runs the interactive protocol with the adversary, catching the adversary's committed polynomial via the DEEP evaluation at the random point  $\zeta$ . The RBR knowledge error is  $\varepsilon_{\text{rbr}} + k/|F|$ , and the BCS compilation preserves knowledge soundness with the same  $Q$ -linear loss as for plain soundness.  $\square$

**Corollary 79** (Non-Interactive PCS Knowledge Soundness). *Under the hypotheses of Theorem 78, with the Fiat–Shamir / Merkle-tree compiler instantiated by a random oracle  $H : \{0, 1\}^* \rightarrow \{0, 1\}^\kappa$ , the compiled DEEP-FRI PCS has knowledge soundness against  $Q$ -query adversaries:*

$$\varepsilon_{\text{PCS}}^{\text{NI}} \leq Q \cdot \left( \frac{k + nR}{|F|} + \left( 1 - \frac{\delta}{2} \right)^q \right) + \frac{3(Q^2 + 1)}{2^\kappa}.$$

**Remark 80** (Scope of the knowledge-soundness claim). *Theorem 78 and Corollary 79 give knowledge soundness for the DEEP-FRI polynomial commitment scheme: the extractor recovers a degree- $< k$  polynomial consistent with the committed oracle. The full SNARK (e.g., Plonk with FRI-PCS) additionally requires showing that the extracted polynomial satisfies the circuit constraints; this is the arithmetization step, which is protocol-specific and orthogonal to the proximity-testing soundness. Our contribution is the proximity-testing layer; the arithmetization-to-PCS composition follows standard arguments from [37].*

## References

- [1] Ethereum Foundation, *The Proximity Prize: \$1M Reed–Solomon Challenge* (Grand MCA Challenge and Grand List Decoding Challenge with  $\varepsilon^* = 2^{-128}$  at rates  $\rho \in \{1/2, 1/4, 1/8, 1/16\}$ ), <https://proximityprize.org>, accessed May 2026.
- [2] R. Chai, X. Fan, *Action–Orbit Approaches to Reed–Solomon Proximity on Cyclic Domains*, companion paper in preparation, 2026.
- [3] R. Chai, X. Fan, *Codim-2( $c-1$ ) Reed–Solomon Proximity Bounds via the Berlekamp Realizer*, companion paper in preparation, 2026.
- [4] G. Arnon, D. Boneh, G. Fenzi, *Open problems in list decoding and correlated agreement*, ePrint 2026/680.
- [5] E. Crites, A. Stewart, *On Reed–Solomon proximity gaps conjectures*, ePrint 2025/2046.
- [6] E. Ben-Sasson, D. Carmon, U. Haböck, S. Kopparty, S. Saraf, *On proximity gaps for Reed–Solomon codes*, ePrint 2025/2055, 2025.
- [7] E. Ben-Sasson, D. Carmon, Y. Ishai, S. Kopparty, S. Saraf, *Proximity gaps for Reed–Solomon codes*, FOCS 2020.
- [8] R. Goyal, V. Guruswami, *Optimal proximity gaps for subspace-design codes and (random) Reed–Solomon codes*, ePrint 2025/2054.

- [9] F. G. Jeronimo, L. Liu, P. Rajpal, *Optimal proximity gap for folded Reed–Solomon codes via subspace designs*, arXiv:2601.10047, 2026.
- [10] G. Arnon, A. Chiesa, G. Fenzi, E. Yegorov, *STIR: Reed–Solomon proximity testing with fewer queries*, CRYPTO 2024; ePrint 2024/390.
- [11] G. Arnon, A. Chiesa, G. Fenzi, E. Yegorov, *WHIR: Reed–Solomon proximity testing with super-fast verification*, ePrint 2024/1586.
- [12] A. Kambiré, *Proximity gaps conjecture fails near capacity over prime fields*, arXiv:2604.09724, 2026.
- [13] R. Lidl, H. Niederreiter, *Finite Fields*, Cambridge, 2nd ed., 1997.
- [14] P. Frankl, R. M. Wilson, *Intersection theorems with geometric consequences*, Combinatorica 1(4):357–368, 1981.
- [15] M. Sudan, *Decoding of Reed–Solomon codes beyond the error-correction bound*, J. Complexity 13(1):180–193, 1997.
- [16] V. Guruswami, M. Sudan, *Improved decoding of Reed–Solomon and algebraic-geometric codes*, IEEE Trans. Inform. Theory 45(6):1757–1767, 1999.
- [17] T. Helleseth, *Some results about the cross-correlation function between two maximal linear sequences*, Discrete Math. 16(3):209–232, 1976.
- [18] S. W. Golomb, G. Gong, *Signal Design for Good Correlation: For Wireless Communication, Cryptography, and Radar*, Cambridge University Press, 2005.
- [19] O. Moreno, C. J. Moreno, *Exponential sums and Goppa codes: I*, Proc. Amer. Math. Soc. 111(2):523–531, 1991.
- [20] E. Ben-Sasson, I. Bentov, Y. Horesh, M. Riabzev, *Fast Reed–Solomon interactive oracle proofs of proximity*, in ICALP 2018, LIPIcs vol. 107, 2018. Full version: ePrint 2017/134.
- [21] B. E. Diamond, J. Posen, *Succinct Arguments over Towers of Binary Fields*, IACR ePrint 2023/1784, 2023.
- [22] U. Haböck, D. Levit, S. Papini, *Circle STARKs*, ePrint 2024/278, 2024.
- [23] D. Carmon, S. Goldberg, U. Haböck, V. Lerer, S. Lesokhin, *Stwo whitepaper*, ePrint 2026/532, 2026.
- [24] R. Chai, X. Fan, *Structural results on per-word list size for Reed–Solomon codes*, notes/0088-structural-list-size.md in [39], 2026.
- [25] R. Chai, X. Fan, *Equal-threshold CA bound:  $\varepsilon_{\text{ca}} = \binom{n}{w}/|F|$  via erasure-pattern counting*, notes/0089-op1-resolution.md in [39], 2026.
- [26] R. Chai, X. Fan, *Birthday-bound at codimension excess  $c = 2$ : Bernstein concentration via Möbius exclusion, and the  $k$ -wise common-core obstruction*, notes/0092-open-set-rank-lemma-failure.md in [39], 2026.

- [27] R. Chai, X. Fan, *Open-Set Rank Lemma at codimension excess  $c \geq 3$ : phase diagram, triangle and tetrahedron obstructions at small primes*, `notes/0094-proof-strategy-c3.md` in [39], 2026.
- [28] S. Janson, *On concentration of probability*, in Contemporary Combinatorics, Bolyai Soc. Math. Stud., vol. 10, Springer, 2002, pp. 289–301.
- [29] D. Dubhashi, D. Ranjan, *Balls and bins: A study in negative dependence*, Random Structures & Algorithms, vol. 13, no. 2, 1998, pp. 99–124.
- [30] K. Joag-Dev, F. Proschan, *Negative association of random variables, with applications*, Annals of Statistics, vol. 11, no. 1, 1983, pp. 286–295.
- [31] G. N. Rothblum, S. Vadhan, A. Wigderson, *Interactive proofs of proximity: delegating computation in sublinear time*, in STOC 2013. Full version: ECCC TR12-182.
- [32] S. Ames, C. Hazay, Y. Ishai, M. Venkatasubramanian, *Ligero: Lightweight sublinear arguments without a trusted setup*, in Proceedings of the 24th ACM Conference on Computer and Communications Security (CCS), 2017.
- [33] R. M. Roth and G. Zémor, personal communication, 2018. Cited in [7].
- [34] E. Ben-Sasson, S. Kopparty, S. Saraf, *Worst-case to average-case reductions for the distance to a code*, in 33rd Computational Complexity Conference (CCC), pages 24:1–24:23, 2018.
- [35] E. Ben-Sasson, A. Chiesa, N. Spooner, *Interactive oracle proofs*, in TCC 2016-B, LNCS 9986, pages 31–60, 2016.
- [36] R. Canetti, Y. Chen, J. Holmgren, A. Lombardi, G. N. Rothblum, R. D. Rothblum, D. Wichs, *Fiat–Shamir: from practice to theory*, in STOC 2019, pages 1082–1090, 2019.
- [37] A. Block, A. Garreta, J. Katz, J. Thaler, P. N. Tiwari, M. Zajac, *Fiat–Shamir security of FRI and related SNARKs*, in Asiacrypt 2023, LNCS 14444, pages 3–40, 2023.
- [38] E. Ben-Sasson, L. Goldberg, S. Kopparty, S. Saraf, *DEEP-FRI: sampling outside the box improves soundness*, ePrint 2019/336, 2019.
- [39] R. Chai, X. Fan, *Computational companion: Lean 4 formalization and verification scripts for FRI soundness above the Johnson bound*, <https://github.com/iotexproject/rs-proximity-gaps>, 2026.

## Acknowledgments

We thank Professor Guang Gong (University of Waterloo) for reviewing an earlier draft and for the sequence-school perspective, and the Ethereum Foundation for posing the Proximity Prize that motivated this work.




Figure 1: Proof map. **Top half** (mainline): RVW13 yields the half-threshold CA bound (Theorem 6), which composes with the FRI even/odd coupling to give the round-1 proximity gap (Theorem 16); BCIKS supplies the round- $\geq 2$  distance lock (the structural observation  $\delta/2 < (1-\rho)/2$ ); the two combine into the FRI soundness theorem. **Bottom half** (parallel branches from the same CA bound): STIR via batch CA (Theorem 48); WHIR via batch CA on linear subcodes (Corollary 50 proves the proximity-layer; Conjecture 54 states the end-to-end soundness conjecture); characteristic-2 FRI via additive coupling; circle FRI via the Stwo coupling; equal-threshold CA (the OP1 ceiling) showing the  $\binom{n}{w}/|F|$  bound is vacuous at FRI scale, so the  $2\times$  overhead is intrinsic to any CA-based proof. The list-size structural track (§7) is independent of the soundness proof and is not shown.